ICML 2020 Accepted Papers 1088

If you are an author on a paper here and your institution is missing, you should immediately update your CMT profile and the corresponding profile at https://ICML.cc/EditProfile.

Please read this guide on how to design for color blindness for tips to make your poster friendlier to the color-blind.

Reverse-engineering deep ReLU networks
David Rolnick (University of Pennsylvania) · Konrad Kording (Upenn)

My Fair Bandit: Distributed Learning of Max-Min Fairness with Multi-player Bandits
Ilai Bistritz (Stanford University) · Tavor Baharav (Stanford University) · Amir Leshem (Bar-Ilan University) · Nicholas Bambos ()

Scalable Differentiable Physics for Learning and Control
Yi-Ling Qiao (University of Maryland, College Park) · Junbang Liang (University of Maryland, College Park) · Vladlen Koltun (Intel Labs) · Ming Lin (UMD-CP & UNC-CH)

Generalization to New Actions in Reinforcement Learning
Ayush Jain (University of Southern California) · Andrew Szot (University of Southern California) · Joseph Lim (Univ. of Southern California)

Randomized Block-Diagonal Preconditioning for Parallel Learning
Celestine Mendler-Dünner (University of California, Berkeley) · Aurelien Lucchi (ETH Zurich)

Stochastic Flows and Geometric Optimization on the Orthogonal Group
Krzysztof Choromanski (Google Brain Robotics) · Valerii Likhosherstov (University of Cambridge) · Jared Q Davis (Google Research) · David Cheikhi (Columbia University) · Achille Nazaret (Columbia University) · Xingyou Song (Google Brain) · Achraf Bahamou (Columbia University) · Jack Parker-Holder (University of Oxford) · Mrugank Akarte (Columbia University) · YUAN GAO (Columbia University) · Jacob Bergquist (Columbia University) · Aldo Pacchiano (UC Berkeley) · Vikas Sindhwani (Google) · Tamas Sarlos (Google) · Adrian Weller (University of Cambridge, Alan Turing Institute)

PackIt: A Virtual Environment for Geometric Planning
Ankit Goyal (Princeton University) · Jia Deng (Princeton University)

Soft Threshold Weight Reparameterization for Learnable Sparsity
Aditya Kusupati (University of Washington) · Vivek Ramanujan (Allen Institute for Artificial Intelligence) · Raghav Somani (University of Washington) · Mitchell Wortsman (University of Washington) · Prateek Jain (Microsoft Research) · Sham Kakade (University of Washington) · Ali Farhadi (University of Washington, Allen Institue for AI)

Stochastic Latent Residual Video Prediction
Jean-Yves Franceschi (Sorbonne Université) · Edouard Delasalles (Sorbonne Université) · Mickael Chen (Sorbonne Université) · Sylvain Lamprier (LIP6 - Sorbonne Universités) · Patrick Gallinari (Sorbonne Universite, Criteo AI Lab)

Fractional Underdamped Langevin Dynamics: Retargeting SGD with Momentum under Heavy-Tailed Gradient Noise
Umut Simsekli (Institut Polytechnique de Paris / University of Oxford) · Lingjiong Zhu (FSU) · Yee Whye Teh (Oxford and DeepMind) · Mert Gurbuzbalaban (Rutgers University)

Context Aware Local Differential Privacy
Jayadev Acharya (Cornell University) · Keith Bonawitz (Google) · Peter Kairouz (Google) · Daniel Ramage (Google) · Ziteng Sun (Cornell/Google)

Privately Learning Markov Random Fields
Gautam Kamath (University of Waterloo) · Janardhan Kulkarni (Microsoft Research, Redmond) · Steven Wu (University of Minnesota) · Huanyu Zhang (Cornell University)

A Mean Field Analysis Of Deep ResNet And Beyond: Towards Provably Optimization Via Overparameterization From Depth
Yiping Lu (Stanford University) · Chao Ma (Princeton University) · Yulong Lu (Duke University) · Jianfeng Lu (Duke University) · Lexing Ying (Stanford University)

Provable Smoothness Guarantees for Black-Box Variational Inference
Justin Domke (University of Massachusetts, Amherst)

Enhancing Simple Models by Exploiting What They Already Know
Amit Dhurandhar (IBM Research) · Karthikeyan Shanmugam (IBM Research, T. J. Watson Research Center) · Ronny Luss (IBM Research)

Fiduciary Bandits
Gal Bahar (Technion – Israel Institute of Technology) · Omer Ben-Porat (Technion--Israel Institute of Technology) · Kevin Leyton-Brown (University of British Columbia) · Moshe Tennenholtz (Technion – Israel Institute of Technology)

Training Deep Energy-Based Models with f-Divergence Minimization
Lantao Yu (Stanford University) · Yang Song (Stanford University) · Jiaming Song (Stanford) · Stefano Ermon (Stanford University)

Progressive Graph Learning for Open-Set Domain Adaptation
Yadan Luo (University of Queensland) · Zijian Wang (University of Queensland) · Mahsa Baktashmotlagh (University of Queensland) · Zi Huang (University of Queensland)

Learning De-biased Representations with Biased Representations
Hyojin Bahng (Korea University) · SANGHYUK CHUN (Naver corp.) · Sangdoo Yun (Clova AI Research, NAVER Corp.) · Jaegul Choo (Korea University) · Seong Joon Oh (Clova AI Research, NAVER Corp.)

Generalized Neural Policies for Relational MDPs
Sankalp Garg (Indian Institute of Technology Delhi) · Aniket Bajpai (Indian Institute of Technology, Delhi) · Mausam (IIT Delhi)

Feature-map-level Online Adversarial Knowledge Distillation
Inseop Chung (Seoul National University) · Seonguk Park (Seoul National University) · Kim Jangho (Seoul National University) · NOJUN KWAK (Seoul National University)

DRWR: A Differentiable Renderer without Rendering for Unsupervised 3D Structure Learning from Silhouette Images
Zhizhong Han (University of Maryland, College Park) · Chao Chen (Tsinghua University) · Yu-Shen Liu (Tsinghua University) · Matthias Zwicker (University of Maryland)

Towards Accurate Post-training Network Quantization via Bit-Split and Stitching
Peisong Wang (Institute of Automation, Chinese Academy of Sciences) · Qiang Chen (CASIA) · Xiangyu He (CASIA) · Jian Cheng ("Chinese Academy of Sciences, China")

Hybrid Stochastic-Deterministic Minibatch Proximal Gradient: Less-Than-Single-Pass Optimization with Nearly Optimal Generalization
Pan Zhou (Salesforce) · Xiao-Tong Yuan (Nanjing University of Information Science & Technology)

Reserve Pricing in Repeated Second-Price Auctions with Strategic Bidders
Alexey Drutsa (Yandex)

On Gradient Descent Ascent for Nonconvex-Concave Minimax Problems
Tianyi Lin (UC Berkeley) · Chi Jin (Princeton University) · Michael Jordan (UC Berkeley)

Learning Binary Neurons with Noisy Supervision
Kai Han (Noah’s Ark Lab, Huawei Technologies) · Yunhe Wang (Peking University) · Yixing Xu (Huawei Technologies) · Chunjing Xu (Huawei Noah's Ark Lab) · Enhua Wu (CAS) · Chang Xu (University of Sydney)

Stochastic Frank-Wolfe for Constrained Finite-Sum Minimization
Geoffrey Negiar (UC Berkeley) · Gideon Dresdner (ETH Zürich) · Alicia Yi-Ting Tsai (University of California, Berkeley) · Laurent El Ghaoui (UC Berkeley) · Francesco Locatello (ETH Zurich - Max Planck Institute) · Fabian Pedregosa (Google)

Do We Really Need to Access the Source Data? Source Hypothesis Transfer for Unsupervised Domain Adaptation
Jian Liang (NUS) · Dapeng Hu (NUS) · Jiashi Feng (National University of Singapore)

Acceleration through spectral density estimation
Fabian Pedregosa (Google) · Damien Scieur (INRIA - ENS)

Graph Structure of Neural Networks
Jiaxuan You (Stanford University) · Kaiming He (Facebook AI Research) · Jure Leskovec (Stanford University) · Saining Xie (Facebook AI Research)

Optimal Continual Learning has Perfect Memory and is NP-hard
Jeremias Knoblauch (Warwick University) · Hisham Husain (Australian National University) · Tom Diethe (Amazon)

Clinician-in-the-Loop Decision Making: Reinforcement Learning with Near-Optimal Set-Valued Policies
Shengpu Tang (University of Michigan) · Aditya Modi (University of Michigan) · Michael Sjoding (University of Michigan) · Jenna Wiens (University of Michigan)

Computational and Statistical Tradeoffs in Inferring Combinatorial Structures of Ising Model
Ying Jin (Stanford University) · Zhaoran Wang (Northwestern) · Junwei Lu ()

On the Number of Linear Regions of Convolutional Neural Networks
Huan Xiong (Mohamed bin Zayed University of Artificial Intelligence (MBZUAI)) · Lei Huang (Inception Institute of Artificial Intelligence) · Mengyang Yu (Inception Institute of Artificial Intelligence) · Li Liu (the inception institute of artificial intelligence) · Fan Zhu (Inception Institute of Artificial Intelligence) · Ling Shao (Inception Institute of Artificial Intelligence)

Deep Streaming Label Learning
Zhen Wang (University of Sydney) · Liu Liu (The University of Sydney) · Dacheng Tao (The University of Sydney)

From Importance Sampling to Doubly Robust Policy Gradient
Jiawei Huang (University of Illinois at Urbana-Champaign) · Nan Jiang (University of Illinois at Urbana-Champaign)

Loss Function Search for Face Recognition
Xiaobo Wang (JD AI Research) · Shuo Wang (JD AI Research) · Shifeng Zhang (CBSR, NLPR, CASIA) · Cheng Chi (University of Chinese Academy of Sciences) · Tao Mei (AI Research of JD.com)

Breaking the Curse of Space Explosion: Towards Efficient NAS with Curriculum Search
Yong Guo (South China University of Technology) · Yaofo Chen (South China University of Technology) · Yin Zheng (Tencent AI Lab) · Peilin Zhao (Artificial Intelligence Department, Ant Financial) · Jian Chen ("South China University of Technology, China") · Junzhou Huang (University of Texas at Arlington / Tencent AI Lab) · Mingkui Tan (South China University of Technology)

Automatic Reparameterisation of Probabilistic Programs
Maria Gorinova (University of Edinburgh) · Dave Moore (Google) · Matthew Hoffman (Google)

Kernel Methods for Cooperative Multi-Agent Learning with Delays
Abhimanyu Dubey (Massachusetts Institute of Technology) · Alex `Sandy' Pentland (MIT)

Robust Multi-Agent Decision-Making with Heavy-Tailed Payoffs
Abhimanyu Dubey (Massachusetts Institute of Technology) · Alex `Sandy' Pentland (MIT)

Learning the Valuations of a k -demand Agent
Hanrui Zhang (Duke University) · Vincent Conitzer (Duke)

Rigging the Lottery: Making All Tickets Winners
Utku Evci (Google) · Trevor Gale (Google Brain) · Jacob Menick (DeepMind) · Pablo Samuel Castro (Google Brain) · Erich Elsen (Google)

Active Learning on Attributed Graphs via Graph Cognizant Logistic Regression and Preemptive Query Generation
Florence Regol (McGill University) · Soumyasundar Pal (McGill University) · Yingxue Zhang (Huawei Technologies Canada) · Mark Coates (McGill University)

Performative Prediction
Juan Perdomo (University of California, Berkeley) · Tijana Zrnic (University of California, Berkeley) · Celestine Mendler-Dünner (University of California, Berkeley) · University of California Moritz Hardt (University of California, Berkeley)

On Layer Normalization in the Transformer Architecture
Ruinbin Xiong (Institute of Computing Technology) · Yunchang Yang (Peking University) · Di He (Peking University) · Kai Zheng (Peking University) · Shuxin Zheng (microsoft.com) · Chen Xing (Nankai University) · Huishuai Zhang (Microsoft) · Yanyan Lan (Institute of Computing Technology) · Liwei Wang (Peking University) · Tie-Yan Liu (Microsoft Research Asia)

The many Shapley values for model explanation
Mukund Sundararajan (Google Inc.) · Amir Najmi (Google)

Linear Convergence of Randomized Primal-Dual Coordinate Method for Large-scale Linear Constrained Convex Programming
Daoli Zhu (Shanghai Jiao Tong University) · Lei Zhao (Shanghai Jiao Tong University)

New Oracle-Efficient Algorithms for Private Synthetic Data Release
Giuseppe Vietri (University of Minnesota) · Steven Wu (University of Minnesota) · Mark Bun (Princeton University) · Thomas Steinke (IBM, Almaden) · Grace Tian (Harvard)

Oracle Efficient Private Non-Convex Optimization
Seth Neel (University of Pennsylvania) · Aaron Roth (University of Pennsylvania) · Giuseppe Vietri (University of Minnesota) · Steven Wu (University of Minnesota)

Universal Asymptotic Optimality of Polyak Momentum
Damien Scieur (INRIA - ENS) · Fabian Pedregosa (Google)

Adversarial Robustness via Runtime Masking and Cleansing
Yi-Hsuan Wu (National Tsing Hua University) · Chia-Hung Yuan (National Tsing Hua University) · Shan-Hung (Brandon) Wu (National Tsing Hua University)

Implicit Euler Skip Connections: Enhancing Adversarial Robustness via Numerical Stability
Mingjie Li (Peking University) · Lingshen He (Peking University) · Zhouchen Lin (Peking University)

Best Arm Identification for Cascading Bandits in the Fixed Confidence Setting
Zixin Zhong (National University of Singapore) · Wang Chi Cheung (National University of Singapore) · Vincent Tan (National University of Singapore)

Robustness to Programmable String Transformations via Augmented Abstract Training
Yuhao Zhang (University of Wisconsin-Madison) · Aws Albarghouthi (University of Wisconsin-Madison) · Loris D'Antoni (University of Wisconsin-Madison)

The Complexity of Finding Stationary Points with Stochastic Gradient Descent
Yoel Drori (Google Research) · Ohad Shamir (Weizmann Institute of Science)

Sample Complexity Bounds for 1-bit Compressive Sensing and Binary Stable Embeddings with Generative Priors
Zhaoqiang Liu (NUS) · Selwyn Gomes (BITS Pilani, K K Birla Goa Campus) · Avtansh Tiwari (IIT Kanpur) · Jonathan Scarlett (National University of Singapore)

Class-Weighted Classification: Trade-offs and Robust Approaches
Ziyu Xu (Carnegie Mellon University) · Chen Dan (Carnegie Mellon University) · Justin Khim (Carnegie Mellon University) · Pradeep Ravikumar (Carnegie Mellon University)

Neural Architecture Search in a Proxy Validation Loss Landscape
Yanxi Li (University of Sydney) · Minjing Dong (The University of Sydney) · Yunhe Wang (Huawei Noah's Ark Lab) · Chang Xu (University of Sydney)

Almost Tune-Free Variance Reduction
Bingcong Li (University of Minnesota) · Lingda Wang (University of Illinois at Urbana-Champaign) · Georgios B. Giannakis (University of Minnesota)

Uniform Convergence of Rank-weighted Learning
Liu Leqi (Carnegie Mellon University) · Justin Khim (Carnegie Mellon University) · Adarsh Prasad (Carnegie Mellon University) · Pradeep Ravikumar (Carnegie Mellon University)

Parallel Machine Translation with Disentangled Context Transformer
Jungo Kasai (University of Washington) · James Cross (Facebook) · Marjan Ghazvininejad (Facebook AI Research) · Jiatao Gu (Facebook AI Research)

More Information Supervised Probabilistic Deep Face Embedding Learning
Ying Huang (huya dopamine team) · Shangfeng Qiu (Guangzhou Huya Information Technologies Co., Limited) · Wenwei Zhang (Guangzhou Huya Information Technologies Co., Limited) · Xianghui Luo (Guangzhou Huya Information Technologies Co., Limited) · Jinzhuo Wang (University of Oxford)

Parameter-Free Learning for Evolving Markov Decision Processes: The Blessing of (More) Optimism
Wang Chi Cheung (National University of Singapore) · David Simchi-Levi (MIT) · Ruihao Zhu (MIT)

Improved Sleeping Bandits with Stochastic Action Sets and Adversarial Rewards
Aadirupa Saha (Indian Institute of Science (IISc), Bangalore) · Pierre Gaillard () · Michal Valko (DeepMind)

From PAC to Instance-Optimal Sample Complexity in the Plackett-Luce Model
Aadirupa Saha (Indian Institute of Science (IISc), Bangalore) · Aditya Gopalan (Indian Institute of Science)

Reliable Fidelity and Diversity Metrics for Generative Models
Muhammad Ferjad Naeem (Technical University of Munich) · Seong Joon Oh (Clova AI Research, NAVER Corp.) · Yunjey Choi (Clova AI Research, NAVER Corp.) · Youngjung Uh (Clova AI Research, NAVER Corp.) · Jaejun Yoo (EPFL)

Learning Factorized Weight Matrix for Joint Image Filtering
Xiangyu Xu (Tsinghua University) · Yongrui Ma (SenseTime) · Wenxiu Sun (SenseTime Research)

Likelihood-free MCMC with Amortized Approximate Ratio Estimators
Joeri Hermans (University of Liège) · Volodimir Begy (CERN) · Gilles Louppe (University of Liège)

Attacks Which Do Not Kill Training Make Adversarial Learning Stronger
Jingfeng Zhang (National University of Singapore) · Xilie Xu (Shandong University) · Bo Han (HKBU / RIKEN) · Gang Niu (RIKEN) · Lizhen Cui (ShanDong University) · Masashi Sugiyama (RIKEN / The University of Tokyo) · Mohan Kankanhalli (National University of Singapore,)

GradientDICE: Rethinking Generalized Offline Estimation of Stationary Values
Shangtong Zhang (University of Oxford) · Bo Liu (Auburn University) · Shimon Whiteson (University of Oxford)

Provably Convergent Two-Timescale Off-Policy Actor-Critic with Function Approximation
Shangtong Zhang (University of Oxford) · Bo Liu (Auburn University) · Hengshuai Yao (Huawei Technologies) · Shimon Whiteson (University of Oxford)

Adversarial Attacks on Probabilistic Autoregressive Forecasting Models
Raphaël Dang-Nhu (ETH Zürich) · Gagandeep Singh (ETH Zurich) · Pavol Bielik (ETH Zurich) · Martin Vechev (ETH Zurich)

Informative Dropout for Robust Representation Learning: A Shape-bias Perspective
Baifeng Shi (Peking University) · Dinghuai Zhang (Peking University) · Qi Dai (Microsoft Research) · Jingdong Wang (Microsoft) · Zhanxing Zhu (Peking University) · Yadong Mu (Peking University)

Graph Convolutional Network for Recommendation with Low-pass Collaborative Filters
Wenhui Yu (Tsinghua University) · Zheng Qin (Tsinghua University)

SoftSort: A Differantiable Continuous Relaxation of the argsort Operator
Sebastian Prillo (UC Berkeley) · Julian Eisenschlos (Google)

Too Relaxed to Be Fair
Michael Lohaus (Universität Tübingen) · Michael Perrot (Université Jean Monnet) · Ulrike von Luxburg (University of Tübingen)

Lorentz Group Equivariant Neural Network for Particle Physics
Alexander Bogatskiy (University of Chicago) · Brandon Anderson (University of Chicago) · Jan Offermann (University of Chicago) · Marwah Roussi (University of Chicago) · David Miller (University of Chicago) · Risi Kondor (The University of Chicago)

One-shot distributed ridge regression in high dimensions
Yue Sheng (University of Pennsylvania) · Edgar Dobriban (University of Pennsylvania)

Streaming k-Submodular Maximization under Noise subject to Size Constraint
Lan Nguyen (University of Florida) · My Thai (University of Florida)

Variational Imitation Learning with Diverse-quality Demonstrations
Voot Tangkaratt (RIKEN AIP) · Bo Han (HKBU / RIKEN) · Mohammad Emtiyaz Khan (RIKEN) · Masashi Sugiyama (RIKEN / The University of Tokyo)

Task Understanding from Confusing Mulit-task Data
Xin Su (Tsinghua University) · yizhou Jiang (Tsinghua University) · Shangqi Guo (Tsinghua University) · Feng Chen (Tsinghua University)

Cost-effective Interactive Attention Learning with Neural Attention Process
Jay Heo (KAIST) · Junhyeon Park (KAIST) · Hyewon Jeong (KAIST) · Kwang Joon Kim (Yonsei University College of Medicine) · Juho Lee (AITRICS) · Eunho Yang (KAIST,AITRICS) · Sung Ju Hwang (KAIST, AITRICS)

Channel Equilibrium Networks for Learning Deep Representation
Wenqi Shao (The Chinese University of HongKong) · Shitao Tang (Simon Fraser University) · Xingang Pan (The Chinese University of Hong Kong) · Ping Tan (Simon Fraser University) · Xiaogang Wang (Chinese University of Hong Kong, Hong Kong) · Ping Luo (The University of Hong Kong)

Optimal Non-parametric Learning in Repeated Contextual Auctions with Strategic Buyer
Alexey Drutsa (Yandex)

Topological Autoencoders
Michael Moor (ETH Zurich) · Max Horn (MLCB, D-BSSE, ETH Zurich) · Bastian Rieck (ETH Zurich) · Karsten Borgwardt (ETH Zurich)

An Accelerated DFO Algorithm for Finite-sum Convex Functions
Yuwen Chen (ETH Zurich) · Antonio Orvieto (ETH Zurich) · Aurelien Lucchi (ETH Zurich)

The Shapley Taylor Interaction Index
Mukund Sundararajan (Google Inc.) · Kedar Dhamdhere (Google LLC) · Ashish Agarwal (Google Brain)

Privately detecting changes in unknown distributions
Rachel Cummings (Georgia Tech) · Sara Krehbiel (Santa Clara University) · Yuliia Lut (Georgia Institute of Technology) · Wanrong Zhang (Georgia Institute of Technology)

CAUSE: Learning Granger Causality from Event Sequences using Attribution Methods
Wei Zhang (University of Wisconsin-Madison) · Thomas Panum (Aalborg University) · Somesh Jha (University of Wisconsin, Madison) · Prasad Chalasani (MediaMath) · David Page (Duke)

Efficient Continuous Pareto Exploration in Multi-Task Learning
Pingchuan Ma (MIT) · Tao Du (MIT) · Wojciech Matusik (MIT)

WaveFlow: A Compact Flow-based Model for Raw Audio
Wei Ping (Baidu Research) · Kainan Peng (Baidu Research) · Kexin Zhao (Baidu) · Zhao Song (Baidu Research)

Multi-Agent Determinantal Q-Learning
Yaodong Yang (AIG) · Ying Wen (UCL) · Jun Wang (UCL) · Liheng Chen (Shanghai Jiao Tong University) · Kun Shao (Huawei Noah's Ark Lab) · David Mguni (Noah's Ark Laboratory, Huawei) · Weinan Zhang (Shanghai Jiao Tong University)

Revisiting Spatial Invariance with Low-Rank Local Connectivity
Gamaleldin Elsayed (Google Brain) · Prajit Ramachandran (Google) · Jon Shlens (Google Brain) · Simon Kornblith (Google Brain)

Minimax Weight and Q-Function Learning for Off-Policy Evaluation
Masatoshi Uehara (Harvard University) · Jiawei Huang (University of Illinois at Urbana-Champaign) · Nan Jiang (University of Illinois at Urbana-Champaign)

Tensor denoising and completion based on ordinal observations
Chanwoo Lee (University of Wisconsin - Madison) · Miaoyan Wang (University of Wisconsin - Madison)

Learning Human Objectives by Evaluating Hypothetical Behavior
Siddharth Reddy (University of California, Berkeley) · EECS Anca Dragan (EECS Department, University of California, Berkeley) · Sergey Levine (UC Berkeley) · Shane Legg (DeepMind) · Jan Leike (DeepMind)

Counterfactual Cross-Validation: Stable Model Selection Procedure for Causal Inference Models
Yuta Saito (Tokyo Institute of Technology.) · Shota Yasui (Cyberagent)

Learning Efficient Multi-agent Communication: An Information Bottleneck Approach
Rundong Wang (Nanyang Technological University) · Xu He (Nanyang Technological University) · Runsheng Yu (Nanyang Technological University) · Wei Qiu (Nanyang Technological University) · Bo An (Nanyang Technological University) · Zinovi Rabinovich (Nanyang Technological University)

MoNet3D: Towards Accurate Monocular 3D Object Localization in Real Time
XICHUAN ZHOU (Chongqing University) · YiChong Peng (Chongqing University) · Chunqiao Long (Chongqing University) · Fengbo Ren (Arizona State University) · Cong Shi (Chongqing University)

S2GA: Robust Deep Learning with Noisy Labels without Early Stopping
Bo Han (HKBU / RIKEN) · Gang Niu (RIKEN) · Xingrui Yu (University of Technology Sydney) · QUANMING YAO (4Paradigm) · Miao Xu (University of Queensland/ RIKEN AIP) · Ivor Tsang (University of Technology Sydney) · Masashi Sugiyama (RIKEN / The University of Tokyo)

Multinomial Logit Bandit with Low Switching Cost
Kefan Dong (Tsinghua University) · Yingkai Li (Northwestern University) · Qin Zhang (Indiana University Bloomington) · Yuan Zhou (UIUC)

Deep Reasoning Networks for Unsupervised Pattern De-mixing with Constraint Reasoning
Di Chen (Cornell University) · Yiwei Bai (Cornell University) · Wenting Zhao (Cornell University) · Sebastian Ament (Cornell University) · John Gregoire (Caltech) · Carla Gomes (Cornell University)

Uncertainty-Aware Lookahead Factor Models for Improved Quantitative Investing
Lakshay Chauhan (University of Michigan) · John Alberg (Euclidean Technologies) · Zachary Lipton (Carnegie Mellon University)

On the Unreasonable Effectiveness of the Greedy Algorithm: Greedy Adapts to Sharpness
Sebastian Pokutta (ZIB) · Mohit Singh (Georgia Institute of Technology) · Alfredo Torrico (Polytechnique Montreal)

Stronger and Faster Wasserstein Adversarial Attacks
Kaiwen Wu (University of Waterloo) · Allen Wang (University of Waterloo) · Yaoliang Yu (University of Waterloo)

Optimizing Multiagent Cooperation via Policy Evolution and Shared Experiences
Somdeb Majumdar (Intel AI Lab) · Shauharda Khadka (Intel AI) · Santiago Miret (Intel AI Products Group) · Stephen Mcaleer (UC Irvine) · Kagan Tumer (Oregon State University US)

Why are learned indexes so effective?
Paolo Ferragina (Università di Pisa) · Fabrizio Lillo (Università di Bologna) · Giorgio Vinciguerra (University of Pisa)

Fast OSCAR and OWL with Safe Screening Rules
Runxue Bao (University of Pittsburgh) · Bin Gu (Nanjing University of Information Science & Technology) · Heng Huang (University of Pittsburgh)

Which Tasks Should Be Learned Together in Multi-task Learning?
Trevor Standley (Stanford University) · Amir Zamir (Stanford, UC Berkeley) · Dawn Chen (Google) · Leonidas Guibas (Stanford University) · Jitendra Malik (University of California at Berkeley) · Silvio Savarese (Stanford University)

Inertial Block Proximal Methods for Non-Convex Non-Smooth Optimization
Hien Le (University of Mons, Belgium.) · Nicolas Gillis (Université de Mons) · Panagiotis Patrinos (KU Leuven)

Adversarial Neural Pruning with Latent Vulnerability Suppression
Divyam Madaan (KAIST) · Jinwoo Shin (KAIST) · Sung Ju Hwang (KAIST, AITRICS)

Lifted Disjoint Paths with Application in Multiple Object Tracking
Andrea Hornakova (Max Planck Institute for Informatics) · Roberto Henschel (Leibniz University of Hannover) · Bodo Rosenhahn (Leibniz University Hannover) · Paul Swoboda (MPI fuer Informatik, Saarbruecken)

Being Bayesian, Even Just a Bit, Fixes Overconfidence in ReLU Networks
Agustinus Kristiadi (University of Tuebingen) · Matthias Hein (University of Tübingen) · Philipp Hennig (University of Tuebingen)

SCAFFOLD: Stochastic Controlled Averaging for Federated Learning
Sai Praneeth Reddy Karimireddy (EPFL) · Satyen Kale (Google) · Mehryar Mohri (Google Research and Courant Institute of Mathematical Sciences) · Sashank Jakkam Reddi (Google) · Sebastian Stich (EPFL) · Ananda Theertha Suresh (Google Research)

Statistically Preconditioned Accelerated Gradient Method for Distributed Optimization
Hadrien Hendrikx (INRIA) · Lin Xiao (Microsoft Research) · Sebastien Bubeck (Microsoft Research) · Francis Bach (INRIA - Ecole Normale Supérieure) · Laurent Massoulié (MSR-INRIA Joint Center)

Pretrained Generalized Autoregressive Model with Adaptive Probabilistic Label Cluster for Extreme Multi-label Text Classification
Hui Ye (Lehigh University) · Zhiyu Chen (Lehigh University) · Da-Han Wang (Xiamen University of Technology) · Brian Davison (Lehigh University)

Frequentist Uncertainty in Recurrent Neural Networks via Blockwise Influence Functions
Ahmed Alaa (UCLA) · M van der Schaar (UCLA)

Disentangling Trainability and Generalization in Deep Neural Networks
Lechao Xiao (Google Brain) · Jeffrey Pennington (Google Brain) · Samuel Schoenholz (Google Brain)

Moniqua: Modulo Quantized Communication in Decentralized SGD
Yucheng Lu (Cornell University) · Christopher De Sa (Cornell)

Expectation Maximization with Bias-Corrected Calibration is Hard-To-Beat at Label Shift Adaptation
Amr Mohamed Alexandari (Stanford University) · Anshul Kundaje (Stanford University) · Avanti Shrikumar (Stanford University)

Expert Learning through Generalized Inverse Multiobjective Optimization: Models, Insights and Algorithms
Chaosheng Dong (Amazon) · Bo Zeng (University of Pittsburgh)

Random Matrix Theory Proves that Deep Learning Representations of GAN-data Behave as Gaussian Mixtures
Mohamed El Amine Seddik (CEA) · Cosme Louart (CEA) · Mohamed Tamaazousti (CEA Saclay) · Romain COUILLET (CentraleSupélec)

Optimizing Data Usage via Differentiable Rewards
Xinyi Wang (Carnegie Mellon University) · Hieu Pham (Carnegie Mellon University) · Paul Michel (Carnegie Mellon University) · Antonios Anastasopoulos (Carnegie Mellon University) · Jaime Carbonell (Carnegie Mellon University) · Graham Neubig (Carnegie Mellon University)

Optimistic Policy Optimization with Bandit Feedback
Lior Shani (Technion) · Yonathan Efroni (Technion) · Aviv Rosenberg (Tel Aviv University) · Shie Mannor (Technion)

Maximum-and-Concatenation Networks
Xingyu Xie (Peking Unversity) · Hao Kong (Peking University) · Jianlong Wu (Peking University) · Wayne Zhang (SenseTime Research) · Guangcan Liu (Nanjing University of Information Science and Technology) · Zhouchen Lin (Peking University)

Learning Adversarial Markov Decision Processes with Bandit Feedback and Unknown Transition
Chi Jin (Princeton University) · Tiancheng Jin (University of Southern California) · Haipeng Luo (University of Southern California) · Suvrit Sra (MIT) · Tiancheng Yu (MIT)

Kernelized Stein Discrepancy Tests of Goodness-of-fit for Time-to-Event Data
Wenkai Xu (Gatsby Unit, UCL) · Tamara Fernandez (University College London) · Nicolas Rivera (University of Cambridge) · Arthur Gretton (Gatsby Computational Neuroscience Unit)

Efficient Intervention Design for Causal Discovery with Latents
Raghavendra Addanki (University of Massachusetts Amherst) · Shiva Kasiviswanathan (Amazon) · Andrew McGregor (University of Massachusetts Amherst) · Cameron Musco (UMass)

Certified Data Removal from Machine Learning Models
Chuan Guo (Cornell University) · Tom Goldstein (University of Maryland) · Awni Hannun (Facebook AI Research) · Laurens van der Maaten (Facebook)

One Size Fits All: Can We Train One Denoiser for All Noise Levels?
Abhiram Gnanasambandam (Purdue University) · Stanley Chan (Purdue University, USA)

GNN-FiLM: Graph Neural Networks with Feature-wise Linear Modulation
Marc Brockschmidt (Microsoft Research)

Sparse Gaussian Processes with Spherical Harmonic Features
Vincent Dutordoir (PROWLER.io) · Nicolas Durrande (PROWLER.io) · James Hensman (PROWLER.io)

Asynchronous Coagent Networks
James Kostas (University of Massachusetts Amherst) · Chris Nota (University of Massachusetts Amherst) · Philip Thomas (University of Massachusetts Amherst)

Adaptive Checkpoint Adjoint Method for Gradient Estimation in Neural ODE
Juntang Zhuang (Yale University) · Nicha Dvornek (Yale University) · Xiaoxiao Li (Yale University) · Sekhar Tatikonda (Yale) · Xenophon Papademetris (Yale University) · James Duncan (Yale University)

Understanding the Curse of Horizon in Off-Policy Evaluation via Conditional Importance Sampling
Yao Liu (Stanford University) · Pierre-Luc Bacon (Stanford University) · Emma Brunskill (Stanford University)

Taylor Expansion Policy Optimization
Yunhao Tang (Columbia University) · Michal Valko (DeepMind) · Remi Munos (DeepMind)

Reinforcement Learning for Integer Programming: Learning to Cut
Yunhao Tang (Columbia University) · Shipra Agrawal (Columbia University) · Yuri Faenza (Columbia University)

Safe Reinforcement Learning in Constrained Markov Decision Processes
Akifumi Wachi (IBM Research AI) · Yanan Sui (Tsinghua University)

Layered Sampling for Robust Optimization Problems
Hu Ding (University of Science and Technology of China) · Zixiu Wang (University of Science and Technology of China)

Learning to Encode Position for Transformer with Continuous Dynamical Model
Xuanqing Liu (University of California Los Angeles) · Hsiang-Fu Yu (Amazon) · Inderjit Dhillon (UT Austin & Amazon) · Cho-Jui Hsieh (UCLA)

Do RNN and LSTM have Long Memory?
Jingyu Zhao (The University of Hong Kong) · Feiqing Huang (University of Hong Kong) · Jia Lv (Huawei Noah's Ark Lab) · Yanjie Duan (Huawei Noah's Ark Lab) · Zhen Qin (Huawei Noah's Ark Lab) · Guodong Li (University of Hong Kong) · Guangjian Tian (Huawei Noah's Ark Lab)

Training Linear Neural Networks: Non-Local Convergence and Complexity Results
Armin Eftekhari (Umea University)

On Validation and Planning of An Optimal Decision Rule with Application in Healthcare Studies
Hengrui Cai (North Carolina State University) · Wenbin Lu () · Rui Song ()

Graph Optimal Transport for Cross-Domain Alignment
Liqun Chen (Duke University) · Zhe Gan (Microsoft) · Yu Cheng (Microsoft) · Linjie Li (Microsoft) · Lawrence Carin (Duke) · Jingjing Liu (Microsoft)

Approximation Capabilities of Neural ODEs and Invertible Residual Networks
Han Zhang (Virginia Commonwealth University) · Xi Gao (Virginia Commonwealth University) · Jacob Unterman (Virginia Commonwealth University) · Tomasz Arodz (Virginia Commonwealth University)

Refined bounds for algorithm configuration: The knife-edge of dual class approximability
Nina Balcan (Carnegie Mellon University) · Tuomas Sandholm (Carnegie Mellon University) · Ellen Vitercik (Carnegie Mellon University)

Teaching with Limited Information on the Learner's Behaviour
Ferdinando Cicalese (University of Verona) · Sergio Filho (PUC-RIO) · Eduardo Laber (PUC-RIO) · Marco Molinaro (PUC-RIO)

Interpretations are Useful: Penalizing Explanations to Align Neural Networks with Prior Knowledge
Laura Rieger (Technical University of Denmark) · Chandan Singh (UC Berkeley) · William Murdoch (UC Berkeley) · Bin Yu (University of California, Berkeley)

DeltaGrad: Rapid retraining of machine learning models
Yinjun Wu (university of pennsylvania) · Edgar Dobriban (University of Pennsylvania) · Susan Davidson (University of Pennsylvania)

The Cost-free Nature of Optimally Tuning Tikhonov Regularizers and Other Ordered Smoothers
Pierre Bellec (rutgers) · Dana Yang (Duke University)

Approximation Guarantees of Local Search Algorithms via Localizability of Set Functions
Kaito Fujii (National Institute of Informatics)

Fine-Grained Analysis of Stability and Generalization for Stochastic Gradient Descent
Yunwen Lei (University of Kaiserslautern) · Yiming Ying (SUNY Albany)

Online Dense Subgraph Discovery via Blurred-Graph Feedback
Yuko Kuroki (The University of Tokyo /RIKEN) · Atsushi Miyauchi (University of Tokyo) · Junya Honda (University of Tokyo / RIKEN) · Masashi Sugiyama (RIKEN / The University of Tokyo)

Lazylter: A Fast Algorithm for Counting Markov Equivalent DAGs and Designing Experiments
Ali AhmadiTeshnizi (Sharif University of Technology) · Saber Salehkaleybar (Sharif University of Technology) · Negar Kiyavash (École Polytechnique Fédérale de Lausanne)

Perceptual Generative Autoencoders
Zijun Zhang (University of Calgary) · Ruixiang ZHANG (Mila/UdeM) · Zongpeng Li (Wuhan University) · Yoshua Bengio (Mila / U. Montreal) · Liam Paull (Université de Montréal)

Towards Understanding the Regularization of Adversarial Robustness on Neural Networks
Yuxin Wen (South China University of Technology) · Shuai Li (South China University of Technology) · Kui Jia (South China University of Technology)

Stochastic Gradient and Langevin Processes
Xiang Cheng (UC Berkeley) · Dong Yin (UC Berkeley) · Peter Bartlett (Berkeley) · Michael Jordan (UC Berkeley)

ROMA: Multi-Agent Reinforcement Learning with Emergent Roles
Tonghan Wang (Tsinghua University) · Heng Dong (Tsinghua) · Victor Lesser (UMASS) · Chongjie Zhang (Tsinghua University)

Minimax Pareto Fairness: A Multi Objective Perspective
Martin Bertran (Duke University) · Natalia Martinez (Duke University) · Guillermo Sapiro (Duke University)

Online Pricing with Offline Data: Phase Transition and Inverse Square Law
Jinzhi Bu (MIT) · David Simchi-Levi (MIT) · Yunzong Xu (MIT)

Explicit Gradient Learning for Black-Box Optimization
Elad Sarafian (Bar-Ilan University) · Mor Sinay (Bar-Ilan University) · yoram louzoun (Bar Ilan University) · Noa Agmon (Bar-Ilan University) · Sarit Kraus (Bar-Ilan University)

Optimization and Analysis of the pAp@k Metric for Recommender Systems
Gaurush Hiranandani (University of Illinois at Urbana-Champaign) · Warut Vijitbenjaronk (University of Illinois, Urbana-Champaign) · Sanmi Koyejo (Illinois / Google) · Prateek Jain (Microsoft Research)

When Explanations Lie: Why Many Modified BP Attributions Fail
Leon Sixt (Frei Universität Berlin) · Maximilian Granz (Freie Universität Berlin) · Tim Landgraf (Freie Universität Berlin)

Naive Exploration is Optimal for Online LQR
Max Simchowitz (UC Berkeley) · Dylan Foster (MIT)

Learning Structured Latent Factors from Dependent Data:A Generative Model Framework from Information-Theoretic Perspective
Ruixiang ZHANG (Mila/UdeM) · Katsuhiko Ishiguro (NTT Docomo) · Masanori Koyama (Preferred Networks Inc.)

Implicit Generative Modeling for Efficient Exploration
Neale Ratzlaff (Oregon State University) · Qinxun Bai (Horizon Robotics) · Fuxin Li (Oregon State University) · Wei Xu (Horizon Robotics)

Prediction-Guided Multi-Objective Reinforcement Learning for Continuous Robot Control
Jie Xu (Massachusetts Institute of Technology) · Yunsheng Tian (Massachusetts Institute of Technology) · Pingchuan Ma (MIT) · Daniela Rus (MIT CSAIL) · Shinjiro Sueda (Texas A&M University) · Wojciech Matusik (MIT)

Goodness-of-Fit Tests for Inhomogeneous Random Graphs
Soham Dan (University of Pennsylvania) · Bhaswar B. Bhattacharya (University of Pennsylvania)

Few-shot Domain Adaptation by Causal Mechanism Transfer
Takeshi Teshima (The University of Tokyo / RIKEN) · Issei Sato (University of Tokyo / RIKEN) · Masashi Sugiyama (RIKEN / The University of Tokyo)

Adaptive Adversarial Multi-task Representation Learning
YUREN MAO (School of Computer Science and Engineering, University of New South Wales) · Weiwei Liu (Wuhan University) · Xuemin Lin (University of New South Wales)

Streaming Submodular Maximization under a k-Set System Constraint
Ran Haba (Open University of Israel) · Ehsan Kazemi (Yale) · Moran Feldman (University of Haifa) · Amin Karbasi (Yale)

A Generic First-Order Algorithmic Framework for Bi-Level Programming Beyond Lower-Level Singleton
Risheng Liu (Dalian University of Technology) · Pan Mu (Dalian University of Technology) · Xiaoming Yuan (The University of Hong Kong) · Shangzhi Zeng (The University of Hong Kong) · Jin Zhang (Southern University of Science and Technology)

Optimal approximation for unconstrained non-submodular minimization
Marwa El Halabi (MIT) · Stefanie Jegelka (Massachusetts Institute of Technology)

Generating Programmatic Referring Expressions via Program Synthesis
Jiani Huang (University of Pennsylvania) · Calvin Smith (University of Wisconsin at Madison) · Osbert Bastani (University of Pennsylvania) · Rishabh Singh (Google Brain) · Aws Albarghouthi (University of Wisconsin-Madison) · Mayur Naik (University of Pennsylvania)

Nearly Linear Row Sampling Algorithm for Quantile Regression
Yi Li (Nanyang Technological University) · Ruosong Wang (Carnegie Mellon University) · Lin Yang (UCLA) · Hanrui Zhang (Duke University)

On Leveraging Pretrained GANs for Limited-Data Generation
Miaoyun Zhao (Duke University) · Yulai Cong (Duke University) · Lawrence Carin (Duke)

More Data Can Expand The Generalization Gap Between Adversarially Robust and Standard Models
Lin Chen (Yale University) · Yifei Min (Yale University) · Mingrui Zhang (Yale University) · Amin Karbasi (Yale)

Double Reinforcement Learning for Efficient and Robust Off-Policy Evaluation
Nathan Kallus (Cornell University) · Masatoshi Uehara (Harvard University)

Statistically Efficient Off-Policy Policy Gradients

Nathan Kallus (Cornell University) · Masatoshi Uehara (Harvard University)

Self-PU: Self Boosted and Calibrated Positive-Unlabeled Training
Xuxi Chen (University of Science and Technology of China) · Wuyang Chen (Texas A&M University) · Tianlong Chen (Texas A&M University) · Ye Yuan (Texas A&M University) · Chen Gong (Nanjing University of Science and Technology) · Kewei Chen (Green Valley Pharmaceutical LLC) · Zhangyang Wang (Texas A&M University)

When Does Self-Supervision Help Graph Convolutional Networks?
Yuning You (Texas A&M University) · Tianlong Chen (Texas A&M University) · Zhangyang Wang (Texas A&M University) · Yang Shen (Texas A&M University)

On Differentially Private Stochastic Convex Optimization with Heavy-tailed Data
Di Wang (State University of New York at Buffalo) · Hanshen Xiao (MIT CSAIL) · Srinivas Devadas (MIT) · Jinhui Xu (SUNY Buffalo)

Variance Reduced Coordinate Descent with Acceleration: New Method With a Surprising Application to Finite-Sum Problems
Filip Hanzely (KAUST) · Dmitry Kovalev (KAUST) · Peter Richtarik (KAUST)

Stochastic Subspace Cubic Newton Method
Filip Hanzely (KAUST) · Nikita Doikov (Université catholique de Louvain) · Yurii Nesterov (Universite catholique de Louvain) · Peter Richtarik (KAUST)

Ready Policy One: World Building Through Active Learning
Philip Ball (University of Oxford) · Jack Parker-Holder (University of Oxford) · Aldo Pacchiano (UC Berkeley) · Krzysztof Choromanski (Google) · Stephen Roberts (University of Oxford)

Structural Language Models of Code
Uri Alon (Technion) · Roy Sadaka (Technion) · Omer Levy (University of Washington) · Eran Yahav (Technion)

PEGASUS: Pre-training with Extracted Gap-sentences for Abstractive Summarization
Jingqing Zhang (Imperial College London) · Yao Zhao (Google) · Mohammad Saleh (Google) · Peter Liu (Google Brain)

Aggregation of Multiple Knockoffs
Tuan-Binh Nguyen (INRIA Saclay Ile-de-France) · Jerome-Alexis Chevalier (INRIA Saclay Ile-de-France) · Sylvain Arlot (University Paris Sud) · Thirion Bertrand (inria)

Off-Policy Actor-Critic with Shared Experience Replay
Simon Schmitt (DeepMind) · Matteo Hessel (Deep Mind) · Karen Simonyan (DeepMind)

Graph-based Nearest Neighbor Search: From Practice to Theory
Liudmila Prokhorenkova (Yandex) · Aleksandr Shekhovtsov (Yandex)

Policy Teaching via Environment Poisoning: Training-time Adversarial Attacks against Reinforcement Learning
Amin Rakhsha (MPI-SWS) · Goran Radanovic (Max Planck Institute for Software Systems) · Rati Devidze (Max Planck Institute for Software Systems) · Jerry Zhu (University of Wisconsin-Madison) · Adish Singla (Max Planck Institute (MPI-SWS))

Semismooth Newton Algorithm for Efficient Projections onto $\ell_{1,\infty}$ -norm Ball
Dejun Chu (Tsinghua University) · Changshui Zhang (Tsinghua University) · Shiliang Sun (East China Normal University) · Qing Tao (Army Academy of Artillery and Air Defense)

Influenza forecasting framework based on Gaussian processes
Christoph Zimmer (Bosch Center for Artificial Intelligence BCAI) · Reza Yaesoubi (Health Policy and Management, Yale School of Public Health)

Unique Properties of Wide Minima in Deep Networks
Rotem Mulayoff (Technion) · Tomer Michaeli (Technion)

Does the Markov Decision Process Fit the Data: Testing for the Markov Property in Sequential Decision Making
Chengchun Shi (London School of Economics and Political Science) · Runzhe Wan (North Carolina State University) · Rui Song () · Wenbin Lu () · Ling Leng (Amazon)

LTF: A Label Transformation Framework for Correcting Label Shift
Jiaxian Guo (The University of Sydney) · Mingming Gong (University of Melbourne) · Tongliang Liu (The University of Sydney) · Kun Zhang (Carnegie Mellon University) · Dacheng Tao (The University of Sydney)

Divide, Conquer, and Combine: a New Inference Strategy for Probabilistic Programs with Stochastic Support
Yuan Zhou (University of Oxford) · Hongseok Yang (KAIST) · Yee Whye Teh (Oxford and DeepMind) · Tom Rainforth (University of Oxford)

Duality in RKHSs with Infinite Dimensional Outputs: Application to Robust Losses
Pierre Laforgue (Télécom ParisTech) · Alex Lambert (Télécom ParisTech) · Luc Brogat-Motte (Télécom Paris) · Florence d'Alche-Buc (Télécom ParisTech, Université Paris-Saclay,Paris, France)

Causal Effect Estimation and Optimal Dose Suggestions in Mobile Health
Liangyu Zhu (North Carolina State University) · Wenbin Lu () · Rui Song ()

Towards Understanding the Dynamics of the First-Order Adversaries
Zhun Deng (Harvard) · Hangfeng He (University of Pennsylvania) · Jiaoyang Huang (Institute of Advanced Study) · Weijie Su (University of Pennsylvania)

Interpreting Robust Optimization via Adversarial Influence Functions
Zhun Deng (Harvard) · Cynthia Dwork (Harvard) · Jialiang Wang (Harvard University) · Linjun Zhang (Rutgers University)

Multilinear Latent Conditioning for Generating Unseen Attribute Combinations
Markos Georgopoulos (Imperial College London) · Grigorios Chrysos (Imperial College London) · Yannis Panagakis (Imperial College London) · Maja Pantic (Samsung AI Centre Cambridge/ Imperial College London)

No-Regret Exploration in Goal-Oriented Reinforcement Learning
Jean Tarbouriech (Facebook AI Research Paris & Inria Lille) · Evrard Garcelon (Facebook AI Research) · Michal Valko (DeepMind) · Matteo Pirotta (Facebook AI Research) · Alessandro Lazaric (Facebook AI Research)

OPtions as REsponses: Grounding behavioural hierarchies in multi-agent reinforcement learning
Alexander Vezhnevets (DeepMind) · Yuhuai Wu (University of Toronto) · Maria Eckstein (UC Berkeley) · Rémi Leblond (DeepMind) · Joel Z Leibo (DeepMind)

Feature Noise Induces Loss Discrepancy Across Groups
Fereshte Khani (Stanford University) · Percy Liang (Stanford University)

Reinforcement Learning for Molecular Design Guided by Quantum Mechanics
Gregor Simm (Cambridge University) · Robert Pinsler (University of Cambridge) · Jose Hernandez-Lobato (University of Cambridge)

Small-GAN: Speeding up GAN Training using Core-Sets
Samrath Sinha (University of Toronto) · Han Zhang (Google) · Anirudh Goyal (Université de Montréal) · Yoshua Bengio (Mila / U. Montreal) · Hugo Larochelle (Google Brain) · Augustus Odena (Google Brain)

Conditional gradient methods for stochastically constrained convex minimization
Maria-Luiza Vladarean (EPFL) · Ahmet Alacaoglu (EPFL) · Ya-Ping Hsieh (EPFL) · Volkan Cevher (EPFL)

Undirected Graphical Models as Approximate Posteriors
Arash Vahdat (NVIDIA) · Evgeny Andriyash (D-Wave Systems Inc.) · William Macready (D-Wave)

Dynamics of Deep Neural Networks and Neural Tangent Hierarchy
Jiaoyang Huang (IAS) · Horng-Tzer Yau (Harvard)

Measuring Non-Expert Comprehension of Machine Learning Fairness Metrics
Debjani Saha (University of Maryland) · Candice Schumann (University of Maryland) · Duncan McElfresh (University of Maryland) · John P Dickerson (University of Maryland) · Michelle Mazurek (University of Maryland) · Michael Tschantz (International Computer Science Institute)

Encoding Musical Style with Transformer Autoencoders
Kristy Choi (Stanford University) · Curtis Hawthorne (Google Brain) · Ian Simon (Google Brain) · Monica Dinculescu (Google Brain) · Jesse Engel (Google Brain)

Min-Max Optimization without Gradients: Convergence and Applications to Black-Box Evasion and Poisoning Attacks
Sijia Liu (MIT-IBM Watson AI Lab) · Songtao Lu (IBM Research) · XIANGYI CHEN (University of Minnesota) · Yao Feng (Tsinghua University) · Kaidi Xu (Northeastern University) · Abdullah Al-Dujaili (CSAIL) · Mingyi Hong (University of Minnesota) · Una-May O'Reilly (MIT)

ConQUR: Mitigating Delusional Bias in Deep Q-Learning
DiJia Su (Princeton University) · Jayden Ooi (Google) · Tyler Lu (Google) · Dale Schuurmans (Google / University of Alberta) · Craig Boutilier (Google)

Self-Modulating Nonparametric Event-Tensor Factorization
Zheng Wang (University of Utah) · Xinqi Chu (Xjera Labs, Pte.Ltd) · Shandian Zhe (University of Utah)

Extreme Multi-label Classification from Aggregated Labels
Yanyao Shen (UT Austin) · Hsiang-Fu Yu (Amazon) · Sujay Sanghavi (UT Austin) · Inderjit Dhillon (UT Austin & Amazon)

Full Law Identification In Graphical Models Of Missing Data: Completeness Results
Razieh Nabi (Johns Hopkins University) · Rohit Bhattacharya (Johns Hopkins University) · Ilya Shpitser (Johns Hopkins University)

Self-Attentive Associative Memory
Hung Le (Deakin University) · Truyen Tran (Deakin University) · Svetha Venkatesh (Deakin University)

Imputer: Sequence Modelling via Imputation and Dynamic Programming
William Chan (Google Brain) · Chitwan Saharia (Google) · Geoffrey Hinton (Google) · Mohammad Norouzi (Google Brain) · Navdeep Jaitly (D. E. Shaw)

Continuously Indexed Domain Adaptation
Hao Wang (MIT) · Hao He (Massachusetts Institute of Technology) · Dina Katabi (MIT)

Evolving Machine Learning Algorithms From Scratch
Esteban Real (Google Inc.) · Chen Liang (Google Brain) · David So (Google Brain) · Quoc Le (Google Brain)

Self-Attentive Hawkes Process
Qiang Zhang (University College London) · Aldo Lipani (University College London) · Omer Kirnap (University College London) · Emine Yilmaz (University College London)

On hyperparameter tuning in general clustering problemsm
Xinjie Fan (UT Austin) · Yuguang Yue (University of Texas at Austin) · Purnamrita Sarkar (UT Austin) · Y. X. Rachel Wang (University of Sydney)

Communication-Efficient Distributed Stochastic AUC Maximization with Deep Neural Networks
Zhishuai Guo (The University of Iowa) · Mingrui Liu (The University of Iowa) · Zhuoning Yuan (The University of Iowa) · Li Shen (Tencent AI Lab) · Wei Liu (Tencent AI Lab) · Tianbao Yang (The University of Iowa)

Adaptive Region-Based Active Learning
Corinna Cortes (Google Research) · Giulia DeSalvo (Google Research) · Claudio Gentile (INRIA and Google) · Mehryar Mohri (Google Research and Courant Institute of Mathematical Sciences) · Ningshan Zhang (New York University)

Robust Outlier Arm Identification
Yinglun Zhu (University of Wisconsin-Madison) · Sumeet Katariya (UW-Madison and Amazon) · Robert Nowak (University of Wisconsion-Madison)

Provably Efficient Exploration in Policy Optimization
Qi Cai (Northwestern University) · Zhuoran Yang (Princeton University) · Chi Jin (Princeton University) · Zhaoran Wang (Northwestern U)

Striving for simplicity and performance in off-policy DRL: Output Normalization and Non-Uniform Sampling
Che Wang (New York University) · Yanqiu Wu (New York University) · Quan Vuong (University of California San Diego) · Keith Ross (New York University Shanghai)

Multidimensional Shape Constraints
Maya Gupta (Google) · Erez Louidor (Google, Inc.) · Oleksandr Mangylov (Google Research) · Nobu Morioka (Google Research) · Tamann Narayan (Google) · Sen Zhao (Google Research)

Fast Deterministic CUR Matrix Decomposition with Accuracy Assurance
Yasutoshi Ida (NTT) · Sekitoshi Kanai (NTT Software Innovation Center) · Yasuhiro Fujiwara (NTT Communication Science Laboratories) · Tomoharu Iwata (NTT) · Koh Takeuchi (NTT) · Hisashi Kashima (Kyoto University/RIKEN Center for AIP)

Operation-Aware Soft Channel Pruning using Differentiable Masks
Minsoo Kang (Seoul National University) · Bohyung Han (Seoul National University)

Normalized Loss Functions for Deep Learning with Noisy Labels
Xingjun Ma (The University of Melbourne) · Hanxun Huang (University of Melbourne) · Yisen Wang (Tsinghua University) · Simone Romano (University of Melbourne) · Sarah Erfani (University of Melbourne) · James Bailey (The University of Melbourne)

Learning Deep Kernels for Non-Parametric Two-Sample Tests
Feng Liu (UTS/UCL) · Wenkai Xu (Gatsby Unit, UCL) · Jie Lu (University of Technology Sydney) · Guangquan Zhang (University of Technology Sydney) · Arthur Gretton (Gatsby Computational Neuroscience Unit) · D.J. Sutherland (TTI-Chicago)

DeBayes: a Bayesian method for debiasing network embeddings
Maarten Buyl (Ghent University) · Tijl De Bie (Ghent University)

Principled learning method for Wasserstein distributionally robust optimization with local perturbations
Yongchan Kwon (Stanford University) · Wonyoung Kim (Seoul National University) · Joong-Ho Won (Seoul National University) · Myunghee Cho Paik (Seoul National University)

Low-Variance and Zero-Variance Baselines for Extensive-Form Games
Trevor Davis (University of Alberta) · Martin Schmid (DeepMind) · Michael Bowling (DeepMind)

Converging to Team-Maxmin Equilibria in Zero-Sum Multiplayer Games
Youzhi Zhang (Nanyang Technological University) · Bo An (Nanyang Technological University)

Landscape Connectivity and Dropout Stability of SGD Solutions for Over-parameterized Neural Networks
Alexander Shevchenko (IST Austria) · Marco Mondelli (IST Austria)

Leveraging Frequency Analysis for Deep Fake Image Recognition
Joel Frank (Ruhr-University Bochum) · Thorsten Eisenhofer (Ruhr University Bochum) · Lea Schönherr (Ruhr-Universität Bochum) · Dorothea Kolossa (Ruhr University Bochum) · Thorsten Holz (Ruhr-Universität Bochum) · Asja Fischer (Ruhr University Bochum)

Tails of Lipschitz Triangular Flows
Priyank Jaini (University of Waterloo, Vector Institute) · Ivan Kobyzev (Borealis AI) · Yaoliang Yu (University of Waterloo) · Marcus Brubaker (Borealis AI)

Deep Coordination Graphs
Wendelin Boehmer (University of Oxford) · Vitaly Kurin (University of Oxford) · Shimon Whiteson (University of Oxford)

Voice Separation with an Unknown Number of Multiple Speakers
Eliya Nachmani (Tel-Aviv University & Facebook AI Research) · Yossi Adi (Bar-Ilan University) · Lior Wolf (Facebook AI Research and Tel Aviv University)

Predicting Choice with Set-Dependent Aggregation
Nir Rosenfeld (Harvard University) · Kojin Oshiba (Harvard University) · Yaron Singer (Harvard)

Thompson Sampling Algorithms for Mean-Variance Bandits
Qiuyu Zhu (National University of Singapore) · Vincent Tan (National University of Singapore)

Differentiable Likelihoods for Fast Inversion of 'Likelihood-Free' Dynamical Systems
Hans Kersting (University of Tuebingen) · Nicholas Krämer (University of Tübingen) · Martin Schiegg (Bosch Center for Artificial Intelligence) · Christian Daniel (Bosch Center for Artificial Intelligence) · Michael Schober (Bosch Center for Artificial Intelligence) · Philipp Hennig (University of Tuebingen)

Debiased Sinkhorn barycenters
Hicham Janati (INRIA) · Marco Cuturi (Google and CREST/ENSAE) · Alexandre Gramfort (Inria)

Double Trouble in Double Descent: Bias and Variance(s) in the Lazy Regime
Stéphane d'Ascoli (ENS) · Maria Refinetti (Laboratoire de Physique de l'Ecole Normale Supérieure Paris) · Giulio Biroli (ENS) · Florent Krzakala (ENS)

Explore, Discover and Learn: Unsupervised Discovery of State-Covering Skills
Victor Campos (Barcelona Supercomputing Center) · Alexander Trott (Salesforce Research) · Caiming Xiong (Salesforce) · Richard Socher (Salesforce) · Xavier Giro-i-Nieto (Universitat Politecnica de Catalunya) · Jordi Torres (Barcelona Supercomputing Center)

Sparsified Linear Programming for Zero-Sum Equilibrium Finding
Brian Zhang (Carnegie Mellon University) · Tuomas Sandholm (Carnegie Mellon University)

Extra-gradient with player sampling for faster convergence in n-player games
Samy Jelassi (Princeton University) · Carles Domingo-Enrich (NYU) · Damien Scieur (INRIA - ENS) · Arthur Mensch (ENS) · Joan Bruna (New York University)

Entropy Minimization In Emergent Languages
Evgeny Kharitonov (FAIR) · Rahma Chaabouni (Facebook/ENS/INRIA) · Diane Bouchacourt (Facebook AI) · Marco Baroni (Facebook Artificial Intelligence Research)

Spectral Clustering with Graph Neural Networks for Graph Pooling
Filippo Maria Bianchi (NORCE) · Daniele Grattarola (Università della Svizzera Italiana) · Cesare Alippi (Università della Svizzera Italiana)

VFlow: More Expressive Generative Flows with Variational Data Augmentation
Jianfei Chen (University of California, Berkeley) · Cheng Lu (Tsinghua University) · Biqi Chenli (Tsinghua University) · Jun Zhu (Tsinghua University) · Tian Tian (RealAI)

Fully Parallel Hyperparameter Search: Reshaped Space-Filling
Marie-Liesse Cauwet (Université Paris-Est, LIGM (UMR 8049), CNRS, ESIEE Paris) · Camille Couprie (FAIR) · Julien Dehos (LISIC, Université du Littoral Côte d'Opale) · Pauline Luc (Facebook AI Research) · Jeremy Rapin (Facebook AI Research) · Morgane Riviere (Facebook Artificial Intelligence Research) · Fabien Teytaud (LISIC, Université du Littoral Côte d'Opale) · Olivier Teytaud (Facebook) · Nicolas Usunier (Facebook AI Research)

Discount Factor as a Regularizer in Reinforcement Learning
Ron Amit (Technion – Israel Institute of Technology) · Kamil Ciosek (Microsoft) · Ron Meir (Technion Israeli Institute of Technology)

On Learning Sets of Symmetric Elements
Haggai Maron (Weizmann Institute of Science) · Or Litany (Stanford University) · Gal Chechik (NVIDIA / Bar-Ilan University) · Ethan Fetaya (Bar Ilan University)

Non-convex Learning via Replica Exchange Stochastic Gradient MCMC
Wei Deng (Purdue University) · Qi Feng (University of Southern California) · Liyao Gao (Purdue University) · Faming Liang (Purdue University) · Guang Lin (Purdue University)

Learning Similarity Metrics for Numerical Simulations
Georg Kohl (Technical University of Munich) · Kiwon Um (Technical University of Munich) · Nils Thuerey (Technical University of Munich)

FR-Train: A mutual information-based approach to fair and robust training
Yuji Roh (KAIST) · Kangwook Lee (UW Madison) · Steven Whang (KAIST) · Changho Suh (KAIST)

Real-Time Optimisation for Online Learning in Auctions
Lorenzo Croissant (Criteo) · Marc Abeille (Criteo) · Clement Calauzenes (Criteo)

Graph Random Neural Features for Distance-Preserving Graph Representations
Daniele Zambon (Università della Svizzera Italiana) · Cesare Alippi (Università della Svizzera Italiana) · Lorenzo Livi (University of Manitoba)

Modulating Surrogates for Bayesian Optimization
Erik Bodin (University of Bristol) · Markus Kaiser (Technical University Munich) · Ieva Kazlauskaite (University of Bath) · Zhenwen Dai (Spotify) · Neill Campbell (University of Bath) · Carl Henrik Ek (University of Bristol)

Convolutional Kernel Networks for Graph-Structured Data
Dexiong Chen (Inria) · Laurent Jacob (CNRS) · Julien Mairal (Inria)

Improving the Sample and Communication Complexity for Decentralized Non-Convex Optimization: Joint Gradient Estimation and Tracking
Haoran Sun (University of Minnesota) · Songtao Lu (IBM Research) · Mingyi Hong (University of Minnesota)

Proper Network Interpretability Helps Adversarial Robustness in Classification
Akhilan Boopathy (MIT) · Sijia Liu (MIT-IBM Watson AI Lab) · Gaoyuan Zhang (IBM Research) · Cynthia Liu (1998) · Pin-Yu Chen (IBM Research AI) · Shiyu Chang (MIT-IBM Watson AI Lab) · Luca Daniel (Massachusetts Institute of Technology)

Generalization Guarantees for Sparse Kernel Approximation with Entropic Optimal Features
Liang Ding (Texas A&M University) · Rui Tuo (Texas A&M University) · Shahin Shahrampour (Texas A&M University)

Understanding the Impact of Model Incoherence on Convergence of Incremental SGD with Random Reshuffle
Shaocong Ma (University of Utah) · Yi Zhou (University of Utah)

Calibration, Entropy Rates, and Memory in Language Models
Mark Braverman (Princeton University) · Xinyi Chen () · Sham Kakade (University of Washington) · Karthik Narasimhan (Princeton) · Cyril Zhang (Princeton University) · Yi Zhang (Princeton University)

Learning Opinions in Social Networks
Vincent Conitzer (Duke) · Debmalya Panigrahi (Duke University) · Hanrui Zhang (Duke University)

Latent Variable Modelling with Hyperbolic Normalizing Flows
Joey Bose (McGill/Mila) · Ariella Smofsky (McGill University and Mila) · Renjie Liao (University of Toronto) · Prakash Panangaden () · Will Hamilton (McGill University and Mila)

StochasticRank: Global Optimization of Scale-Free Discrete Functions
Aleksei Ustimenko (Yandex) · Liudmila Prokhorenkova (Yandex)

Working Memory Graphs
Ricky Loynd (Microsoft Research) · Roland Fernandez (Microsoft Research) · Asli Celikyilmaz (Microsoft Research) · Adith Swaminathan (Microsoft Research) · Matthew Hausknecht (Microsoft Research)

Learning to Combine Top-Down and Bottom-Up Signals in Recurrent Neural Networks with Attention over Modules
Sarthak Mittal (Uber ATG) · Alex Lamb (Universite de Montreal) · Anirudh Goyal (Université de Montréal) · Vikram Voleti (Mila, University of Montreal) · Murray Shanahan (DeepMind / Imperial College London) · Guillaume Lajoie (Mila, Université de Montréal) · Michael Mozer (Google Research / University of Colorado) · Yoshua Bengio (Mila / U. Montreal)

Spread Divergence
Mingtian Zhang (UCL) · Peter Hayes (University College London) · Thomas Bird (UCL) · Raza Habib (UCL) · David Barber (University College London)

Optimizing Black-box Metrics with Adaptive Surrogates
Qijia Jiang (Stanford University) · Olaoluwa Adigun (University of Southern California Los Angeles) · Harikrishna Narasimhan (Google Research) · Mahdi Milani Fard (Google) · Maya Gupta (Google)

Domain Adaptive Imitation Learning
Kuno Kim (Stanford University) · Yihong Gu (Tsinghua University) · Jiaming Song (Stanford) · Shengjia Zhao (Stanford University) · Stefano Ermon (Stanford University)

A general recurrent state space framework for modeling neural dynamics during decision-making
David Zoltowski (Princeton University) · Jonathan Pillow (Princeton University) · Scott Linderman (Stanford)

An Imitation Learning Approach for Cache Replacement
Evan Liu (Google) · Milad Hashemi (Google) · Kevin Swersky (Google Brain) · Parthasarathy Ranganathan (Google, USA) · Junwhan Ahn (Google)

Revisiting Training Strategies and Generalization Performance in Deep Metric Learning
Karsten Roth (Heidelberg University) · Timo Milbich (Heidelberg University) · Samrath Sinha (University of Toronto) · Prateek Gupta (University of Oxford) · Bjorn Ommer (Heidelberg University) · Joseph Paul Cohen (Mila, University of Montreal)

Temporal Phenotyping using Deep Predicting Clustering of Disease Progression
Changhee Lee (UCLA) · M van der Schaar (UCLA)

Countering Language Drift with Seeded Iterated Learning
Yuchen Lu (Mila & University of Montreal) · Soumye Singhal (Mila, University of Montreal) · Florian Strub (DeepMind) · Aaron Courville (Université de Montréal) · Olivier Pietquin (Google Brain)

Stochastic Gauss-Newton Algorithms for Nonconvex Compositional Optimization
Quoc Tran-Dinh (The University of North Carolina at Chapel Hill) · Nhan H Pham (University of North Carolina at Chapel Hill) · Lam Nguyen (IBM Research, Thomas J. Watson Research Center)

Strategyproof Mean Estimation from Multiple-Choice Questions
Anson Kahng (Carnegie Mellon University) · Gregory Kehne (Carnegie Mellon University) · Ariel Procaccia (Harvard University)

Sequential Cooperative Bayesian Inference
Junqi Wang (Rutgers University-Newark) · Pei Wang (Rutgers University-Newark) · Patrick Shafto (Rutgers University-Newark)

Spectral Graph Matching and Regularized Quadratic Relaxations: Algorithm and Theory
Zhou Fan (Yale Univ) · Cheng Mao (Georgia Institute of Technology) · Yihong Wu (Yale University) · Jiaming Xu (Duke University)

Zeno++: Robust Fully Asynchronous SGD
Cong Xie (UIUC) · Sanmi Koyejo (Illinois / Google) · Indranil Gupta (UIUC)

Network Pruning by Greedy Subnetwork Selection
Mao Ye (PURDUE UNIVERSITY) · Chengyue Gong (university of texas at austin) · Lizhen Nie (The University of Chicago) · Denny Zhou (Google Brain) · Adam Klivans (University of Texas at Austin) · Qiang Liu (UT Austin)

Logarithmic Regret for Learning Linear Quadratic Regulators Efficiently
Asaf Cassel (Tel Aviv University) · Alon Cohen (Technion, Google) · Tomer Koren (Google Brain)

Hierarchical Verification for Adversarial Robustness
Cong Han Lim (Uber ATG) · Raquel Urtasun (Uber ATG) · Ersin Yumer (Uber ATG)

BINOCULARS for efficient, nonmyopic sequential experimental design
Shali Jiang (Washington University in St. Louis) · Henry Chai (Washington University in St. Louis) · Javier Gonzalez (Microsoft Research) · Roman Garnett (Washington University in St. Louis)

On the Global Optimality of Model-Agnostic Meta-Learning
Lingxiao Wang (Northwestern University) · Qi Cai (Northwestern University) · Zhuoran Yang (Princeton University) · Zhaoran Wang (Northwestern U)

Breaking the Curse of Many Agents: Provable Mean Embedding Q -Iteration for Mean-Field Reinforcement Learning
Lingxiao Wang (Northwestern University) · Zhuoran Yang (Princeton University) · Zhaoran Wang (Northwestern U)

Learning with Bounded Instance- and Label-dependent Label Noise
Jiacheng Cheng (University of California, San Diego) · Tongliang Liu (The University of Sydney) · Kotagiri Ramamohanarao (The University of Melbourne) · Dacheng Tao (The University of Sydney)

Transparency Promotion with Model-Agnostic Linear Competitors
Hassan Rafique (The University of Iowa) · Tong Wang (University of Iowa) · Qihang Lin (University of Iowa) · Arshia Singhani (BASIS Independent Silicon Valley)

Learning Mixtures of Graphs from Epidemic Cascades
Jessica Hoffmann (University of Texas at Austin) · Soumya Basu (University of Texas at Austin) · Surbhi Goel (UT Austin) · Constantine Caramanis (University of Texas)

Implicit differentiation of Lasso-type models for hyperparameter optimization
Quentin Bertrand (INRIA) · Quentin Klopfenstein (Université de Bourgogne) · Mathieu Blondel (NTT) · Samuel Vaiter (CNRS) · Alexandre Gramfort (Inria) · Joseph Salmon (Université de Montpellier)

Latent Space Factorisation and Manipulation via Matrix Subspace Projection
Xiao Li (University of Aberdeen) · Chenghua Lin (University of Sheffield) · Ruizhe Li (The University of Sheffield) · Chaozheng Wang (University of Aberdeen) · Frank Guerin (University of Surrey)

Active World Model Learning in Agent-rich Environments with Progress Curiosity
Kuno Kim (Stanford University) · Megumi Sano (Stanford University) · Julian De Freitas (Harvard University) · Nick Haber (Stanford University) · Daniel Yamins (Stanford University)

SDE-Net: Equipping Deep Neural Networks with Uncertainty Estimates
Lingkai Kong (Georgia Institute of Technoloy) · Jimeng Sun (UIUC) · Chao Zhang (Georgia Institute of Technology)

GANs May Have No Nash Equilibria
Farzan Farnia (Massachusetts Institute of Technology) · Asuman Ozdaglar (MIT)

Gradient Temporal-Difference Learning with Regularized Corrections
Sina Ghiassian (University of Alberta) · Andrew Patterson (University of Alberta) · Shivam Garg (University of alberta) · Dhawal Gutpa (University of Alberta) · Adam White (University of Alberta) · Martha White (University of Alberta)

Online mirror descent and dual averaging: keeping pace in the dynamic case
Huang Fang (University of British Columbia) · Victor Sanches Portella (University of British Columbia) · Nick Harvey (University of British Columbia) · Michael Friedlander (University of British Columbia)

Choice Set Optimization Under Discrete Choice Models of Group Decisions
Kiran Tomlinson (Cornell University) · Austin Benson (Cornell University)

Complexity of Finding Stationary Points of Nonconvex Nonsmooth Functions
Jingzhao Zhang (MIT) · Hongzhou Lin (MIT) · Stefanie Jegelka (Massachusetts Institute of Technology) · Suvrit Sra (MIT) · Ali Jadbabaie (Massachusetts Institute of Technology)

Multi-Agent Routing Value Iteration Network
Quinlan Sykora (Uber ATG) · Mengye Ren (Uber ATG / University of Toronto) · Raquel Urtasun (Uber ATG)

Adversarial Attacks on Copyright Detection Systems
Parsa Saadatpanah (University of Maryland) · Ali Shafahi (University of Maryland) · Tom Goldstein (University of Maryland)

Differentiating through the Fréchet Mean
Aaron Lou (Cornell University) · Isay Katsman (Cornell University) · Qingxuan Jiang (Cornell University) · Serge Belongie (Cornell University) · Ser Nam Lim (Facebook) · Christopher De Sa (Cornell)

Online Learning for Active Cache Synchronization
Andrey Kolobov (Microsoft Research) · Sebastien Bubeck (Microsoft Research) · Julian Zimmert (University of Copenhagen)

PoKED: A Semi-Supervised System for Word Sense Disambiguation
Feng Wei (York University)

A Finite-Time Analysis of Q-Learning with Neural Network Function Approximation
Pan Xu (University of California, Los Angeles) · Quanquan Gu (University of California, Los Angeles)

Understanding and Stabilizing GANs' Training Dynamics Using Control Theory
Kun Xu (Tsinghua University) · Chongxuan Li (Tsinghua University) · Jun Zhu (Tsinghua University) · Bo Zhang (Tsinghua University)

Scalable Nearest Neighbor Search for Optimal Transport
Arturs Backurs (TTIC) · Yihe Dong (Microsoft Research) · Piotr Indyk (MIT) · Ilya Razenshteyn (Microsoft Research Redmond) · Tal Wagner (MIT)

Supervised learning: no loss no cry
Richard Nock (Data61, The Australian National University and the University of Sydney) · Aditya Menon (Google Research)

Label-Noise Robust Domain Adaptation
Xiyu Yu (Baidu Inc.) · Tongliang Liu (The University of Sydney) · Mingming Gong (University of Melbourne) · Kun Zhang (Carnegie Mellon University) · Kayhan Batmanghelich (University of Pittsburgh) · Dacheng Tao (The University of Sydney)

Description Based Text Classification with Reinforcement Learning
Wei Wu (Shannon.AI) · Duo Chai (Shannon.AI) · Qinghong Han (Shannon.AI) · Fei Wu (Zhejiang University, China) · Jiwei Li (Shannon.AI)

Bandits for BMO Functions
Tianyu Wang (Duke University) · Cynthia Rudin (Duke)

Cost-effectively Identifying Causal Effect When Only Response Variable Observable
Tian-Zuo Wang (Nanjing University) · Xi-Zhu Wu (Nanjing University) · Sheng-Jun Huang (Nanjing University of Aeronautics and Astronautics) · Zhi-Hua Zhou (Nanjing University)

Learning with Multiple Complementary Labels
LEI FENG (Nanyang Technological University) · Takuo Kaneko (The University of Tokyo) · Bo Han (HKBU / RIKEN) · Gang Niu (RIKEN) · Bo An (Nanyang Technological University) · Masashi Sugiyama (RIKEN / The University of Tokyo)

Graph Representation Learning by Maximizing Mutual Information Between Spatial and Spectral Views
Kaveh Hassani (Autodesk) · Amir Hosein Khasahmadi (University of Toronto)

A Chance-Constrained Generative Framework for Sequence Optimization
Liu Xianggen (Tsinghua University) · Jian Peng (UIUC) · Qiang Liu (UT Austin) · Sen Song (Tsinghua University)

dS^2LBI: Exploring Structural Sparsity on Deep Network via Differential Inclusion Paths
Yanwei Fu (Fudan university) · Chen Liu (Fudan University) · Donghao Li (HKUST) · Xinwei Sun (MSRA) · Jinshan ZENG (Hongkong University of Science and Technology) · Yuan Yao (HongKong University of Science and Technology)

Sparse Subspace Clustering with Entropy-Norm
Liang Bai (Shanxi University, China) · Jiye Liang (Shanxi University)

On the Generalization Effects of Linear Transformations in Data Augmentation
Sen Wu (Stanford University) · Hongyang Zhang (University of Pennsylvania) · Gregory Valiant (Stanford University) · Christopher Re (Stanford)

Sparse Shrunk Additive Models
Hong Chen (Huazhong Agricultural University) · guodong liu (university of Pittsburgh) · Heng Huang (University of Pittsburgh)

Unsupervised Discovery of Interpretable Directions in the GAN Latent Space
Andrey Voynov (Yandex) · Artem Babenko (Yandex)

DropNet: Reducing Neural Network Complexity via Iterative Pruning
Chong Min John Tan (National University of Singapore) · Mehul Motani (NUS)

Self-supervised Label Augmentation via Input Transformations
Hankook Lee (KAIST) · Sung Ju Hwang (KAIST, AITRICS) · Jinwoo Shin (KAIST)

Mapping natural-language problems to formal-language solutions using structured neural representations
Kezhen Chen (Northwestern University) · Qiuyuan Huang (Microsoft Research, Redmond) · Hamid Palangi (Microsoft Research) · Paul Smolensky (Microsoft Research) · Ken Forbus (Northwestern University) · Jianfeng Gao (Microsoft Research AI)

Transformation of ReLU-based recurrent neural networks from discrete-time to continuous-time
Zahra Monfared (ZI Mannheim) · Daniel Durstewitz (ZI Mannheim)

Implicit Geometric Regularization for Learning Shapes
Amos Gropp (Weizmann Institute of Science) · Lior Yariv (Weizmann Institute of Science) · Niv Haim (Weizmann Institute of Science) · Matan Atzmon (Weizmann Institute of Science) · Yaron Lipman (Weizmann Institute of Science)

Influence Diagram Bandits
Tong Yu (Carnegie Mellon University) · Branislav Kveton (Google Research) · Zheng Wen (DeepMind) · Ruiyi Zhang (Duke University) · Ole J. Mengshoel (Carnegie Mellon University)

Information Particle Filter Tree: An Online Algorithm for POMDPs with Belief-Based Rewards on Continuous Domains
Johannes Fischer (Karlsruhe Institute of Technology (KIT)) · Ömer Sahin Tas (Karlsruhe Institute of Technology (KIT))

Convergence Rates of Variational Inference in Sparse Deep Learning
Badr-Eddine Chérief-Abdellatif (CREST)

Unsupervised Transfer Learning for Spatiotemporal Predictive Networks
Zhiyu Yao (Tsinghua University) · Yunbo Wang (Tsinghua University) · Mingsheng Long (Tsinghua University) · Jianmin Wang (Tsinghua University)

DINO: Distributed Newton-Type Optimization Method
Rixon Crane (The University of Queensland) · Fred Roosta (University of Queensland)

Quantum Expectation-Maximization for Gaussian Mixture Models
Alessandro Luongo (Université Paris Diderot) · Iordanis Kerenidis (Université Paris Diderot) · Anupam Prakash (Université Paris Diderot)

Consistent Structured Prediction with Max-Min Margin Markov Networks
Alex Nowak (INRIA - Ecole Normale Supérieure) · Francis Bach (INRIA - Ecole Normale Supérieure) · Alessandro Rudi (INRIA, École Normale Supérieure)

Concentration bounds for CVaR estimation: The cases of light-tailed and heavy-tailed distributions
Prashanth L.A. (IIT Madras) · Krishna Jagannathan (Indian Institute of Technology Madras) · Ravi Kolla (1989)

Robust Pricing in Dynamic Mechanism Design
Yuan Deng (Duke University) · Sébastien Lahaie (Google) · Vahab Mirrokni (Google Research)

Nested Subspace Arrangement for Representation of Relational Data
Nozomi Hata (Kyushu University) · Shizuo Kaji (Kyushu University) · Akihiro Yoshida (Kyushu University) · Katsuki Fujisawa (Kyushu University)

Equivariant Neural Rendering
Emilien Dupont (University of Oxford) · Miguel Bautista Martin (Apple Inc.) · Alex Colburn (Apple Inc.) · Aditya Sankar (Apple Inc.) · Joshua M Susskind (Apple, Inc.) · Qi Shan (Apple Inc)

Bounding the fairness and accuracy of classifiers from population statistics
Sivan Sabato (Ben-Gurion University of the Negev) · Elad Yom-Tov (Microsoft Research)

Healing Gaussian Process Experts
samuel cohen (University College London) · Rendani Mbuvha (University of Witwatersrand) · Tshilidzi Marwala (University of Johannesburg) · Marc Deisenroth (University College London)

Beyond UCB: Optimal and Efficient Contextual Bandits with Regression Oracles
Dylan Foster (MIT) · Alexander Rakhlin (MIT)

Simple and Deep Graph Convolutional Networks
Ming Chen (Renmin University of China) · Zhewei Wei (Renmin University of China) · Zengfeng Huang (Fudan University) · Bolin Ding ("Data Analytics and Intelligence Lab, Alibaba Group") · Yaliang Li (Alibaba Group)

Projection-free Distributed Online Convex Optimization with $O(\sqrt{T})$ Communication Complexity
Yuanyu Wan (Nanjing University) · Wei-Wei Tu (4Paradigm Inc.) · Lijun Zhang (Nanjing University)

Meta Variance Transfer: Learning to Augment from the Others
Seong-Jin Park (Samsung Advanced Institute of Technology) · Seungju Han (Samsung Advanced Institute of Technology) · Ji-won Baek (Samsung Advanced Institute of Technology) · Insoo Kim (Samsung Advanced Institute of Technology) · Juhwan Song (Samsung Advanced Institute of Technology) · Hae Beom Lee (KAIST) · Jae-Joon Han (Samsung) · Sung Ju Hwang (KAIST, AITRICS)

Coresets for Clustering in Graphs of Bounded Treewidth
Daniel Baker (Johns Hopkins University) · Vladimir Braverman (Johns Hopkins University) · Lingxiao Huang (Yale University) · Shaofeng H.-C. Jiang (Weizmann Institute of Science) · Robert Krauthgamer (Weizmann Institute of Science) · Xuan Wu (Johns Hopkins University)

On Breaking Deep Generative Model-based Defenses and Beyond
Yanzhi Chen (University of Edinburgh) · Renjie Xie (Southeast University) · Zhanxing Zhu (Peking University)

Exploration Through Bias: Revisiting Biased Maximum Likelihood Estimation in Stochastic Multi-Armed Bandits
Xi Liu (Texas A&M University) · Ping-Chun Hsieh (National Chiao Tung University) · Yu Heng Hung (NCTU) · Anirban Bhattacharya (Texas A&M University) · P. Kumar (Texas A&M University)

Bisection-Based Pricing for Repeated Contextual Auctions against Strategic Buyer
Anton Zhiyanov (Yandex Research, MSU im. Lomonosova) · Alexey Drutsa (Yandex)

Haar Graph Pooling
Yuguang Wang (University of New South Wales) · Ming Li (Zhejiang Normal University) · Zheng Ma (Princeton University) · Guido Montufar (UCLA Math / Stat; MPI MIS) · Xiaosheng Zhuang (City University of Hong Kong) · Yanan Fan (University of New South Wales)

Explaining Groups of Points in Low-Dimensional Representations
Gregory Plumb (Carnegie Mellon University) · Jonathan Terhorst (U-M LSA) · Sriram Sankararaman (UCLA) · Ameet Talwalkar (Carnegie Mellon University)

Learning Portable Representations for High-Level Planning
Steven James (University of the Witwatersrand) · Benjamin Rosman (University of the Witwatersrand / CSIR, South Africa) · George Konidaris (Brown)

Adaptive Estimator Selection for Off-Policy Evaluation
Yi Su (Cornell University) · Pavithra Srinath (Microsoft Research) · Akshay Krishnamurthy (Microsoft Research)

Doubly Stochastic Variational Inference for Neural Processes with Hierarchical Latent Variables
Qi Wang (Informatics Institute, University of Amsterdam) · Herke van Hoof (University of Amsterdam)

Generative Flows with Matrix Exponential
Changyi Xiao (University of Science and Technology of China) · Ligang Liu (University of Science and Technology of China)

Composable Sketches for Functions of Frequencies: Beyond the Worst Case
Edith Cohen (Google Research and Tel Aviv University) · Ofir Geri (Stanford University) · Rasmus Pagh (IT University of Copenhagen)

Self-concordant analysis of Frank-Wolfe algorithm
Mathias Staudigl (Maastricht University) · Pavel Dvurechenskii (Weierstrass Institute) · Shimrit Shtein (Technion) · Kamil Safin (Moscow Institute of Physics and Technology) · Petr Ostroukhov (Moscow Institute of Physics and Technology)

Towards non-parametric drift detection via Dynamic Adapting Window Independence Drift Detection (DAWIDD)
Fabian Hinder (CITEC, Bielefeld University) · André Artelt (Bielefeld University) · CITEC Barbara Hammer (CITEC, Bielefeld University)

Non-Stationary Bandits with Intermediate Observations
Claire Vernade (DeepMind) · Andras Gyorgy (DeepMind) · Timothy Mann (DeepMind)

Does label smoothing mitigate label noise?
Lukasik Michal (Google Research) · Srinadh Bhojanapalli (Google AI) · Aditya Menon (Google Research) · Sanjiv Kumar (Google Research, NY)

Proving the Lottery Ticket Hypothesis: Pruning is All You Need
Gilad Yehudai (Weizmann Institute of Science) · Eran Malach (Hebrew University Jerusalem Israel) · Shai Shalev-Schwartz (Hebrew University of Jerusalem) · Ohad Shamir (Weizmann Institute of Science)

Linear bandits with Stochastic Delayed Feedback
Claire Vernade (DeepMind) · Alexandra Carpentier (Otto-von-Guericke University) · Tor Lattimore (DeepMind) · Giovanni Zappella (Amazon) · Beyza Ermis (Amazon Research) · Michael Brueckner (Amazon Research Berlin)

Time Series Deconfounder: Estimating Treatment Effects over Time in the Presence of Hidden Confounders
Ioana Bica (University of Oxford) · Ahmed Alaa (UCLA) · Mihaela van der Schaar (University of Cambridge, The Alan Turing Institute, UCLA)

Negative Sampling in Semi-Supervised learning
John Chen (Rice University) · Vatsal Shah (University of Texas at Austin) · Anastasios Kyrillidis (Rice University)

Adaptive Sketching for Fast and Convergent Canonical Polyadic Decomposition
Alex Gittens (Rensselaer Polytechnic Institute) · Kareem Aggour (GE Research) · Bülent Yener (Rensselaer Polytechnic Institute)

Private Counting from Anonymous Messages: Near-Optimal Accuracy with Vanishing Communication Overhead
Badih Ghazi (Google) · Ravi Kumar (Google) · Pasin Manurangsi (Google) · Rasmus Pagh (IT University of Copenhagen)

On the generalization benefit of noise in stochastic gradient descent
Samuel L Smith (DeepMind) · Erich Elsen (Google) · Soham De (DeepMind)

Momentum-Based Policy Gradient Methods
Feihu Huang (University of Pittsburgh) · Shangqian Gao (University of Pittsburgh) · Jian Pei (Simon Fraser University) · Heng Huang (University of Pittsburgh)

Knowing The What But Not The Where in Bayesian Optimization
Vu Nguyen (University of Oxford) · Michael A Osborne (U Oxford)

Robust Bayesian Classification Using An Optimistic Score Ratio
Viet Anh Nguyen (Stanford University) · Nian Si (Stanford University) · Jose Blanchet (Stanford University)

Boosted Histogram Transform for Regression
Yuchao Cai (AI Lab, Samsung Research China - Beijing) · Hanyuan Hang (AI Lab, Samsung Research China - Beijing) · Hanfang Yang (Renmin university of China) · Zhouchen Lin (Peking University)

Stochastic bandits with arm-dependent delays
Anne Gael Manegueu (Otto-von-Guerricke University) · Claire Vernade (Amazon Research) · Alexandra Carpentier (Otto-von-Guericke University) · Michal Valko (DeepMind)

Projective Preferential Bayesian Optimization
Petrus Mikkola (Aalto University) · Milica Todorović (Aalto University) · Jari Järvi (Aalto University) · Patrick Rinke (Aalto University) · Samuel Kaski (Aalto University)

On Relativistic f-Divergences
Alexia Jolicoeur-Martineau (Mila)

A Flexible Framework for Nonparametric Graphical Modeling that Accommodates Machine Learning
Yunhua Xiang (University of Washington) · Noah Simon (University of Washington)

The Natural Lottery Ticket Winner: Reinforcement Learning with Ordinary Neural Circuits
Ramin Hasani (TU Wien) · Mathias Lechner (IST Austria) · Alexander Amini (MIT) · Daniela Rus (MIT CSAIL) · Radu Grosu (TU Wien)

Schatten Norms in Matrix Streams: Hello Sparsity, Goodbye Dimension
Aditya Krishnan (Johns Hopkins University) · Roi Sinoff (Weizmann Institute of Science, Israel) · Robert Krauthgamer (Weizmann Institute of Science) · Vladimir Braverman (Johns Hopkins University)

Control Frequency Adaptation via Action Persistence in Batch Reinforcement Learning
Alberto Maria Metelli (Politecnico di Milano) · Flavio Mazzolini (Politecnico di Milano) · Lorenzo Bisi (Politecnico di Milano) · Luca Sabbioni (Politecnico di Milano) · Marcello Restelli (Politecnico di Milano)

Minimax Rate for Learning From Pairwise Comparisons in the BTL Model
Julien Hendrickx (University of Catholique de Louvain) · Alex Olshevsky (Boston University) · Venkatesh Saligrama (Boston University)

Interferometric Graph Transform:a Deep Unsupervised Graph Invariant Representation
Edouard Oyallon (CNRS/LIP6)

Stochastic Differential Equations with variational Wishart diffusions
Martin Jørgensen (Technical University of Denmark) · Marc Deisenroth (University College London) · Hugh Salimbeni (Imperial College London)

What Can Learned Intrinsic Rewards Capture?
Zeyu Zheng (University of Michigan) · Junhyuk Oh (DeepMind) · Matteo Hessel (Deep Mind) · Zhongwen Xu (DeepMind) · Manuel Kroiss (DeepMind) · Hado van Hasselt (DeepMind) · David Silver (Google DeepMind) · Satinder Singh (DeepMind)

Random extrapolation for primal-dual coordinate descent
Ahmet Alacaoglu (EPFL) · Olivier Fercoq (Telecom Paris) · Volkan Cevher (EPFL)

Reinforcement Learning with Differential Privacy
Giuseppe Vietri (University of Minnesota) · Borja de Balle Pigem (Amazon Research) · Steven Wu (University of Minnesota) · Akshay Krishnamurthy (Microsoft Research)

Median Matrix Completion: from Embarrassment to Optimality
Weidong Liu (Shanghai Jiaotong University) · Xiaojun Mao (Fudan University) · Raymond K. W. Wong (Texas A&M University)

Improved Optimistic Algorithms for Logistic Bandits
Louis Faury (Criteo) · Marc Abeille (Criteo) · Clement Calauzenes (Criteo) · Olivier Fercoq (Telecom Paris)

Learning to Rank Learning Curves
Martin Wistuba (IBM Research) · Tejaswini Pedapati (IBM Research)

Model Fusion with Kullback--Leibler Divergence
Sebastian Claici (MIT) · Mikhail Yurochkin (IBM Research AI) · Soumya Ghosh (IBM Research) · Justin Solomon (MIT)

Randomization matters How to defend against strong adversarial attacks
Rafael Pinot (Dauphine University - CEA LIST) · Raphael Ettedgui (Paris-Dauphine University) · Geovani Rizk (Université Paris Dauphine) · Yann Chevaleyre (Univ. Paris Dauphine) · Jamal Atif (Université Paris-Dauphine)

Evolutionary Topology Search for Tensor Network Decomposition
Chao Li (RIKEN Center for Advanced Intelligence Project) · Zhun Sun (RIKEN)

Quadratically Regularized Subgradient Methods for Weakly Convex Optimization with Weakly Convex Constraints
Runchao Ma (University of Iowa) · Qihang Lin (University of Iowa) · Tianbao Yang (The University of Iowa)

Scalable and Efficient Comparison-based Search without Features
Daniyar Chumbalov (EPFL) · Lucas Maystre (Spotify) · Matthias Grossglauser (EPFL)

Error-Bounded Correction of Noisy Labels
Songzhu Zheng (Stony Brook University) · Pengxiang Wu (Rutgers University) · Aman Goswami (Bain and Company) · Mayank Goswami (Queens College of CUNY) · Dimitris Metaxas (Rutgers) · Chao Chen (Stony Brook University)

Learning with Feature and Distribution Evolvable Streams
Zhen-Yu Zhang (Nanjing University) · Peng Zhao (Nanjing University) · Yuan Jiang (Nanjing University) · Zhi-Hua Zhou (Nanjing University)

On Unbalanced Optimal Transport: An Analysis of Sinkhorn Algorithm
Khiem Pham (VinAI Research) · Khang Le (VinAI Research) · Nhat Ho (University of California, Berkeley) · Tung Pham (VinAI Research) · Hung Bui (VinAI Research)

Learning Optimal Tree Models under Beam Search
Jingwei Zhuo (Alibaba Group) · Ziru Xu (Alibaba Group) · Wei Dai (Alibaba Group) · Han Zhu (Tsinghua University) · HAN LI (Alibaba Group) · Jian Xu (Alibaba Group) · Kun Gai (Alibaba Group)

Estimating the number and effect sizes of non-null hypotheses
Jennifer Brennan (University of Washington) · Ramya Korlakai Vinayak (University of Washington) · Kevin Jamieson (University of Washington)

Estimating Model Uncertainty of Neural Network in Sparse Information Form
Jongseok Lee (German Aerospace Center (DLR)) · Matthias Humt (Deutsches Zentrum für Luft- und Raumfahrt (DLR) - German Aerospace Center) · Jianxiang Feng (German Aerospace Center (DLR)) · Rudolph Triebel (German Aerospace Center (DLR))

Double-Loop Unadjusted Langevin Algorithm
Paul Rolland (Ecole Polytechnique Fédérale de Lausanne) · Armin Eftekhari (Umea University) · Ali Kavis (EPFL) · Volkan Cevher (EPFL)

Growing Action Spaces
Gregory Farquhar (University of Oxford) · Laura Gustafson (Facebook AI Research) · Zeming Lin (Facebook AI Reseach) · Shimon Whiteson (Oxford University) · Nicolas Usunier (Facebook AI Research) · Gabriel Synnaeve (Facebook AI Research)

Analytic Marching: An Analytic Meshing Solution from Deep Implicit Surface Networks
Jiabao Lei (South China University of Technology) · Kui Jia (South China University of Technology)

Anderson Acceleration of Proximal Gradient Methods
Vien Van Mai (KTH Royal Institute of Technology) · Mikael Johansson (KTH Royal Institute of Technology)

Interpretable, Multidimensional, Multimodal Anomaly Detection with Negative Sampling for Detection of Device Failure
John Sipple (Google Inc)

Certified Robustness to Label-Flipping Attacks via Randomized Smoothing
Elan Rosenfeld (Carnegie Mellon University) · Ezra Winston () · Pradeep Ravikumar (Carnegie Mellon University) · Zico Kolter (Carnegie Mellon University / Bosch Center for AI)

Responsive Safety in Reinforcement Learning
Adam Stooke (UC Berkeley) · Joshua Achiam (OpenAI) · Pieter Abbeel (UC Berkeley & Covariant)

Deep k-NN for Noisy Labels
Dara Bahri (Google) · Heinrich Jiang (Google Research) · Maya Gupta (Google)

Learning the piece-wise constant graph structure of a varying Ising model
Batiste Le Bars (CMLA - Ens Paris-Saclay) · Pierre Humbert (Ecole Normale Supérieure Paris-Saclay, Université Paris Saclay) · Argyris Kalogeratos (CMLA - ENS Paris-Saclay) · Nicolas Vayatis (ENS Cachan)

Stabilizing Transformers for Reinforcement Learning
Emilio Parisotto (Carnegie Mellon University) · Francis Song (DeepMind) · Jack Rae (DeepMind) · Razvan Pascanu (DeepMind) · Caglar Gulcehre (DeepMind) · Siddhant Jayakumar (DeepMind) · Max Jaderberg (DeepMind) · Raphael Lopez Kaufman (Deepmind) · Aidan Clark (DeepMind) · Seb Noury (DeepMind) · Matthew Botvinick (DeepMind) · Nicolas Heess (DeepMind) · Raia Hadsell (DeepMind)

An Explicitly Relational Neural Network Architecture
Murray Shanahan (DeepMind / Imperial College London) · Kyriacos Nikiforou (DeepMind) · Antoina Creswell (Imperial College London) · Christos Kaplanis (DeepMind Technologies Ltd) · David GT Barrett (DeepMind) · Marta Garnelo (DeepMind)

Harmonic Decompositions of Convolutional Networks
Meyer Scetbon (ENSAE) · Zaid Harchaoui (University of Washington)

Discriminative Jackknife: Quantifying Uncertainty in Deep Learning via Higher-Order Influence Functions
Ahmed Alaa (UCLA) · M van der Schaar (UCLA)

Robust Graph Representation Learning via Neural Sparsification
Cheng Zheng (UCLA) · Bo Zong (NEC Labs) · Wei Cheng (NEC Laboratories America) · Dongjin Song (NEC Labs America) · Jingchao Ni (NEC Laboratories America, Inc.) · Wenchao Yu (UCLA) · Haifeng Chen (NEC Labs) · Wei Wang (UCLA)

Semiparametric Nonlinear Bipartite Graph Representation Learning with Provable Guarantees
Sen Na (The University of Chicago) · Yuwei Luo (The University of Chicago) · Zhuoran Yang (Princeton University) · Zhaoran Wang (Northwestern U) · Mladen Kolar (University of Chicago Booth School of Business)

Forecasting sequential data using Consistent Koopman Autoencoders
Omri Azencot (UCLA) · N. Benjamin Erichson (University of California, Berkeley) · Vanessa Lin (University of California Berkeley) · Michael Mahoney (UC Berkeley)

Scalable Identification of Partially Observed Systems with Certainty-Equivalent EM
Kunal Menda (Stanford University) · Jean de Becdelievre (Stanford University) · Jayesh Gupta (Stanford University) · Ilan Kroo (Stanford University) · Mykel Kochenderfer (Stanford University) · Zachary Manchester (Stanford)

Learning to Score Behaviors for Guided Policy Optimization
Aldo Pacchiano (UC Berkeley) · Jack Parker-Holder (University of Oxford) · Yunhao Tang (Columbia University) · Krzysztof Choromanski (Google) · Anna Choromanska (NYU Tandon School of Engineering) · Michael Jordan (UC Berkeley)

Improved Communication Cost in Distributed PageRank Computation – A Theoretical Study
Siqiang Luo (Harvard)

Learning Autoencoders with Relational Regularization
Hongteng Xu (InfiniaML, Inc.) · Dixin Luo (Duke University) · Ricardo Henao (Duke University) · Svati Shah (Duke University) · Lawrence Carin (Duke)

Neural Contextual Bandits with UCB-based Exploration
Dongruo Zhou (UCLA) · Lihong Li (Google Research) · Quanquan Gu (University of California, Los Angeles)

Super-efficiency of automatic differentiation for functions defined as a minimum
Pierre Ablin (CNRS and ENS) · Gabriel Peyré (CNRS and ENS) · Thomas Moreau (Inria)

Rethinking Batch Normalization in Transformers
Sheng Shen (University of California, Berkeley) · Zhewei Yao (University of California, Berkeley) · Amir Gholaminejad (University of California, Berkeley) · Michael Mahoney (UC Berkeley) · Kurt Keutzer (UC Berkeley)

Invertible generative models for inverse problems: mitigating representation error and dataset bias
Muhammad Asim (Information Technology University, Lahore) · Max Daniels (Northeastern University) · Oscar Leong (Rice University) · Paul Hand (Northeastern University) · Ali Ahmed (Information Technology University)

Acceleration for Compressed Gradient Descent in Distributed Optimization
Zhize Li (King Abdullah University of Science and Technology) · Dmitry Kovalev (KAUST) · Xun Qian (KAUST) · Peter Richtarik (KAUST)

Neural Networks are Convex Regularizers: Exact Polynomial-time Convex Optimization Formulations for Two-Layer Networks
Mert Pilanci (Stanford) · Tolga Ergen (Stanford University)

Learning Quadratic Games on Networks
Yan Leng (MIT Media Lab) · Xiaowen Dong (University of Oxford) · Junfeng Wu (Zhejiang University) · Alex `Sandy' Pentland (MIT)

Margin-aware Adversarial Domain Adaptation with Optimal Transport
Sofien Dhouib (CREATIS UMR CNRS 5220) · Ievgen Redko (Laboratoire Hubert Curien) · Carole Lartizien (CREATIS)

The Sample Complexity of Best-k Items Selection from Pairwise Comparisons
Wenbo Ren (The Ohio State University) · Jia Liu (Iowa State University) · Ness Shroff (The Ohio State University)

GraphOpt: Learning Optimization Models of Graph Formation
Rakshit Trivedi (Georgia Institute of Technology) · Jiachen Yang (Georgia Institute of Technology) · Hongyuan Zha (Georgia Institute of Technology)

Distributionally Robust Policy Evaluation and Learning in Offline Contextual Bandits
Nian Si (Stanford University) · Fan Zhang (Stanford University) · Zhengyuan Zhou (Stanford University) · Jose Blanchet (Stanford University)

Incremental Sampling Without Replacement for Sequence Models
Kensen Shi (Google) · David Bieber (Google Brain) · Charles Sutton (Google)

Variable Skipping for Autoregressive Range Density Estimation
Eric Liang (University of California, Berkeley) · Zongheng Yang (UC Berkeley) · Ion Stoica (UC Berkeley) · Pieter Abbeel (UC Berkeley & Covariant) · Yan Duan (COVARIANT.AI) · Peter Chen (COVARIANT.AI)

TaskNorm: Rethinking Batch Normalization for Meta-Learning
John Bronskill (University of Cambridge) · Jonathan Gordon (University of Cambridge) · James Requeima (University of Cambridge) · Sebastian Nowozin (Google Research) · Richard E Turner (University of Cambridge)

Scalable Gaussian Process Regression for Kernels with a Non-Stationary Phase
Jan Graßhoff (Universität zu Lübeck) · Alexandra Jankowski (Institute for Electrical Engineering in Medicine, Universität zu Lübeck) · Philipp Rostalski (Universität zu Lübeck)

Transformer Hawkes Process
Simiao Zuo (Georgia Institute of Technology) · Haoming Jiang (Georgia Tech) · Zichong Li (University of Science and technology of China) · Tuo Zhao (Gatech) · Hongyuan Zha (Georgia Institute of Technology)

An EM Approach to Non-autoregressive Conditional Sequence Generation
Zhiqing Sun (Carnegie Mellon University) · Yiming Yang (Carnegie Mellon University)

Variance Reduction in Stochastic Particle-Optimization Sampling
Jianyi Zhang (Duke University) · Yang Zhao (University at Buffalo) · Changyou Chen (SUNY Buffalo)

CLUB: A Contrastive Log-ratio Upper Bound of Mutual Information
Pengyu Cheng (Duke University) · Weituo Hao (Duke University) · Shuyang Dai (Duke University) · Jiachang Liu (Duke University) · Zhe Gan (Microsoft) · Lawrence Carin (Duke)

State Space Expectation Propagation: Efficient Inference Schemes for Temporal Gaussian Processes
William Wilkinson (Aalto University) · Paul Chang (Aalto University) · Michael Andersen (Technical University of Denmark) · Arno Solin (Aalto University)

Training Neural Networks for and by Interpolation
Leonard Berrada (University of Oxford) · M. Pawan Kumar (University of Oxford) · Andrew Zisserman (University of Oxford, DeepMind)

Learning Representations that Support Extrapolation
Taylor Webb (UCLA Psychology Department) · Zachary Dulberg (Princeton University) · Steven M Frankland (Princeton University) · Alexander Petrov (The Ohio State University) · Randall O'Reilly (University of California Davis) · Jonathan Cohen (Princeton University)

Topic Modeling via Full Dependence Mixtures
Dan Fisher (Technion) · Mark Kozdoba (Technion) · Shie Mannor (Technion)

Instance-hiding Schemes for Private Distributed Learning
Yangsibo Huang (Princeton University) · Zhao Song (IAS/Princeton) · Sanjeev Arora (Princeton University and Institute for Advanced Study) · Kai Li (Princeton University)

The Implicit Regularization of Stochastic Gradient Flow for Least Squares
Alnur Ali (Carnegie Mellon University) · Edgar Dobriban (University of Pennsylvania) · Ryan Tibshirani (Carnegie Mellon University)

Decentralised Learning with Random Features and Distributed Gradient Descent
Dominic Richards (University of Oxford) · Patrick Rebeschini (University of Oxford) · Lorenzo Rosasco (unige, mit, iit)

Hierarchical Generation of Molecular Graphs using Structural Motifs
Wengong Jin (MIT) · Regina Barzilay (MIT CSAIL) · Tommi Jaakkola (MIT)

Composing Molecules with Multiple Property Constraints
Wengong Jin (MIT) · Regina Barzilay (MIT CSAIL) · Tommi Jaakkola (MIT)

Data preprocessing to mitigate bias: A maximum entropy based approach
Elisa Celis (Yale University) · Vijay Keswani (Yale University) · Nisheeth Vishnoi (Yale)

On Efficient Low Distortion Ultrametric Embedding
Vincent Cohen-Addad (CNRS & Sorbonne Université) · Karthik C. S. (Tel Aviv University) · Guillaume Lagarde (LaBRI)

Global Concavity and Optimization in a Class of Dynamic Discrete Choice Models
Yiding Feng (Northwestern University) · Ekaterina Khmelnitskaya (University of Virginia) · Denis Nekipelov ()

Efficient Policy Learning from Surrogate-Loss Classification Reductions
Andrew Bennett (Cornell University) · Nathan Kallus (Cornell University)

On Contrastive Learning for Likelihood-free Inference
Conor Durkan (University of Edinburgh) · Iain Murray (University of Edinburgh) · George Papamakarios (DeepMind)

Obtaining Adjustable Regularization for Free via Iterate Averaging
Jingfeng Wu (Johns Hopkins University) · Vladimir Braverman (Johns Hopkins University) · Lin Yang (UCLA)

Invariant Risk Minimization Games
Kartik Ahuja (IBM Research) · Karthikeyan Shanmugam (IBM Research, T. J. Watson Research Center) · Kush Varshney (IBM Research AI) · Amit Dhurandhar (IBM Research)

Video Prediction via Example Guidance
Jingwei Xu (Shanghai Jiao Tong University) · Harry (Huazhe) Xu (UC Berkeley) · Bingbing Ni (Shanghai Jiao Tong University) · Xiaokang Yang (Shanghai Jiao Tong University of China) · Trevor Darrell (University of California at Berkeley)

Learning Discrete Structured Representations by Adversarially Maximizing Mutual Information
Karl Stratos (Toyota Technological Institute at Chicago) · Sam Wiseman (TTIC)

Reinforcement Learning in Feature Space: Matrix Bandit, Kernels, and Regret Bound
Lin Yang (UCLA) · Mengdi Wang (Princeton University)

Frequency Bias in Neural Networks for Input of Non-Uniform Density
Ronen Basri (Weizmann Institute of Science) · Meirav Galun (Weizmann Institute of Science) · Amnon Geifman (Weizmann Institute) · David Jacobs (University of Maryland, USA) · Yoni Kasten (Weizmann Institute) · Shira Kritchman (Weizmann Institute)

Constrained Markov Decision Processes via Backward Value Functions
harsh satija (McGill University/FAIR) · Philip Amortila (McGill University) · Joelle Pineau (McGill University / Facebook)

Adding seemingly uninformative labels helps in low data regimes
Christos Matsoukas (KTH Royal Institute of Technology) · Albert Bou Hernandez (Universitat Pompeu Fabra) · Yue Liu (KTH, Royal Institute of Technology) · Karin Dembrower (Karolinska Institute) · Gisele Miranda (Science for Life Laboratory) · Emir Konuk (KTH) · Johan Fredin Haslum (KTH Royal Institute of Technology) · Athanasios Zouzos (Karolinska Institute) · Peter Lindholm (Karolinska Institute) · Fredrik Strand (Karolinska Institutet) · Kevin Smith (KTH Royal Institute of Technology)

When are Non-Parametric Methods Robust?
Robi Bhattacharjee (UCSD) · Kamalika Chaudhuri (University of California at San Diego)

Learning Calibratable Policies using Programmatic Style-Consistency
Eric Zhan (California Institute of Technology) · Albert Tseng (Caltech) · Yisong Yue (Caltech) · Adith Swaminathan (Microsoft Research) · Matthew Hausknecht (Microsoft Research)

Momentum Improves Normalized SGD
Ashok Cutkosky (Google) · Harsh Mehta (Google Research)

Parameter-free, Dynamic, and Strongly-Adaptive Online Learning
Ashok Cutkosky (Google)

PENNI: Pruned Kernel Sharing for Efficient CNN Inference
Shiyu Li (Duke University) · Edward Hanson (Duke University) · Hai Li (Duke University) · Yiran Chen (Duke University)

Optimal transport mapping via input convex neural networks
Ashok Vardhan Makkuva (UIUC) · Amirhossein Taghvaei (University of Illinois at Urbana-Champaign) · Sewoong Oh (University of Washington) · Jason Lee (Princeton)

All in the (Exponential) Family: Information Geometry and Thermodynamic Variational Inference
Robert Brekelmans (University of Southern California) · Vaden W Masrani (University of British Columbia) · Frank Wood (University of British Columbia) · Greg Ver Steeg (University of Southern California) · Aram Galstyan (USC ISI)

SimGANs: Simulator-Based Generative Adversarial Networks for ECG Synthesis to Improve Deep ECG Classification
Tomer Golany (Technion - Israel Institute of Technology) · Kira Radinsky (Technion-Israel institute of technology) · Daniel Freedman (Google Israel)

Is There a Trade-Off Between Fairness and Accuracy? A Perspective Using Mismatched Hypothesis Testing
Sanghamitra Dutta (Carnegie Mellon University) · Dennis Wei (IBM Research) · Hazar Yueksel (IBM Research) · Pin-Yu Chen (IBM Research AI) · Sijia Liu (MIT-IBM Watson AI Lab) · Kush Varshney (IBM Research AI)

Convex Calibrated Surrogates for the Multi-Label F-Measure
Mingyuan Zhang (University of Pennsylvania) · Harish Guruprasad Ramaswamy (IIT Madras) · Shivani Agarwal (University of Pennsylvania)

Learning Robot Skills with Temporal Variational Inference
Tanmay Shankar (Facebook AI Research) · Abhinav Gupta (Carnegie Mellon University)

Adaptive Gradient Descent without Descent
Konstantin Mishchenko (King Abdullah University of Science & Technology (KAUST)) · Yura Malitsky (EPFL)

An end-to-end Differentially Private Latent Dirichlet Allocation Using a Spectral Algorithm
Christopher R DeCarolis (University of Maryland) · Mukul A Ram (University of Maryland) · Seyed Esmaeili (University of Maryland, College Park) · Yu-Xiang Wang (UC Santa Barbara) · Furong Huang (University of Maryland College Park)

Dual Mirror Descent for Online Allocation Problems
Haihao Lu (MIT) · Santiago Balseiro (Columbia University) · Vahab Mirrokni (Google Research)

Optimal Robust Learning of Discrete Distributions from Batches
Ayush Jain (UC San Diego) · Alon Orlitsky (UCSD)

BoXHED: Boosted eXact Hazard Estimator with Dynamic covariates
Xiaochen Wang (Facebook) · Arash Pakbin (Texas A&M University) · Bobak Mortazavi (Texas A&M University) · Hongyu Zhao (Yale University) · Donald Lee (Emory University)

Unlabelled Data Improves Bayesian Uncertainty Calibration under Covariate Shift
Alexander Chan (University of Cambridge) · Ahmed Alaa (UCLA) · Zhaozhi Qian (University of Cambridge) · M van der Schaar (UCLA)

Universal Equivariant Multilayer Perceptrons
Siamak Ravanbakhsh (UBC)

Improving generalization by controlling label-noise information in neural network weights
Hrayr Harutyunyan (University of Southern California) · Kyle Reing (USC Information Sciences Institute) · Greg Ver Steeg (University of Southern California) · Aram Galstyan (USC ISI)

DeepMatch: Balancing Deep Covariate Representations for Causal Inference Using Adversarial Training
Nathan Kallus (Cornell University)

Bayesian Optimisation over Multiple Continuous and Categorical Inputs
Binxin Ru (University of Oxford) · Ahsan Alvi (University of Oxford) · Vu Nguyen (University of Oxford) · Michael A Osborne (U Oxford) · Stephen Roberts (University of Oxford)

Generalization and Representational Limits of Graph Neural Networks
Vikas K Garg (Massachusetts Institute of Technology) · Stefanie Jegelka (Massachusetts Institute of Technology) · Tommi Jaakkola (MIT)

Multi-Precision Policy Enforced Training (MuPPET) : A Precision-Switching Strategy for Quantised Fixed-Point Training of CNNs
Aditya Rajagopal (Imperial College London) · Diederik Vink (Imperial College London) · Stylianos Venieris (Samsung AI) · Christos-Savvas Bouganis (Imperial College London)

LowFER: Low-rank Bilinear Pooling for Link Prediction
Saadullah Amin (German Research Center for Aritificial Intelligence (DFKI)) · Stalin Varanasi (German Research Center for Artificial Intelligence (DFKI)) · Katherine Ann Dunfield (German Research Center for Artificial Intelligence (DFKI)) · Günter Neumann (German Research Center for Artificial Intelligence (DFKI))

Parameterized Rate-Distortion Stochastic Encoder
Quan Hoang (Monash University) · Trung Le (Monash University) · Dinh Phung (Monash University, Australia)

Incidence Networks for Geometric Deep Learning
Marjan Albooyeh (UBC) · Daniele Bertolini (N/A) · Siamak Ravanbakhsh (UBC)

Energy-Based Processes for Exchangeable Data
Sherry Yang (Google) · Bo Dai (Google Brain) · Hanjun Dai (Google Brain) · Dale Schuurmans (Google / University of Alberta)

Deep Isometric Learning for Visual Recognition
Haozhi Qi (UC Berkeley) · Chong You (University of California, Berkeley) · Xiaolong Wang (UCSD/UC Berkeley) · Yi Ma (UC Berkeley) · Jitendra Malik (University of California at Berkeley)

Second-Order Provable Defenses against Adversarial Attacks
Sahil Singla (University of Maryland) · Soheil Feizi (University of Maryland)

Transformers are RNNs: Fast Autoregressive Transformers with Linear Attention
Angelos Katharopoulos (Idiap & EPFL) · Apoorv Vyas (Idiap Research Institute) · Nikolaos Pappas (University of Washington) · Francois Fleuret (Idiap research institute)

Overfitting in adversarially robust deep learning
Eric Wong (Carnegie Mellon University) · Leslie Rice (Carnegie Mellon University) · Zico Kolter (Carnegie Mellon University / Bosch Center for AI)

Rethinking Bias-Variance Trade-off for Generalization of Neural Networks
Zitong Yang (University of California, Berkeley) · Yaodong Yu (University of California, Berkeley) · Chong You (University of California, Berkeley) · Jacob Steinhardt (UC Berkeley) · Yi Ma (UC Berkeley)

Boosting for Control of Dynamical Systems
Naman Agarwal (Google AI Princeton) · Nataly Brukhim (Princeton University) · Elad Hazan (Princeton University) · Zhou Lu (Princeton University)

Frustratingly Simple Few-Shot Object Detection
Xin Wang (UC Berkeley) · Thomas Huang (University of Michigan) · Joseph Gonzalez (UC Berkeley) · Trevor Darrell (University of California at Berkeley) · Fisher Yu (UC Berkeley)

Data-Dependent Differentially Private Parameter Learning for Directed Graphical Models
Amrita Roy Chowdhury (University of Wisconsin-Madison) · Theodoros Rekatsinas (University of Wisconsin-Madison) · Somesh Jha (University of Wisconsin, Madison)

Adversarial Risk via Optimal Transport and Optimal Couplings
Muni Sreenivas Pydi (University of Wisconsin - Madison) · Varun Jog (University of Wisconsin - Madison)

Decoupled Greedy Learning of CNNs
Eugene Belilovsky (Mila) · Michael Eickenberg (UC Berkeley) · Edouard Oyallon (CNRS/LIP6)

ACFlow: Flow Models for Arbitrary Conditional Likelihoods
Yang Li (Department of Computer Science, University of North Carolina at Chapel Hill) · Shoaib Akbar (North Carolina state university) · Junier Oliva (UNC-Chapel Hill)

Can autonomous vehicles identify, recover from, and adapt to distribution shifts?
Angelos Filos (University of Oxford) · Panagiotis Tigkas (Oxford University) · Rowan McAllister (UC Berkeley) · Nicholas Rhinehart (Carnegie Mellon University) · Sergey Levine (UC Berkeley) · Yarin Gal (University of Oxford)

Leveraging Procedural Generation to Benchmark Reinforcement Learning
Karl Cobbe (OpenAI) · Chris Hesse (OpenAI) · Jacob Hilton (OpenAI) · John Schulman (OpenAI)

The Tree Ensemble Layer: Differentiability meets Conditional Computation
Hussein Hazimeh (Massachusetts Institute of Technology) · Natalia Ponomareva (Google) · Rahul Mazumder (Massachusetts Institute of Technology) · Zhenyu Tan (Google) · Petros Mol (Google Research)

Near-Tight Margin-Based Generalization Bounds for Support Vector Machines
Allan Grønlund (Aarhus University, MADALGO) · Lior Kamma (Aarhus University) · Kasper Green Larsen (Aarhus University, MADALGO)

Error Estimation for Sketched SVD
Miles Lopes (University of California, Davis) · N. Benjamin Erichson (University of California, Berkeley) · Michael Mahoney (UC Berkeley)

Goal-Aware Prediction: Learning to Model What Matters
Suraj Nair (Stanford University) · Silvio Savarese (Stanford University) · Chelsea Finn (Stanford)

Combinatorial Pure Exploration for Dueling Bandit
Wei Chen (Microsoft) · Yihan Du (IIIS, Tsinghua University) · Longbo Huang (Tsinghua University) · Haoyu Zhao (Tsinghua University)

Optimal Sequential Maximization: One Interview is Enough!
Moein Falahatgar (UCSD) · Alon Orlitsky (UCSD) · Venkatadheeraj Pichapati (University of California San Diego)

What can I do here? A Theory of Affordances in Reinforcement Learning
Khimya Khetarpal (McGill University, Mila Montreal) · Zafarali Ahmed (DeepMind) · Gheorghe Comanici (DeepMind) · David Abel (Brown University) · Doina Precup (DeepMind)

An end-to-end approach for the verification problem: learning the right distance
Joao Monteiro (Institut National de la Recherche Scientifique (INRS)) · Isabela Albuquerque (Institut National de la Recherche Scientifique) · Jahangir Alam (Computer Research Institute of Montreal (CRIM), Montreal (Quebec) Canada) · R Devon Hjelm (Microsoft Research) · Tiago Falk (INRS-EMT)

Data Valuation using Reinforcement Learning
Jinsung Yoon (University of California, Los Angeles) · Sercan O. Arik (Google) · Tomas Pfister (Google)

FormulaZero: Distributionally Robust Online Adaptation via Offline Population Synthesis
Aman Sinha (Stanford University) · Matthew O'Kelly (University of Pennsylvania) · Hongrui Zheng (University of Pennsylvania) · Rahul Mangharam (University of Pennsylvania) · John Duchi (Stanford University) · Russ Tedrake (MIT)

Latent Bernoulli Autoencoder
Jiri Fajtl (Kingston University London) · Vasileios Argyriou (Kingston University) · Dorothy Monekosso (Leeds Beckett University) · Paolo Remagnino (Kingston University)

Learning to Stop While Learning to Predict
Xinshi Chen (Georgia Institution of Technology) · Hanjun Dai (Google Brain) · Yu Li (King Abdullah University of Science and Technology) · Xin Gao (Kaust) · Le Song (Georgia Institute of Technology)

Accelerating the diffusion-based ensemble sampling by non-reversible dynamics
Futoshi Futami (University of Tokyo) · Issei Sato (University of Tokyo / RIKEN) · Masashi Sugiyama (RIKEN / The University of Tokyo)

A unified approach for assessing population feature importance using Shapley values
Brian Williamson (Fred Hutchinson Cancer Research Center) · Jean Feng (University of Washington)

Curse of Dimensionality on Randomized Smoothing for Certifiable Robustness
Aounon Kumar (University of Maryland, College Park) · Alexander Levine (University of Maryland) · Tom Goldstein (University of Maryland) · Soheil Feizi (University of Maryland)

Upper bounds for Model-Free Row-Sparse Principal Component Analysis
Guanyi Wang (Georgia Institute of Technology) · Santanu Dey (Georgia Institute of Technology)

Explainable k-Means and k-Medians Clustering
Michal Moshkovitz (University of California San Diego) · Sanjoy Dasgupta (UC San Diego) · Cyrus Rashtchian (UCSD) · Nave Frost (Tel-Aviv University)

Reward-Free Exploration for Reinforcement Learning
Chi Jin (Princeton University) · Akshay Krishnamurthy (Microsoft Research) · Max Simchowitz (UC Berkeley) · Tiancheng Yu (MIT)

Parametric Gaussian Process Regressors
Martin Jankowiak (Uber AI Labs) · Geoff Pleiss (Cornell University) · Jacob Gardner (Uber AI Labs)

p-Norm Flow Diffusion for Local Graph Clustering
Kimon Fountoulakis (University of Waterloo) · Di Wang (Google Research) · Shenghao Yang (University of Waterloo)

Low-Rank Bottleneck in Multi-head Attention Models
Srinadh Bhojanapalli (Google AI) · Chulhee Yun (MIT) · Ankit Singh Rawat (Google) · Sashank Jakkam Reddi (Google) · Sanjiv Kumar (Google Research, NY)

LEEP: A New Measure to Evaluate Transferability of Learned Representations
Cuong V Nguyen (Amazon) · Tal Hassner (Open University of Israel) · Cedric Archambeau (Amazon) · Matthias W Seeger (Amazon)

The FAST Algorithm for Submodular Maximization
Adam Breuer (Harvard University) · Eric Balkanski (Harvard) · Yaron Singer (Harvard)

On the Relation between Quality-Diversity Evaluation and Distribution-Fitting Goal in Text Generation
Jianing Li (Institute of Computing Technology, Chinese Academy of Sciences) · Yanyan Lan (Institute of Computing Technology) · Jiafeng Guo (Institute of Computing Technology, Chinese Academy of Sciences) · Xueqi Cheng (Institute of Computing Technology, CAS, China)

Designing Optimal Dynamic Treatment Regimes: A Causal Reinforcement Learning Approach
Junzhe Zhang (Columbia University)

Global Decision-Making via Local Economic Transactions
Michael Chang (UC Berkeley) · Sid Kaushik (UCB) · S. Matthew Weinberg (Princeton University) · Sergey Levine (UC Berkeley) · Thomas Griffiths (Princeton University)

Retrieval Augmented Language Model Pre-Training
Kelvin Guu (Google) · Kenton Lee (Google Research) · Zora Tung (Google) · Panupong Pasupat (Google) · Mingwei Chang (Google)

Variational Label Enhancement
Ning Xu (Southeast University) · Yun-Peng Liu (Southeast University) · Jun Shu (Xi'an Jiaotong University) · Xin Geng (Southeast University)

Bandits with Adversarial Scaling
Thodoris Lykouris (Microsoft Research) · Vahab Mirrokni (Google Research) · Renato Leme (Google Research)

Eliminating the Invariance on the Loss Landscape of Linear Autoencoders
Reza Oftadeh (Texas A&M university) · Jiayi Shen (Texas A&M) · Zhangyang Wang (Texas A&M University) · Dylan Shell (Texas A&M University)

What is Local Optimality in Nonconvex-Nonconcave Minimax Optimization?
Chi Jin (Princeton University) · Praneeth Netrapalli (Microsoft Research) · Michael Jordan (UC Berkeley)

Lookahead-Bounded Q-learning
Ibrahim El Shar (University of Pittsburgh) · Daniel Jiang (University of Pittsburgh)

Learning From Irregularly-Sampled Time Series: A Missing Data Perspective
Steven Cheng-Xian Li (UMass Amherst) · Benjamin M Marlin (University of Massachusetts Amherst)

Evaluating the Performance of Reinforcement Learning Algorithms
Scott Jordan (University of Massachusetts Amherst) · Yash Chandak (University of Massachusetts Amherst) · Daniel Cohen (University of Massachusetts Amherst) · Mengxue Zhang (umass Amherst) · Philip Thomas (University of Massachusetts Amherst)

Unbiased Risk Estimators Can Mislead: A Case Study of Learning with Complementary Labels
Yu-Ting Chou (National Taiwan University) · Gang Niu (RIKEN) · Hsuan-Tien Lin (National Taiwan University) · Masashi Sugiyama (RIKEN / The University of Tokyo)

Provable Self-Play Algorithms for Competitive Reinforcement Learning
Yu Bai (Salesforce Research) · Chi Jin (Princeton University)

Optimizing Long-term Social Welfare in Recommender Systems: A Constrained Matching Approach
Martin Mladenov (Google) · Elliot Creager (University of Toronto) · Omer Ben-Porat (Technion--Israel Institute of Technology) · Kevin Swersky (Google Brain) · Richard Zemel (Vector Institute) · Craig Boutilier (Google)

Semi-Supervised StyleGAN for Disentanglement Learning
Weili Nie (Rice University) · Tero Karras (NVIDIA) · Animesh Garg (University of Toronto, Vector Institute, Nvidia) · Shoubhik Debnath (Nvidia) · Anjul Patney (Nvidia) · Ankit Patel (Rice University, Baylor College of Medicine) · Anima Anandkumar (Amazon AI & Caltech)

The Non-IID Data Quagmire of Decentralized Machine Learning
Kevin Hsieh (Microsoft Research) · Amar Phanishayee (Microsoft Research) · Onur Mutlu (ETH Zurich) · Phillip Gibbons (CMU)

On the Noisy Gradient Descent that Generalizes as SGD
Jingfeng Wu (Johns Hopkins University) · Wenqing Hu (Missouri S&T) · Haoyi Xiong (Baidu Research) · Jun Huan (Styling AI) · Vladimir Braverman (Johns Hopkins University) · Zhanxing Zhu (Peking University)

Safe screening rules for L0-regression
Alper Atamturk (UC Berkeley) · Andres Gomez (usc)

Single Point Transductive Prediction
Nilesh Tripuraneni (UC Berkeley) · Lester Mackey (Microsoft Research)

History-Gradient Aided Batch Size Adaptation for Variance Reduced Algorithms
Kaiyi Ji (The Ohio State University) · Zhe Wang (Ohio State University) · Bowen Weng (Ohio State University) · Yi Zhou (University of Utah) · Wei Zhang (Southern University of Science and Technology) · Yingbin LIANG (The Ohio State University)

Batch Stationary Distribution Estimation
Junfeng Wen (University of Alberta) · Bo Dai (Google Brain) · Lihong Li (Google Research) · Dale Schuurmans (University of Alberta)

Optimal Statistical Guaratees for Adversarially Robust Gaussian Classification
Chen Dan (Carnegie Mellon University) · Yuting Wei (CMU) · Pradeep Ravikumar (Carnegie Mellon University)

Generative Adversarial Imitation Learning with Neural Network Parameterization: Global Optimality and Convergence Rate
Yufeng Zhang (Northwestern University) · Qi Cai (Northwestern University) · Zhuoran Yang (Princeton University) · Zhaoran Wang (Northwestern U)

A Game Theoretic Perspective on Model-Based Reinforcement Learning
Aravind Rajeswaran (University of Washington) · Igor Mordatch (OpenAI) · Vikash Kumar (Google)

(Locally) Differentially Private Combinatorial Semi-Bandits
Xiaoyu Chen (Peking University) · Kai Zheng (Peking University) · Zixin Zhou (Peking University) · Yunchang Yang (Peking University) · Wei Chen (Microsoft) · Liwei Wang (Peking University)

Optimizing for the Future in Non-Stationary MDPs
Yash Chandak (University of Massachusetts Amherst) · Georgios Theocharous (Adobe Research) · Shiv Shankar (University of Massachusetts) · Martha White (University of Alberta) · Sridhar Mahadevan (Adobe Research) · Philip Thomas (University of Massachusetts Amherst)

Learning Task-Agnostic Embedding of Multiple Black-Box Experts for Multi-Task Model Fusion
Nghia Hoang (MIT-IBM Watson AI Lab, IBM Research) · Thanh Lam (National University of Singapore) · Bryan Kian Hsiang Low (National University of Singapore) · Patrick Jaillet (MIT)

Dual-Path Distillation: A Unified Framework to Improve Black-Box Attacks
Yonggang Zhang (USTC) · Ya Li (IFLYTEK Research) · Tongliang Liu (The University of Sydney) · Xinmei Tian (USTC)

Safe Deep Semi-Supervised Learning for Unseen-Class Unlabeled Data
Lan-Zhe Guo (Nanjing University) · Zhen-Yu Zhang (Nanjing University) · Yuan Jiang (Nanjing University) · Yufeng Li (Nanjing University) · Zhi-Hua Zhou (Nanjing University)

Generalizing Convolutional Neural Networks for Equivariance to Lie Groups on Arbitrary Continuous Data
Marc Finzi (New York University) · Samuel Stanton (New York University) · Pavel Izmailov (New York University) · Andrew Wilson (New York University)

Dispersed EM-VAEs for Interpretable Text Generation
Wenxian Shi (Bytedance) · Hao Zhou (Bytedance) · Ning Miao (ByteDance AI Lab) · Lei Li (ByteDance AI Lab)

Deep Graph Random Process for Relational-Thinking-Based Speech Recognition
Huang Hengguan (NUS) · Fuzhao Xue (National University of Singapore) · Hao Wang (AWS AI Labs) · Ye Wang (National University of Singapore)

Hypernetwork approach to generating point clouds
Przemysław Spurek (Jagiellonian University) · Sebastian Winczowski (Jagiellonian University) · Jacek Tabor (Jagiellonian University) · Maciej Zamorski (Wrocław University of Science and Technology) · Maciej Zieba (Wroclaw University of Science and Technology) · Tomasz Trzcinski ()

On a projective ensemble approach to two sample test for equality of distributions
Zhimei Li (Shanghai University of Finance and Economics) · Yaowu Zhang (Shanghai University of Finance and Economics)

Coresets for Data-efficient Training of Machine Learning Models
Baharan Mirzasoleiman (Stanford University) · Jeff Bilmes (UW) · Jure Leskovec (Stanford University)

Searching to Exploit Memorization Effect in Learning from Noisy Labels
QUANMING YAO (4Paradigm) · Hansi Yang (Tsinghua) · Bo Han (HKBU / RIKEN) · Gang Niu (RIKEN) · James Kwok (Hong Kong University of Science and Technology)

Randomized Smoothing of All Shapes and Sizes
Greg Yang (Microsoft Research) · Tony Duan (Microsoft Research AI) · J. Edward Hu (Microsoft Research AI) · Hadi Salman (Microsoft Research AI) · Ilya Razenshteyn (Microsoft Research, Redmond) · Jerry Li (Microsoft)

DeepCoDA: personalized interpretability for compositional health
Thomas Quinn (A2I2) · Dang Nguyen (Deakin University) · Santu Rana (Deakin University) · Sunil Gupta (Deakin University) · Svetha Venkatesh (Deakin University)

Private Query Release Assisted by Public Data
Raef Bassily (The Ohio State University) · Albert Cheu () · Shay Moran (IAS, Princeton) · Aleksandar Nikolov () · Jonathan Ullman (Northeastern University) · Steven Wu (University of Minnesota)

Adaptive Droplet Routing in Digital Microfluidic Biochips Using Deep Reinforcement Learning
Tung-Che Liang (Duke University) · Zhanwei Zhong (Duke University) · Yaas Bigdeli (Duke Univsersity) · Tsung-Yi Ho (National Tsing Hua University) · Richard Fair (Duke University) · Krishnendu Chakrabarty (Duke University)

Continuous-time Lower Bounds for Gradient-based Algorithms
Michael Muehlebach (UC Berkeley) · Michael Jordan (UC Berkeley)

A Tree-Structured Decoder for Image-to-Markup Generation
Jianshu Zhang (University of Science and Technology of China) · Jun Du (University of Science and Technology of China) · Yongxin Yang (University of Surrey) · Yi-Zhe Song (University of Surrey) · Si Wei (iFLYTEK) · Lirong Dai (N/A)

Sample Factory: Egocentric 3D Control from Pixels at 100000 FPS with Asynchronous Reinforcement Learning
Aleksei Petrenko (University of Southern California) · Zhehui Huang (University of Southern California) · Tushar Kumar (University of Southern California) · Gaurav Sukhatme (University of Southern California) · Vladlen Koltun (Intel Labs)

Scalable Deep Generative Modeling for Sparse Graphs
Hanjun Dai (Google Brain) · Azade Nazi (Google Brain) · Yujia Li (DeepMind) · Bo Dai (Google Brain) · Dale Schuurmans (Google / University of Alberta)

Closed Loop Neural-Symbolic Learning via Integrating Neural Perception, Grammar Parsing, and Symbolic Reasoning
Qing Li (UCLA) · Siyuan Huang (UCLA) · Yining Hong (University of California, Los Angeles) · Yixin Chen (UCLA) · Ying Nian Wu (UCLA) · Song-Chun Zhu (UCLA)

NGBoost: Natural Gradient Boosting for Probabilistic Prediction
Tony Duan (Microsoft Research AI) · Anand Avati (Stanford University) · Daisy Ding (Stanford University) · Khanh K. Thai (Kaiser Permanente Division of Research) · Sanjay Basu (Harvard Medical School) · Andrew Ng (Stanford U.) · Alejandro Schuler (Stanford University)

Q-value Path Decomposition for Deep Multiagent Reinforcement Learning
Yaodong Yang (Tianjin University) · Jianye Hao (Tianjin University) · Guangyong Chen (Tencent) · Hongyao Tang (Tianjin University) · Yingfeng Chen (NetEase Fuxi AI Lab) · Yujing Hu (NetEase Fuxi AI Lab) · Changjie Fan (Netease) · Zhongyu Wei (Fudan University)

Online Learned Continual Compression with Adaptive Quantization Modules
Lucas Caccia (McGIll) · Eugene Belilovsky (Mila) · Massimo Caccia (MILA) · Joelle Pineau (McGill University / Facebook)

Learning What to Defer for Maximum Independent Sets
Sung-Soo Ahn (KAIST) · Younggyo Seo (KAIST) · Jinwoo Shin (KAIST)

Generalized and Scalable Optimal Sparse Decision Trees
Jimmy Lin (University of British Columbia) · Chudi Zhong (Duke University) · Diane Hu (Duke University) · Cynthia Rudin (Duke) · Margo Seltzer (University of British Columbia)

The Effect of Natural Distribution Shift on Question Answering Models
John Miller (University of California, Berkeley) · Karl Krauth (UC Berkeley) · Ludwig Schmidt (University of California, Berkeley) · Benjamin Recht (Berkeley)

Quantized Decentralized Stochastic Learning over Directed Graphs
Hossein Taheri (UC Santa Barbara) · Aryan Mokhtari (UT Austin) · Hamed Hassani (University of Pennsylvania) · Ramtin Pedarsani (University of California, Santa Barbara)

Semi-Supervised Learning with Normalizing Flows
Pavel Izmailov (New York University) · Polina Kirichenko (New York Univeristy) · Marc Finzi (New York University) · Andrew Wilson (New York University)

Student Specialization in Deep Rectified Networks With Finite Width and Input Dimension
Yuandong Tian (Facebook AI Research)

Sample Amplification: Increasing Dataset Size even when Learning is Impossible
Brian Axelrod (Stanford) · Shivam Garg (Stanford University) · Vatsal Sharan (Stanford University) · Gregory Valiant (Stanford University)

Alleviating Privacy Attacks via Causal Learning
Shruti Tople (Microsoft Research) · Amit Sharma (Microsoft Research) · Aditya Nori (Microsoft Research Cambridge)

The Intrinsic Robustness of Stochastic Bandits to Strategic Manipulation
Zhe Feng (Harvard University) · David Parkes (Harvard University) · Haifeng Xu (University of Virginia)

Normalized Flat Minima: Exploring Scale Invariant Definition of Flat Minima for Neural Networks Using PAC-Bayesian Analysis
Yusuke Tsuzuku (The University of Tokyo / RIKEN) · Issei Sato (University of Tokyo / RIKEN) · Masashi Sugiyama (RIKEN / The University of Tokyo)

Fiedler Regularization: Learning Neural Networks with Graph Sparsity
Edric Tam (Duke University) · David Dunson (Duke University)

Online Learning with Imperfect Hints
Aditya Bhaskara (University of Utah) · Ashok Cutkosky (Google) · Ravi Kumar (Google) · Manish Purohit (Google Research)

Rate-distortion optimization guided autoencoder for isometric embedding in Euclidean latent space
Keizo Kato (Fujitsu Laboratories Ltd.) · Jing Zhou (Fujitsu R&D Center Co., LTD. Shanghai Laboratory) · Tomotake Sasaki (Fujitsu Laboratories Ltd.) · Akira Nakagawa (Fujitsu Laboratories Ltd.)

Optimization from Structured Samples for Coverage Functions
Wei Chen (Microsoft) · Xiaoming Sun (Institute of Computing Technology, Chinese Academy of Sciences) · Jialin Zhang (Institute of Computing Technology, CAS) · Zhijie Zhang (Institute of Computing Technology, Chinese Academy of Sciences)

Optimal Randomized First-Order Methods for Least-Squares Problems
Jonathan Lacotte (Stanford University) · Mert Pilanci (Stanford)

Stochastic Optimization for Non-convex Inf-Projection Problems
Yan Yan (the University of Iowa) · Yi Xu (Alibaba Group (U.S.) Inc.) · Lijun Zhang (Nanjing University) · Wang Xiaoyu (Intellifusion) · Tianbao Yang (The University of Iowa)

Convex Representation Learning for Generalized Invariance in Semi-Inner-Product Space
Yingyi Ma (UIC) · Vignesh Ganapathiraman (University of Illinois at Chicago) · Yaoliang Yu (University of Waterloo) · Xinhua Zhang (University of Illinois at Chicago (UIC))

Neural Kernels Without Tangents
Vaishaal Shankar (UC Berkeley) · Alex Fang (UC Berkeley) · Wenshuo Guo (UC Berkeley) · Sara Fridovich-Keil (UC Berkeley) · Jonathan Ragan-Kelley (UC Berkeley) · Ludwig Schmidt (University of California, Berkeley) · Benjamin Recht (Berkeley)

Linear Lower Bounds and Conditioning of Differentiable Games
Adam Ibrahim (Mila, Université de Montréal) · Waïss Azizian (Ecole Normale Supérieure de Paris) · Gauthier Gidel (MILA) · Ioannis Mitliagkas (MILA, UdeM)

Finite-Time Last-Iterate Convergence for Multi-Agent Learning in Games
Tianyi Lin (UC Berkeley) · Zhengyuan Zhou (Stanford University) · Panayotis Mertikopoulos (CNRS) · Michael Jordan (UC Berkeley)

Communication-Efficient Distributed PCA by Riemannian Optimization
Long-Kai Huang (NTU, Singapore) · Jialin Pan (NTU, Singapore)

Manifold Identification for Ultimately Communication-Efficient Distributed Optimization
Yu-Sheng Li (National Taiwan University) · Wei-Lin Chiang (National Taiwan University) · Ching-pei Lee (University of Wisconsin-Madison)

When Demands Evolve Larger and Noisier: Learning and Earning in a Growing Environment
Feng Zhu (Peking University) · Zeyu Zheng (UC Berkeley)

Being Bayesian about Categorical Probability
Taejong Joo (ESTsoft) · Uijung Chung (ESTsoft) · Min-Gwan Seo (ESTsoft)

Context-aware Dynamics Model for Generalization in Model-Based Reinforcement Learning
Kimin Lee (UC Berkeley) · Younggyo Seo (KAIST) · Seunghyun Lee (KAIST) · Honglak Lee (Google / U. Michigan) · Jinwoo Shin (KAIST)

Learning Reasoning Strategies in End-to-End Differentiable Proving
Pasquale Minervini (University College London) · Tim Rocktäschel (Facebook, UCL) · Sebastian Riedel (UCL) · Edward Grefenstette (Facebook AI Research / UCL) · Pontus Stenetorp (University College London)

Fast and Private Submodular and k -Submodular Functions Maximization with Matroid Constraints
Akbar Rafiey (Simon Fraser University) · Yuichi Yoshida (National Institute of Informatics)

Streaming Coresets for Symmetric Tensor Factorization
Supratim Shit (IIT Gandhinagar) · Anirban Dasgupta (IIT Gandhinagar) · Rachit Chhaya (IIT Gandhinagar) · Jayesh Choudhari (IIT Gandhinagar)

How Good is the Bayes Posterior in Deep Neural Networks Really?
Florian Wenzel (Google Research) · Kevin Roth (ETH Zurich) · Bastiaan Veeling (University of Amsterdam) · Jakub Swiatkowski (University of Warsaw) · Linh Tran (Imperial College London) · Stephan Mandt (UC Irvine) · Jasper Snoek (Google Brain) · Tim Salimans (Google) · Rodolphe Jenatton (Google) · Sebastian Nowozin (MSR Cambridge)

Optimally Solving Two-Agent Decentralized POMDPs Under One-Sided Information Sharing
Yuxuan Xie (INSA de Lyon) · Jilles Dibangoye (INSA Lyon, INRIA) · Olivier Buffet (INRIA - LORIA)

Learning Algebraic Multigrid Using Graph Neural Networks
Ilay Luz (Weizmann Institute of Science) · Meirav Galun (Weizmann Institute of Science) · Haggai Maron (Weizmann Institute of Science) · Ronen Basri (Weizmann Institute of Science) · Irad Yavneh (Technion)

Fractal Gaussian Networks: A sparse random graph model based on Gaussian Multiplicative Chaos
Subhroshekhar Ghosh (National University of Singapore) · Krishna Balasubramanian (University of California, Davis) · Xiaochuan Yang (Université du Luxembourg)

Structured Policy Iteration for Linear Quadratic Regulator
Youngsuk Park (Stanford University) · Ryan Rossi (Adobe Research) · Zheng Wen (DeepMind) · Gang Wu (Adobe Research) · Handong Zhao (Adobe Research)

T-GD: Transferable GAN-generated Images Detection Framework
Hyeonseong Jeon (Sungkyunkwan University) · Young Oh Bang (Sungkyunkwan University) · Junyaup Kim (Sungkyunkwan University) · Simon Woo (SKKU)

Low Bias Low Variance Gradient Estimates for Hierarchical Boolean Stochastic Networks
Adeel Pervez (University of Amsterdam) · Taco Cohen (Qualcomm) · Efstratios Gavves (University of Amsterdam)

Learning Flat Latent Manifolds with VAEs
Nutan Chen (Volkswagen Group) · Alexej Klushyn (ML Research, VW Group) · Francesco Ferroni (Autonomous Intelligent Driving GmbH) · Justin Bayer (Volkswagen Group) · Patrick van der Smagt (Volkswagen Group)

Multi-Task Learning with User Preferences: Gradient Descent with Controlled Ascent in Pareto Optimization
Debabrata Mahapatra (National University of Singapore) · Vaibhav Rajan (National University of Singapore)

Transfer Learning without Knowing: Reprogramming Black-box Machine Learning Models with Scarce Data and Limited Resources
Yun Yun Tsai (National Tsing Hua University) · Pin-Yu Chen (IBM Research AI) · Tsung-Yi Ho (National Tsing Hua University)

On Coresets for Regularized Regression
Rachit Chhaya (IIT Gandhinagar) · Supratim Shit (IIT Gandhinagar) · Anirban Dasgupta (IIT Gandhinagar)

Budgeted Online Influence Maximization
Pierre Perrault (Inria Lille - Nord Europe) · Zheng Wen (DeepMind) · Michal Valko (DeepMind) · Jennifer Healey (Adobe Research)

On the (In)tractability of Computing Normalizing Constants for the Product of Determinantal Point Processes
Naoto Ohsaka (NEC Corporation) · Tatsuya Matsuoka (NEC Corporation)

Monte-Carlo Tree Search as Regularized Policy Optimization
Jean-Bastien Grill (DeepMind) · Florent Altché (DeepMind) · Yunhao Tang (Columbia University) · Thomas Hubert (DeepMind) · Michal Valko (DeepMind) · Ioannis Antonoglou (Deepmind) · Remi Munos (DeepMind)

On the Expressivity of Neural Networks for Deep Reinforcement Learning
Kefan Dong (Tsinghua University) · Yuping Luo (Princeton University) · Tianhe Yu (Stanford University) · Chelsea Finn (Stanford) · Tengyu Ma (Stanford)

The k-tied Normal Distribution: A Compact Parameterization of Gaussian Mean Field Posteriors in Bayesian Neural Networks
Jakub Swiatkowski (University of Warsaw) · Kevin Roth (ETH Zurich) · Bastiaan Veeling (University of Amsterdam) · Linh Tran (Imperial College London) · Joshua V Dillon (Google) · Jasper Snoek (Google Brain) · Stephan Mandt (UC Irvine) · Tim Salimans (Google) · Rodolphe Jenatton (Google) · Sebastian Nowozin (Google Research)

A Generative Model for Molecular Distance Geometry
Gregor Simm (Cambridge University) · Jose Hernandez-Lobato (University of Cambridge)

Why bigger is not always better: on finite and infinite neural networks
Laurence Aitchison (University of Cambridge)

Data-Efficient Image Recognition with Contrastive Predictive Coding
Olivier Henaff (DeepMind)

Intrinsic Reward Driven Imitation Learning via Generative Model
Xingrui Yu (University of Technology Sydney) · Yueming LYU (University of Technology Sydney) · Ivor Tsang (University of Technology Sydney)

Can Increasing Input Dimensionality Improve Deep Reinforcement Learning?
Kei Ota (Mitsubishi Electric Corporation) · Tomoaki Oiki (Mitsubishi Electric) · Devesh Jha (Mitsubishi Electric Research Labs) · Toshisada Mariyama (Mitsubishi Electric) · Daniel Nikovski (Mitsubishi Electric Research Labs)

Batch Reinforcement Learning with Hyperparameter Gradients
Byung-Jun Lee (KAIST) · Jongmin Lee (KAIST) · Peter Vrancx (PROWLER.io) · Dongho Kim (Prowler.io) · Kee-Eung Kim (KAIST)

Sub-Goal Trees -- a Framework for Goal-Based Reinforcement Learning
Tom Jurgenson (Technion) · Or Avner (Technion) · Edward Groshev (Osaro, Inc.) · Aviv Tamar (Technion)

A Geometric Approach to Archetypal Analysis via Sparse Projections
Vinayak Abrol (Mathematical Institute Oxford) · Pulkit Sharma (Quantum Black)

Sequence Generation with Mixed Representations
Lijun Wu (Sun Yat-sen University) · Shufang Xie (Microsoft Research Asia) · Yingce Xia (University of Science and Technology of China) · Yang Fan (University of Science and Technology of China) · Jian-Huang Lai (Sun Yat-sen University) · Tao Qin (Microsoft Research Asia) · Tie-Yan Liu (Microsoft Research Asia)

Agent57: Outperforming the Atari Human Benchmark
Adrià Puigdomenech Badia (Deepmind) · Bilal Piot (DeepMind) · Steven Kapturowski (Deepmind) · Pablo Sprechmann (Google DeepMind) · Oleksandr Vitvitskyi (DeepMind) · Zhaohan Guo (DeepMind) · Charles Blundell (DeepMind)

RIFLE: Backpropagation in Depth for Deep Transfer Learning through Re-Initializing the Fully-connected LayEr
Xingjian Li (Baidu Research) · Haoyi Xiong (Baidu Research) · Haozhe An (Baidu Research) · Dejing Dou (Baidu) · Cheng-Zhong Xu (University of Macau)

Fairwashing explanations with off-manifold detergent
Christopher Anders (TU Berlin) · Plamen Plasiliev (TU Berlin) · Ann-Kathrin Dombrowski (TU Berlin) · Klaus-robert Mueller (Technische Universität Berlin) · Pan Kessel (TU Berlin)

Learning disconnected manifolds: a no GAN's land
Ugo Tanielian (Criteo) · Jeremie Mary (Criteo) · Thibaut Issenhuth (Criteo) · Elvis Dohmatob (Criteo AI Lab)

Sets Clustering
Ibrahim Jubran (The University of Haifa) · Murad Tukan (University of Haifa) · Alaa Maalouf (The University of Haifa) · Dan Feldman (The University of Haifa)

Variational Autoencoders with Riemannian Brownian Motion Priors
Dimitrios Kalatzis (DTU) · David Eklund (Technical University of Denmark) · Georgios Arvanitidis (MPI for Intelligent Systems, Tübingen) · Søren Hauberg (Technical University of Denmark)

Non-separable Non-stationary random fields
Kangrui Wang (The Alan Turing Institute) · Oliver A Hamelijnck (The Alan Turing Institute) · Theodoros Damoulas (University of Warwick) · Mark Steel (University of Warwick)

Nonparametric Score Estimators
Yuhao Zhou (Tsinghua University) · Jiaxin Shi (Tsinghua University) · Jun Zhu (Tsinghua University)

A Free-Energy Principle for Representation Learning
Yansong Gao (University of Pennsylvania) · Pratik Chaudhari (University of Pennsylvania)

Scalable Differential Privacy with Certified Robustness in Adversarial Learning
Hai Phan (New Jersey Institute of Technology) · My Thai (University of Florida) · Han Hu (New Jersey Institute of Technology) · Ruoming Jin (Kent State University) · Tong Sun (Adobe Research) · Dejing Dou (" University of Oregon, USA")

Variational Inference for Sequential Data with Future Likelihood Estimates
Geon-Hyeong Kim (KAIST) · Youngsoo Jang (KAIST) · Hongseok Yang () · Kee-Eung Kim (KAIST)

Implicit Learning Dynamics in Stackelberg Games: Equilibria Characterization, Convergence Analysis, and Empirical Study
Tanner Fiez (University of Washington) · Benjamin Chasnov (University of Washington) · Lillian Ratliff (University of Washington)

Let's Agree to Agree: Neural Networks Share Classification Order on Real Datasets
Guy Hacohen (Hebrew University of Jerusalem) · Leshem Choshen (Hebrew University, Jerusalem) · Daphna Weinshall (Hebrew University of Jerusalem, Israel)

Quantile Causal Discovery
Natasa Tagasovska (University of Lausanne) · Thibault Vatter (Columbia University) · Valérie Chavez-Demoulin (University of Lausanne)

How to Solve Fair k-Center in Massive Data Models
Ashish Chiplunkar (IIT Delhi) · Sagar Kale (University of Vienna) · Sivaramakrishnan Natarajan Ramamoorthy (University of Washington)

Bayesian Learning from Sequential Data using Gaussian Processes with Signature Covariances
Csaba Toth (University of Oxford) · Harald Oberhauser (University of Oxford)

Beyond Signal Propagation: Is Feature Diversity Necessary in Deep Neural Network Initialization?
Yaniv Blumenfeld (Technion) · Dar Gilboa (Columbia University) · Daniel Soudry (Technion)

Dynamic Knapsack Optimization Towards Efficient Multi-Channel Sequential Advertising
Xiaotian Hao (Tianjin University) · Zhaoqing Peng (Alibaba Group) · Yi Ma (Tianjin University) · Guan Wang (Department of Automation, Tsinghua University) · Junqi Jin (Alibaba Group) · Jianye Hao (Tianjin University) · Shan Chen (Alibaba Group) · Rongquan Bai (Alibaba Group) · Mingzhou Xie (Alibaba Group) · Miao Xu (Alibaba Group) · Zhenzhe Zheng (Shanghai Jiao Tong University) · Chuan Yu (Alibaba Group) · HAN LI (Alibaba Group) · Jian Xu (Alibaba Group) · Kun Gai (Alibaba group)

Stochastically Dominant Distributional Reinforcement Learning
John Martin (Stevens Institute of Technology) · Michal Lyskawinski (Stevens Institute of Technology) · Xiaohu Li (Stevens Institute of Technology) · Brendan Englot (Stevens Institute of Technology)

Adversarial Robustness Against the Union of Multiple Threat Models
Pratyush Maini (IIT Delhi) · Eric Wong (Carnegie Mellon University) · Zico Kolter (Carnegie Mellon University / Bosch Center for AI)

Student-Teacher Curriculum Learning via Reinforcement Learning: Predicting Hospital Inpatient Admission Location
Rasheed El-Bouri (University of Oxford) · David Eyre (University of Oxford) · Peter Watkinson (Oxford University Hospitals NHS Foundation Trust) · Tingting Zhu (University of Oxford) · David Clifton (University of Oxford)

Option Discovery in the Absence of Rewards with Manifold Analysis
Amitay Bar (Technion) · Ronen Talmon (Technion - Israel Institute Of Technology) · Ron Meir (Technion Israeli Institute of Technology)

Generalisation error in learning with random features and the hidden manifold model
Federica Gerace (Politecnico di Torino) · Bruno Loureiro (Université de Paris Saclay) · Florent Krzakala (ENS) · Marc Mezard (ENS) · Lenka Zdeborova (CNRS)

Fast and Consistent Learning of Hidden Markov Models by Incorporating Non-Consecutive Correlations
Robert Mattila (KTH Royal Institute of Technology) · Cristian R. Rojas (KTH Royal Institute of Technology) · Eric Moulines (Ecole Polytechnique) · Vikram Krishnamurthy (Cornell University) · Bo Wahlberg (KTH Royal Institute of Technology)

Gradient-free Online Learning in Continuous Games with Delayed Rewards
Amélie Héliou (Criteo) · Panayotis Mertikopoulos (CNRS) · Zhengyuan Zhou (Stanford University)

Pseudo-Masked Language Models for Unified Language Model Pre-Training
Hangbo Bao (Harbin Institute of Technology) · Li Dong (Microsoft Research) · Furu Wei (Microsoft Research Asia) · Wenhui Wang (Microsoft Research) · Nan Yang (Microsoft Research Asia) · Xiaodong Liu (Microsoft Research) · Yu Wang (Microsoft Research) · Jianfeng Gao (Microsoft Research AI) · Songhao Piao (Harbin Institute of Technology) · Ming Zhou (Microsoft Research) · Hsiao-Wuen Hon (Microsoft Research)

Einsum Networks: Fast and Scalable Learning of Tractable Probabilistic Circuits
Robert Peharz (University of Cambridge) · Steven Lang (Technical University of

Darmstadt) · Antonio Vergari (University of California, Los Angeles) · Karl Stelzner (TU Darmstadt) · Alejandro Molina (TU Darmstadt) · Martin Trapp (Graz University of Technology) · Guy Van den Broeck (University of California, Los Angeles) · Kristian Kersting (TU Darmstadt) · Zoubin Ghahramani (University of Cambridge & Uber)

Polynomial Tensor Sketch for Element-wise Function of Low-Rank Matrix
Insu Han (KAIST) · Haim Avron (Tel Aviv University) · Jinwoo Shin (KAIST)

Inexact Tensor Methods with Dynamic Accuracies
Nikita Doikov (Université catholique de Louvain) · Yurii Nesterov (Universite catholique de Louvain)

k-means++: few more steps yield constant approximation
Davin Choo (ETH) · Christoph Grunau (ETH Zürich) · Julian Portmann (ETH Zürich) · Vaclav Rozhon (ETH)

Radioactive data: tracing through training
Alexandre Sablayrolles (Facebook AI Research) · Douze Matthijs (Facebook AI Research) · Cordelia Schmid (Inria/Google) · Herve Jegou (Facebook AI Research)

Doubly robust off-policy evaluation with shrinkage
Yi Su (Cornell University) · Maria Dimakopoulou (Stanford University) · Akshay Krishnamurthy (Microsoft Research) · Miroslav Dudik (Microsoft Research)

Fast Adaptation to New Environments via Policy-Dynamics Value Functions
Roberta Raileanu (NYU) · Max Goldstein (NYU) · Arthur Szlam (Facebook) · Facebook Rob Fergus (Facebook AI Research, NYU)

Neural Clustering Processes
Ari Pakman (Columbia University) · Yueqi Wang (Columbia University) · Catalin Mitelut (Columbia University) · JinHyung Lee (Columbia University) · Department of Statistics Liam Paninski (Department of Statistics, Columbia University)

Topologically Densified Distributions
Christoph Hofer (University of Salzburg) · Florian Graf (University of Salzburg) · Marc Niethammer (UNC) · Roland Kwitt ("University of Salzburg, Austria")

Low-loss connection of weight vectors: distribution-based approaches
Ivan Anokhin (Skolkovo Institute of Science and Technology) · Dmitry Yarotsky (Skolkovo Institute of Science and Technology)

Graph Filtration Learning
Christoph Hofer (University of Salzburg) · Florian Graf (University of Salzburg) · Bastian Rieck (ETH Zurich) · Marc Niethammer (UNC) · Roland Kwitt ("University of Salzburg, Austria")

Differentiable Product Quantization for Learning Compact Embedding Layers
Ting Chen (Google) · Lala Li (Google) · Yizhou Sun (UCLA)

Scalable Exact Inference in Multi-Output Gaussian Processes
Wessel Bruinsma (Invenia Labs) · Eric Perim Martins (Invenia Labs) · William Tebbutt (University of Cambridge) · Scott Hosking (British Antarctic Survey) · Arno Solin (Aalto University) · Richard E Turner (University of Cambridge)

Lower Complexity Bounds for Finite-Sum Convex-Concave Minimax Optimization Problems
Guangzeng Xie (Peking University) · Luo Luo (Shanghai Jiao Tong University) · yijiang lian (baidu) · Zhihua Zhang (Peking University)

Near-optimal Regret Bounds for Stochastic Shortest Path
Aviv Rosenberg (Tel Aviv University) · Alon Cohen (Technion and Google) · Yishay Mansour (Google and Tel Aviv University) · Haim Kaplan (TAU, GOOGLE)

The Usual Suspects? Reassessing Blame for VAE Posterior Collapse
Bin Dai (Samsung Research China - Beijing) · Ziyu Wang (Tsinghua University) · David Wipf (Microsoft Research)

It's Not What Machines Can Learn, It's What We Cannot Teach
Gal Yehuda (Technion, I.I.T) · Moshe Gabel (University of Toronto) · Assaf Schuster (Technion)

Guided Learning of Nonconvex Models through Successive Functional Gradient Optimization
Rie Johnson (RJ Research Consulting) · Tong Zhang (Hong Kong University of Science and Technology)

A Markov Decision Process Model for Socio-Economic Systems Impacted by Climate Change
Salman Sadiq Shuvo (University of South Florida) · Yasin Yilmaz (University of South Florida) · Alan Bush (University of South Florida) · Mark Hafen (University of South Florida)

Can Stochastic Zeroth-Order Frank-Wolfe Method Converge Faster for Non-Convex Problems?
Hongchang Gao (University of Pittsburgh) · Heng Huang (University of Pittsburgh)

Distance Metric Learning with Joint Representation Diversification
Xu Chu (Peking University) · Yang Lin (Peking University) · Xiting Wang (Microsoft Research Asia) · Xin Gao (Peking University) · Qi Tong (Peking University) · Hailong Yu (Peking University) · Yasha Wang (Peking University)

Meta-Learning with Shared Amortized Variational Inference
Ekaterina Iakovleva (INRIA) · Karteek Alahari (Inria) · Jakob Verbeek (Facebook)

Causal Effect Identifiability under Partial-Observability
Sanghack Lee (Columbia University) · Elias Bareinboim (Columbia)

Continuous Graph Neural Networks
Louis-Pascal Xhonneux (Mila / Université de Montréal) · Meng Qu (MILA) · Jian Tang (HEC Montreal & MILA)

Restarted Bayesian Online Change-point Detector achieves Optimal Detection Delay
REDA ALAMI (Orange Labs, University Paris-Saclay) · Odalric-Ambrym Maillard (Inria Lille - Nord Europe) · Raphael Feraud (Orange Labs)

Robust learning with the Hilbert-Schmidt independence criterion
Daniel Greenfeld (Technion) · Uri Shalit (Technion)

Bayesian Experimental Design for Implicit Models by Mutual Information Neural Estimation
Steven Kleinegesse (University of Edinburgh) · Michael Gutmann (University of Edinburgh)

Fast Differentiable Sorting and Ranking
Mathieu Blondel (NTT) · Olivier Teboul (Google Brain) · Quentin Berthet (Google Brain) · Josip Djolonga (Google AI, Zurich)

Learning for Dose Allocation in Adaptive Clinical Trials with Safety Constraints
Cong Shen (University of Virginia) · Zhiyang Wang (University of Pennsylvania) · Sofia Villar (University of Cambridge) · M van der Schaar (UCLA)

Tuning-free Plug-and-Play Proximal Algorithm for Inverse Imaging Problems
Kaixuan Wei (Beijing Institute of Technology) · Angelica Aviles-Rivero (University of Cambridge) · Jingwei Liang (University of Cambridge) · Ying Fu (Beijing Institute of Technology) · Carola-Bibiane Schönlieb (University of Cambridge) · Hua Huang (Beijing Institute of Technology)

Consistent Estimators for Learning to Defer to an Expert
Hussein Mozannar (Massachusetts Institute of Technology) · David Sontag (Massachusetts Institute of Technology)

A Graph to Graphs Framework for Retrosynthesis Prediction
Chence Shi (Peking University) · Minkai Xu (Shanghai Jiao Tong university) · Hongyu Guo (National Research Council Canada) · Ming Zhang (Peking University) · Jian Tang (HEC Montreal & MILA)

Fast computation of Nash Equilibria in Imperfect Information Games
Remi Munos (DeepMind) · Julien Perolat (DeepMind) · Jean-Baptiste Lespiau (DeepMind) · Mark Rowland (DeepMind) · Bart De Vylder (DeepMind) · Marc Lanctot (DeepMind) · Finbarr Timbers (DeepMind) · Daniel Hennes (DeepMind) · Shayegan Omidshafiei (DeepMind) · Audrunas Gruslys (DeepMind) · Mohammad Gheshlaghi Azar (Deepmind) · Edward Lockhart (DeepMind) · Karl Tuyls (DeepMind)

Invariant Rationalization
Shiyu Chang (MIT-IBM Watson AI Lab) · Yang Zhang (IBM-MIT Research Lab) · Mo Yu (IBM T. J. Watson) · Tommi Jaakkola (MIT)

Accelerated Stochastic Gradient-free and Projection-free Methods
Feihu Huang (Nanjing University of Aeronautics and Astronautics) · Lue Lue (Nanjing University of Aeronautics and Astronautics) · Songcan Chen (Nanjing University of Aeronautics and Astronautics)

Efficient Optimistic Exploration in Linear-Quadratic Regulators via Lagrangian Relaxation
Marc Abeille (Criteo) · Alessandro Lazaric (Facebook AI Research)

Implicit Regularization of Random Feature Models
Arthur Jacot (EPFL) · berfin simsek (EPFL) · Francesco Spadaro (EPFL) · Clement Hongler (EPFL) · Franck Gabriel (EPFL)

Missing Data Imputation using Optimal Transport
Boris Muzellec (ENSAE, Institut Polytechnique de Paris) · Julie Josse (Polytechnique) · Claire Boyer (LPSM, Sorbonne Université) · Marco Cuturi (Google and CREST/ENSAE)

Unsupervised Speech Decomposition via Triple Information Bottleneck
Kaizhi Qian (UIUC) · Yang Zhang (IBM-MIT Research Lab) · Shiyu Chang (MIT-IBM Watson AI Lab) · Mark Hasegawa-Johnson (University of Illinois) · David Cox (MIT-IBM Watson AI Lab)

Provable Representation Learning for Imitation Learning via Bi-level Optimization
Sanjeev Arora (Princeton University and Institute for Advanced Study) · Simon Du (Institute for Advanced Study) · Sham Kakade (University of Washington) · Yuping Luo (Princeton University) · Nikunj Umesh Saunshi (Princeton University)

Convergence of a Stochastic Gradient Method with Momentum for Non-Smooth Non-Convex Optimization
Vien Van Mai (KTH Royal Institute of Technology) · Mikael Johansson (KTH Royal Institute of Technology)

XTREME: A Massively Multilingual Multi-task Benchmark for Evaluating Cross-lingual Generalisation
Junjie Hu (Carnegie Mellon University) · Sebastian Ruder (DeepMind) · Aditya Siddhant (Google Research) · Graham Neubig (Carnegie Mellon University) · Orhan Firat (Google) · Melvin Johnson (Google)

Fair k-Centers via Maximum Matching
Matthew Jones (Khoury College of Computer Sciences) · Thy Nguyen (Northeastern University) · Huy Nguyen (Northeastern University)

Efficiently sampling functions from Gaussian process posteriors
James Wilson (Imperial College London) · Viacheslav Borovitskiy (St. Petersburg Department of Steklov Mathematical Institute of Russian Academy of Sciences (PDMI RAS)) · Alexander Terenin (Imperial College London) · Peter Mostowsky (Saint Petersburg State University) · Marc Deisenroth (University College London)

Characterizing Distribution Equivalence and Structure Learning for Cyclic and Acyclic Directed Graphs
AmirEmad Ghassami (UIUC) · Alan Yang (University of Illinois at Urbana-Champaign) · Negar Kiyavash (École Polytechnique Fédérale de Lausanne) · Kun Zhang (Carnegie Mellon University)

Inverse Active Sensing: Modeling and Understanding Timely Decision-Making
Daniel Jarrett (University of Cambridge) · Mihaela van der Schaar (University of Cambridge)

On Second-Order Group Influence Functions for Black-Box Predictions
Samyadeep Basu (UMD) · Xuchen You (University of Maryland) · Soheil Feizi (University of Maryland)

Safe Imitation Learning via Fast Bayesian Reward Inference from Preferences
Daniel Brown (University of Texas at Austin) · Scott Niekum (University of Texas at Austin) · Russell Coleman (University of Texas at Austin) · Ravi Srinivasan (University of Texas at Austin)

Randomly Projected Additive Gaussian Processes for Regression
Ian Delbridge (Cornell University) · David S Bindel (Cornell University) · Andrew Wilson (New York University)

Attentive Group Equivariant Convolutional Networks
David W. Romero (Vrije Universiteit Amsterdam) · Erik Bekkers (University of Amsterdam) · Jakub Tomczak (Vrije Universiteit Amsterdam) · Mark Hoogendoorn (Vrije Universiteit Amsterdam)

Learning Compound Tasks without Task-specific Knowledge via Imitation and Self-supervised Learning
Sang-Hyun Lee (Seoul National University) · Seung-Woo Seo (Seoul National University)

Confidence Sets and Hypothesis Testing in a Likelihood-Free Inference Setting
Niccolo Dalmasso (Carnegie Mellon University) · Rafael Izbicki (UFSCar) · Ann Lee (Carnegie Mellon University)

Curvature-corrected learning dynamics in deep neural networks
Dongsung Huh (MIT-IBM Watson AI Lab)

Tightening Exploration in Upper Confidence Reinforcement Learning
Hippolyte Bourel (ENS Rennes) · Odalric-Ambrym Maillard (Inria Lille - Nord Europe) · Mohammad Sadegh Talebi (University of Copenhagen)

Bootstrap Latent-Predictive Representations for Multitask Reinforcement Learning
Zhaohan Guo (DeepMind) · Bernardo Avila Pires (DeepMind) · Mohammad Gheshlaghi Azar (Deepmind) · Bilal Piot (DeepMind) · Florent Altché (DeepMind) · Jean-Bastien Grill (DeepMind) · Remi Munos (DeepMind)

Discriminative Adversarial Search for Abstractive Summarization
Thomas Scialom (reciTAL) · Paul-Alexis Dray (reciTAL) · Sylvain Lamprier (LIP6 - Sorbonne Universités) · Benjamin Piwowarski (Sorbonne Université) · Jacopo Staiano (reciTAL)

A Swiss Army Knife for Minimax Optimal Transport
Sofien Dhouib (CREATIS UMR CNRS 5220) · Ievgen Redko (Laboratoire Hubert Curien) · Tanguy Kerdoncuff (Laboratoire Hubert Curien) · Rémi Emonet (Laboratoire Hubert Curien) · Marc Sebban (Jean Monnet University)

Invariant Causal Prediction for Block MDPs
Clare Lyle (University of Oxford) · Amy Zhang (McGill University) · Angelos Filos (University of Oxford) · Shagun Sodhani (Facebook AI Research) · Marta Kwiatkowska (Oxford University) · Yarin Gal (University of Oxford) · Doina Precup (McGill University / DeepMind) · Joelle Pineau (McGill University / Facebook)

Involutive MCMC: One Way to Derive Them All
Kirill Neklyudov (Samsung) · Max Welling (University of Amsterdam & Qualcomm) · Evgenii Egorov (Skolkovo Institute of Science and Technology) · Dmitry Vetrov (Higher School of Economics, Samsung AI Center Moscow)

Adversarial Learning Guarantees for Linear Hypotheses and Neural Networks
Pranjal Awasthi (Rutgers University) · Natalie Frank (NYU) · Mehryar Mohri (Google Research and Courant Institute of Mathematical Sciences)

Deep Reinforcement Learning with Smooth Policy
Qianli Shen (Peking University) · Yan Li (Georgia Tech) · Haoming Jiang (Georgia Tech) · Zhaoran Wang (Northwestern) · Tuo Zhao (Gatech)

On the Power of Compressed Sensing with Generative Models
Akshay Kamath (University of Texas at Austin) · Eric Price (UT-Austin) · Sushrut Karmalkar (University of Texas at Austin)

Laplacian Regularized Few-Shot Learning
Imtiaz Ziko (ETS Montreal) · Jose Dolz (ETS Montreal) · Eric Granger (ETS Montreal) · Ismail Ben Ayed (ETS Montreal)

Neural Datalog Through Time: Informed Temporal Modeling via Logical Specification
Hongyuan Mei (Johns Hopkins University) · Guanghui Qin (JOHNS HOPKINS UNIVERSITY) · Minjie Xu (Bloomberg LP) · Jason Eisner (Johns Hopkins University + Microsoft Semantic Machines)

Up or Down? Adaptive Rounding for Post-Training Quantization
Markus Nagel (Qualcomm Research) · Rana Ali Amjad (Qualcomm) · Marinus van Baalen (Qualcomm) · Christos Louizos (Qualcomm AI Research) · Tijmen Blankevoort (Qualcomm)

A quantile-based approach for hyperparameter transfer learning
David Salinas (Naverlabs Europe) · Huibin Shen (Amazon) · Valerio Perrone (Amazon)

Inductive Bias-driven Reinforcement Learning For Efficient Schedules in Heterogeneous Clusters
Subho Banerjee (University of Illinois at Urbana-Champaign) · Saurabh Jha (UIUC) · Zbigniew Kalbarczyk (University of Illinois at Urbana-Champaign) · Ravishankar Iyer (University of Illinois at Urbana-Champaign)

Adversarial Robustness for Code
Pavol Bielik (ETH Zurich) · Martin Vechev (ETH Zurich)

Nearly Optimal Risk Bounds for Kernel K-Means
Yong Liu (Institute of Information Engineering, CAS) · Lizhong Ding (Inception Institute of Artificial Intelligence) · Hua Zhang (Institute of Information Engineering,Chinese Academy of Sciences) · Wenqi Ren (IIE, CAS) · Xiao Zhang (Tianjin University) · Shali Jiang (Washington University in St. Louis) · Xinwang Liu (National University of Defense Technology) · Weiping Wang (Institute of Information Engineering, CAS, China)

The Boomerang Sampler
Joris Bierkens (Vrije Universiteit Amsterdam) · Sebastiano Grazzi (Technische Universiteit Delft) · Kengo Kamatani (Osaka University) · Gareth Roberts (University of Warwick)

Weakly-Supervised Disentanglement Without Compromises
Francesco Locatello (ETH Zurich - Max Planck Institute) · Ben Poole (Google Brain) · Gunnar Raetsch (ETH Zurich) · Bernhard Schölkopf (MPI for Intelligent Systems Tübingen, Germany) · Olivier Bachem (Google Brain) · Michael Tschannen (Google Brain)

Predictive Sampling with Forecasting Autoregressive Models
Auke Wiggers (Qualcomm AI Research) · Emiel Hoogeboom (University of Amsterdam)

InfoGAN-CR: Disentangling Generative Adversarial Networks with Contrastive Regularizers
Zinan Lin (Carnegie Mellon University) · Kiran Thekumparampil (University of Illinois at Urbana-Champaign) · Giulia Fanti (CMU) · Sewoong Oh (University of Washington)

TrajectoryNet: A Dynamic Optimal Transport Network for Modeling Cellular Dynamics
Alexander Tong (Yale University) · Jessie Huang (Yale University) · Guy Wolf (Université de Montréal) · David van Dijk (Yale University) · Smita Krishnaswamy (Yale University)

The role of regularization in classification of high-dimensional noisy Gaussian mixture
Francesca Mignacco (CEA Saclay) · Florent Krzakala (ENS) · Yue Lu (Harvard University, USA) · Pierfrancesco Urbani (Institut de Physique Théorique) · Lenka Zdeborova (CNRS)

Normalizing Flows on Tori and Spheres
Danilo J. Rezende (DeepMind) · George Papamakarios (DeepMind) · Sebastien Racaniere (DeepMind) · Michael S Albergo (New York University) · Gurtej Kanwar (Massachusetts Institute of Technology) · Phiala Shanahan (Massachusetts Institute of Technology) · Kyle Cranmer (New York University, CERN)

Structured Linear Contextual Bandits: A Sharp and Geometric Smoothed Analysis
Vidyashankar Sivakumar (Walmart Labs) · Steven Wu (University of Minnesota) · Arindam Banerjee (University of Minnesota)

Simple and sharp analysis of k-means||
Václav Rozhoň (ETH)

Efficient proximal mapping of the path-norm regularizer of shallow networks
Fabian Latorre (EPFL) · Paul Rolland (Ecole Polytechnique Fédérale de Lausanne) · Nadav Hallak (EPFL) · Volkan Cevher (EPFL)

Regularized Optimal Transport is Ground Cost Adversarial
François-Pierre Paty (ENSAE Paris) · Marco Cuturi (Google and CREST/ENSAE)

Automatic Shortcut Removal for Self-Supervised Representation Learning
Matthias Minderer (Google Research) · Olivier Bachem (Google Brain) · Neil Houlsby (Google) · Michael Tschannen (Google Research)

Fair Learning with Private Demographic Data
Hussein Mozannar (Massachusetts Institute of Technology) · Mesrob Ohannessian (University of Illinois at Chicago) · Nati Srebro (Toyota Technological Institute at Chicago)

Deep Divergence Learning
Kubra Cilingir (Boston University) · Rachel Manzelli (Boston University) · Brian Kulis (Boston University)

A new regret analysis for Adam-type algorithms
Ahmet Alacaoglu (EPFL) · Yura Malitsky (EPFL) · Panayotis Mertikopoulos (CNRS) · Volkan Cevher (EPFL)

Accelerated Message Passing for Entropy-Regularized MAP Inference
Jonathan Lee (UC Berkeley) · Aldo Pacchiano (UC Berkeley) · Peter Bartlett (UC Berkeley) · Michael Jordan (UC Berkeley)

Dissecting Non-Vacuous Generalization Bounds based on the Mean-Field Approximation
Konstantinos Pitas (Ecole Polytechnique Federale de Lausanne)

(Individual) Fairness for k-Clustering
Sepideh Mahabadi (Toyota Technological Institute at Chicago) · Ali Vakilian (University of Wisconsin-Madison)

Relaxing Bijectivity Constraints with Continuously Indexed Normalising Flows
Rob Cornish (Oxford) · Anthony Caterini (University of Oxford) · George Deligiannidis (Oxford) · Arnaud Doucet (Oxford University)

Gamification of Pure Exploration for Linear Bandits
Rémy Degenne (Inria) · Pierre Menard (Inria) · Xuedong Shang (Inria) · Michal Valko (DeepMind)

Growing Adaptive Multi-hyperplane Machines
Nemanja Djuric (Uber ATG) · Zhuang Wang (Facebook, Inc.) · Slobodan Vucetic (Temple University)

Generative Teaching Networks: Accelerating Neural Architecture Search by Learning to Generate Synthetic Training Data
Felipe Petroski Such (Uber AI Labs) · Aditya Rawal (Uber AI Labs) · Joel Lehman () · Kenneth Stanley (Uber AI and University of Central Florida) · Jeffrey Clune (Open AI)

Structured Prediction with Partial Labelling through the Infimum Loss
Vivien Cabannnes (INRIA) · Francis Bach (INRIA - Ecole Normale Supérieure) · Alessandro Rudi (École Normale Supérieure)

ControlVAE: Controllable Variational Autoencoder
Huajie Shao (University of Illinois at Urbana-Champaign) · Shuochao Yao (University of Illinois at Urbana-Champaign) · Dachun Sun (University of Illinois at Urbana-Champaign) · Aston Zhang (Amazon AI) · Shengzhong Liu (University of Illinois at Urbana-Champaign) · Dongxin Liu (University of Illinois at Urbana-Champaign) · Jun Wang (Alibaba Group) · Tarek Abdelzaher (University of Illinois at Urbana-Champaign)

On Semi-parametric Inference for BART
Veronika Rockova (University of Chicago)

Simple and Scalable Epistemic Uncertainty Estimation Using a Single Deep Deterministic Neural Network
Joost van Amersfoort (University of Oxford) · Lewis Smith (University of Oxford) · Yee Whye Teh (Oxford and DeepMind) · Yarin Gal (University of Oxford)

Ordinal Non-negative Matrix Factorization for Recommendation
Olivier Gouvert (CNRS, IRIT) · Thomas Oberlin (ISAE-SUPAERO) · Cedric Fevotte (CNRS)

NetGAN without GAN: From Random Walks to Low-Rank Approximations
Luca Rendsburg (Eberhard Karls University of Tübingen) · Holger Heidrich (Eberhard Karls Universität Tübingen) · Ulrike von Luxburg (U Tübingen)

On the Iteration Complexity of Hypergradient Computations
Riccardo Grazzi (Istituto Italiano di Tecnologia - University College London) · Saverio Salzo (Istituto Italiano di Tecnologia) · Massimiliano Pontil (Istituto Italiano di Tecnologia and University College London) · Luca Franceschi (Istituto Italiano di Tecnologia - University College London)

Skew-Fit: State-Covering Self-Supervised Reinforcement Learning
Vitchyr Pong (UC Berkeley) · Murtaza Dalal (UC Berkeley) · Steven Lin (UC Berkeley) · Ashvin Nair (UC Berkeley) · Shikhar Bahl (UC Berkeley/Carnegie Mellon University) · Sergey Levine (UC Berkeley)

Stochastic Optimization for Regularized Wasserstein Estimators
Marin Ballu (University of Cambridge) · Quentin Berthet (Google Brain) · Francis Bach (INRIA - Ecole Normale Supérieure)

LP-SparseMAP: Differentiable Relaxed Optimization for Sparse Structured Prediction
Vlad Niculae (Instituto de Telecomunicações) · Andre Filipe Torres Martins (Instituto de Telecomunicacoes)

Problems with Shapley-value-based explanations as feature importance measures
I. Elizabeth Kumar (University of Utah) · Suresh Venkatasubramanian (University of Utah, USA) · Carlos Scheidegger (The University of Arizona) · Sorelle Friedler (Haverford College)

Model-free Reinforcement Learning in Infinite-horizon Average-reward Markov Decision Processes
Chen-Yu Wei (University of Southern California) · Mehdi Jafarnia (University of Southern California) · Haipeng Luo (University of Southern California) · Hiteshi Sharma (University of Southern California) · Rahul Jain (USC)

Near-linear time Gaussian process optimization with adaptive batching and resparsification
Daniele Calandriello (IIT) · Luigi Carratino (University of Genoa) · Alessandro Lazaric (Facebook AI Research) · Michal Valko (DeepMind) · Lorenzo Rosasco (unige, mit, iit)

Parallel Algorithm for Non-Monotone DR-Submodular Maximization
Alina Ene (Boston University) · Huy Nguyen (Northeastern University)

Structure Adaptive Algorithms for Stochastic Bandits
Rémy Degenne (Inria) · Han Shao (Toyota Technological Institute at Chicago) · Wouter Koolen (Centrum Wiskunde & Informatica, Amsterdam)

Spectrum Dependent Learning Curves in Kernel Regression and Wide Neural Networks
Blake Bordelon (Harvard University) · Abdulkadir Canatar (Harvard University) · Cengiz Pehlevan (Harvard University)

Preference modelling with context-dependent salient features
Amanda Bower (University of Michigan) · Laura Balzano (University of Michigan)

Infinite attention: NNGP and NTK for deep attention networks
Jiri Hron (University of Cambridge) · Yasaman Bahri (Google Brain) · Jascha Sohl-Dickstein (Google Brain) · Roman Novak (Google Brain)

Fast Learning of Graph Neural Networks with Guaranteed Generalizability: One-hidden-layer Case
shuai zhang (Rensselaer Polytechnic Institute) · Meng Wang (Rensselaer Polytechnic Institute) · Sijia Liu (MIT-IBM Watson AI Lab) · Pin-Yu Chen (IBM Research AI) · Jinjun Xiong (IBM Thomas J. Watson Research Center)

Efficient Domain Generalization via Common-Specific Low-Rank Decomposition
Vihari Piratla (IIT Bombay) · Praneeth Netrapalli (Microsoft Research) · Sunita Sarawagi (IIT Bombay)

Identifying the Reward Function by Anchor Actions
Sinong Geng (Princeton University) · Houssam Nassif (amazon) · Carlos Manzanares (Amazon) · Max Reppen (Princeton) · Ronnie Sircar (Princeton)

No-Regret and Incentive-Compatible Online Learning
Rupert Freeman (Microsoft Research) · David Pennock (Rutgers University) · Charikleia Podimata (Harvard University) · Jennifer Wortman Vaughan (Microsoft Research)

Probing Emergent Semantics in Predictive Agents via Question Answering
Abhishek Das (Georgia Tech) · Federico Carnevale (Deepmind) · Hamza Merzic (DeepMind) · Laura Rimell () · Rosalia Schneider (DeepMind) · Josh Abramson (DeepMind) · Alden Hung (DeepMind) · Arun Ahuja (DeepMind) · Stephen Clark (University of Cambridge/Deepmind) · Greg Wayne (DeepMind) · Feilx Hill (Deepmind)

Meta-learning with Stochastic Linear Bandits
Leonardo Cella (University of Milan) · Alessandro Lazaric (Facebook AI Research) · Massimiliano Pontil (Istituto Italiano di Tecnologia and University College London)

A Unified Theory of Decentralized SGD with Changing Topology and Local Updates
Anastasiia Koloskova (EPFL) · Nicolas Loizou (Mila, Université de Montréal) · Sadra Boreiri (EPFL) · Martin Jaggi (EPFL) · Sebastian Stich (EPFL)

AdaScale SGD: A User-Friendly Algorithm for Distributed Training
Tyler Johnson (Apple) · Pulkit Agrawal (Apple) · Haijie Gu (Apple) · Carlos Guestrin (Apple & Univesity of Washington)

Kinematic State Abstraction and Provably Efficient Rich-Observation Reinforcement Learning
Dipendra Misra (Microsoft) · Mikael Henaff (Microsoft) · Akshay Krishnamurthy (Microsoft Research) · John Langford (Microsoft Research)

Logistic Regression for Massive Data with Rare Events
HaiYing Wang (University of Connecticut)

Automated Synthetic-to-Real Generalization
Wuyang Chen (Texas A&M University) · Zhiding Yu (NVIDIA) · Zhangyang Wang (Texas A&M University) · Anima Anandkumar (Caltech)

Online Learning with Dependent Stochastic Feedback Graphs
Corinna Cortes (Google Research) · Giulia DeSalvo (Google Research) · Claudio Gentile (INRIA and Google) · Mehryar Mohri (Google Research and Courant Institute of Mathematical Sciences) · Ningshan Zhang (New York University)

Sparse Sinkhorn Attention
Yi Tay (Google) · Dara Bahri (Google) · Liu Yang (Google) · Donald Metzler (Google) · Da-Cheng Juan (Google)

Online Continual Learning from Imbalanced Data
Aristotelis Chrysakis (KU Leuven) · Marie-Francine Moens (KU Leuven)

Differentially Private Set Union
Pankaj Gulhane (Microsoft) · Sivakanth Gopi (Microsoft) · Janardhan Kulkarni (Microsoft Research) · Judy Hanwen Shen (Microsoft) · Milad Shokouhi (Microsoft) · Sergey Yekhanin (Microsoft)

The continuous categorical: a novel simplex-valued exponential family
Elliott Gordon-Rodriguez (Columbia University) · Gabriel Loaiza-Ganem (Layer 6 AI) · John Cunningham (Columbia)

Minimax-Optimal Off-Policy Evaluation with Linear Function Approximation
Yaqi Duan (Princeton University) · Zeyu Jia (Peking University) · Mengdi Wang (Princeton University)

Enhanced POET: Open-ended Reinforcement Learning through Unbounded Invention of Learning Challenges and their Solutions
Rui Wang (Uber AI) · Joel Lehman () · Aditya Rawal (Uber AI Labs) · Jiale Zhi (Uber AI) · Yulun Li (Uber AI) · Jeffrey Clune (Open AI) · Kenneth Stanley (Uber AI and University of Central Florida)

Set Functions for Time Series
Max Horn (MLCB, D-BSSE, ETH Zurich) · Michael Moor (ETH Zurich) · Christian Bock (ETH Zurich) · Bastian Rieck (ETH Zurich) · Karsten Borgwardt (ETH Zurich)

Individual Calibration with Randomized Forecasting
Shengjia Zhao (Stanford University) · Tengyu Ma (Stanford) · Stefano Ermon (Stanford University)

Bayesian Differential Privacy for Machine Learning
Aleksei Triastcyn (EPFL) · Boi Faltings (EPFL)

Causal Modeling for Fairness In Dynamical Systems
Elliot Creager (University of Toronto) · David Madras (University of Toronto) · Toniann Pitassi (University of Toronto) · Richard Zemel (Vector Institute)

Learning General-Purpose Controllers via Locally Communicating Sensorimotor Modules
Wenlong Huang (UC Berkeley) · Igor Mordatch (OpenAI) · Deepak Pathak (UC Berkeley)

Visual Grounding of Learned Physical Models
Yunzhu Li (MIT) · Toru Lin (MIT) · Kexin Yi (Harvard University) · Daniel Bear (Stanford University) · Daniel Yamins (Stanford University) · Jiajun Wu (Stanford University) · Josh Tenenbaum (MIT) · Antonio Torralba (MIT)

Task-Oriented Active Perception and Planning in Environments with Partially Known Semantics
Mahsa Ghasemi (The University of Texas at Austin) · Erdem Bulgur (University of Texas at Austin) · Ufuk Topcu (University of Texas at Austin)

Test-Time Training for Generalization under Distribution Shifts
Yu Sun () · Xiaolong Wang (UC Berkeley) · Zhuang Liu (UC Berkeley) · John Miller (University of California, Berkeley) · Alexei Efros (UC Berkeley) · University of California Moritz Hardt (University of California, Berkeley)

Auto-GAN-Distiller: Searching to Compress Generative Adversarial Networks
Yonggan Fu (Rice University) · Wuyang Chen (Texas A&M University) · Haotao Wang (Texas A&M University) · Haoran Li (Rice University) · Yingyan Lin (Rice University) · Zhangyang Wang (Texas A&M University)

Associative Memory in Iterated Overparameterized Sigmoid Autoencoders
Yibo Jiang (Harvard University) · Cengiz Pehlevan (Harvard University)

Adaptive Reward-Poisoning Attacks against Reinforcement Learning
Xuezhou Zhang (UW-Madison) · Yuzhe Ma (Univ. of Wisconsin-Madison) · Adish Singla (Max Planck Institute (MPI-SWS)) · Jerry Zhu (University of Wisconsin-Madison)

Planning to Explore via Latent Disagreement
Ramanan Sekar (University of Pennsylvania) · Oleh Rybkin (University of Pennsylvania / UC Berkeley (Visiting)) · Kostas Daniilidis (University of Pennsylvania) · Pieter Abbeel (UC Berkeley & Covariant) · Danijar Hafner (Google Brain & University of Toronto) · Deepak Pathak (UC Berkeley)

Defense Through Diverse Directions
Christopher Bender (The University of North Carolina Computer Science Department) · Yang Li (Department of Computer Science, University of North Carolina at Chapel Hill) · Yifeng Shi (University of North Carolina at Chapel Hill) · Michael K. Reiter (UNC at Chapel Hill) · Junier Oliva (UNC-Chapel Hill)

Beyond Synthetic Noise: Deep Learning on Controlled Noisy Labels
Lu Jiang (Google) · Di Huang (Google) · Mason Liu (Cornell) · Weilong Yang (Google Inc.)

Confidence-Calibrated Adversarial Training: Generalizing to Unseen Attacks
David Stutz (Max Planck Institute for Informatics) · Matthias Hein (University of Tübingen) · Bernt Schiele (MPI Informatics)

Online Control of the False Coverage Rate and False Sign Rate
Asaf Weinstein (The Hebrew University of Jerusalem) · Aaditya Ramdas (Carnegie Mellon University)

Online Convex Optimization in the Random Order Model
Dan Garber (Technion) · Gal Korcia (Technion - Israel Institute of Technology) · Kfir Levy (Technion)

A Flexible Latent Space Model for Multilayer Networks
Xuefei Zhang (University of Michigan) · Songkai Xue (University of Michigan) · Ji Zhu (University of Michigan)

Estimation of Bounds on Potential Outcomes For Decision Making
Maggie Makar (MIT) · Fredrik Johansson (Chalmers University of Technology) · John Guttag (MIT) · David Sontag (Massachusetts Institute of Technology)

Deep Gaussian Markov random fields
Per Sidén (Linköping University) · Fredrik Lindsten (Linköping University)

Generalization Error of Generalized Linear Models in High Dimensions
Melikasadat Emami (University of California Los Angeles) · Mojtaba Sahraee-Ardakan (UCLA) · Parthe Pandit (UCLA) · Sundeep Rangan (NYU) · Alyson Fletcher (UCLA)

Poisson Learning: Graph Based Semi-Supervised Learning At Very Low Label Rates
Jeff Calder (University of Minnesota) · Brendan Cook (University of Minnesota) · Matthew Thorpe (University of Manchester) · Dejan Slepcev (Carnegie Mellon University)

Sequential Transfer in Reinforcement Learning with a Generative Model
Andrea Tirinzoni (Politecnico di Milano) · Riccardo Poiani (Politecnico di Milano) · Marcello Restelli (Politecnico di Milano)

Finite-Time Convergence in Continuous-Time Optimization
Orlando Romero (Rensselaer Polytechnic Institute) · mouhacine Benosman (MERL)

Feature Quantization Improves GAN Training
Yang Zhao (University at Buffalo) · Chunyuan Li (Microsoft Research) · Ping Yu (Sony Interactive Entertainment LLC) · Jianfeng Gao (Microsoft Research AI) · Changyou Chen (SUNY Buffalo)

Temporal Logic Point Processes
Shuang Li (Harvard University) · Lu Wang (East China Normal University) · Ruizhi Zhang (University of Nebraska-Lincoln) · xiaofu Chang (Ant Financial Services Group) · Xuqin Liu (Ant Financial Services Group) · Yao Xie (Georgia Institute of Technology) · Yuan Qi (Ant Financial Services Group) · Le Song (Georgia Institute of Technology)

Hallucinative Topological Memory for Zero-Shot Visual Planning
Thanard Kurutach (UC Berkeley) · Kara Liu (UC Berkeley) · Aviv Tamar (Technion) · Pieter Abbeel (UC Berkeley) · Christine Tung (UC Berkeley)

Learning Attentive Meta-Transfer
Jaesik Yoon (sap labs korea) · Gautam Singh (Rutgers Univerity) · Sungjin Ahn (Rutgers University)

Optimizing Dynamic Structures with Bayesian Generative Search
Minh Hoang (Carnegie Mellon University) · Carleton Kingsford (Carnegie Mellon University)

Amortized Finite Element Analysis for Fast PDE-Constrained Optimization
Tianju Xue (Princeton University) · Alex Beatson (Princeton University) · Sigrid Adriaenssens (Princeton University) · Ryan P. Adams (Princeton University)

Preselection Bandits
Viktor Bengs (University of Paderborn) · Eyke Hüllermeier (Paderborn University)

Peer Loss Functions: Learning from Noisy Labels without Knowing Noise Rates
Yang Liu (UCSC) · Hongyi Guo (Shanghai Jiao Tong University)

Rank Aggregation from Pairwise Comparisons in the Presence of Adversarial Corruptions
Prathamesh Patil (University of Pennsylvania) · Arpit Agarwal (University of Pennsylvania) · Shivani Agarwal (University of Pennsylvania) · Sanjeev Khanna (U of Pennsylvania)

Extrapolation for Large-batch Training in Deep Learning
Tao LIN (EPFL) · Lingjing Kong (EPFL) · Sebastian Stich (EPFL) · Martin Jaggi (EPFL)

VideoOneNet: Bidirectional Convolutional Recurrent OneNet with Trainable Data Steps for Video Processing
Zoltán Milacski (Eötvös Loránd University) · Barnabás Póczos (CMU) · Andras Lorincz (Eotvos Lorand University)

Bio-Inspired Hashing for Unsupervised Similarity Search
Chaitanya Ryali (UC San Diego) · John Hopfield (Princeton University) · Leopold Grinberg (IBM Research) · Dmitry Krotov (IBM Research)

MetaFun: Meta-Learning with Iterative Functional Updates
Jin Xu (University of Oxford) · Jean-Francois Ton (University of Oxford) · Hyunjik Kim (University of Oxford, DeepMind) · Adam Kosiorek (DeepMind) · Yee Whye Teh (Oxford and DeepMind)

Learning and Simulation in Generative Structured World Models
Zhixuan Lin (Zhejiang University) · Yi-Fu Wu (Rutgers University) · Skand Peri (Rutgers University, New Jersey) · Bofeng Fu (Tianjin University) · Jindong Jiang (Rutgers University) · Sungjin Ahn (Rutgers University)

Random Hypervolume Scalarizations for Provable Multi-Objective Black Box Optimization
Richard Zhang (Google Brain) · Daniel Golovin (Google, Inc.)

SGD Learns One-Layer Networks in WGANs
Qi Lei (University of Texas at Austin) · Jason Lee (Princeton) · Alexandros Dimakis (UT Austin) · Constantinos Daskalakis (MIT)

Implicit Class-Conditioned Domain Alignment for Unsupervised Domain Adaptation
Xiang Jiang (Imagia, Dalhousie University) · Qicheng Lao (Mila;Imagia) · Stan Matwin (Dalhousie University) · Mohammad Havaei (Imagia)

Interference and Generalization in Temporal Difference Learning
Emmanuel Bengio (McGill University) · Joelle Pineau (McGill University / Facebook) · Doina Precup (McGill University / DeepMind)

CoMic: Co-Training and Mimicry for Reusable Skills
Leonard Hasenclever (DeepMind) · Fabio Pardo (Imperial College London) · Raia Hadsell (DeepMind) · Nicolas Heess (DeepMind) · Josh Merel (DeepMind)

Provably Efficient Model-based Policy Adaptation
Yuda Song (University of California, San Diego) · Aditi Mavalankar (University of California San Diego) · Wen Sun (Microsoft Research) · Sicun Gao (University of California, San Diego)

Optimizer Benchmarking Needs to Account for Hyperparameter Tuning
Prabhu Teja Sivaprasad (Idiap Research Institute) · Florian Mai (Idiap Research Institute) · Thijs Vogels (EPFL) · Martin Jaggi (EPFL) · Francois Fleuret (Idiap research institute)

From Local SGD to Local Fixed Point Methods for Federated Learning
Grigory Malinovskiy (Moscow Institute of Physics and Technology) · Dmitry Kovalev (KAUST) · Elnur Gasanov (KAUST) · Laurent CONDAT (KAUST) · Peter Richtarik (KAUST)

Unraveling Meta-Learning: Understanding Feature Representations for Few-Shot Tasks
Micah Goldblum (University of Maryland) · Liam Fowl (University of Maryland) · Renkun Ni (University of Maryland) · Steven Reich (University of Maryland) · Valeriia Cherepanova (University of Maryland) · Tom Goldstein (University of Maryland)

Federated Learning with Only Positive Labels
Felix Xinnan Yu (Google AI) · Ankit Singh Rawat (Google) · Aditya Menon (Google Research) · Sanjiv Kumar (Google Research, NY)

Causal Inference using Gaussian Processes with Structured Latent Confounders
Sam Witty (University of Massachusetts, Amherst) · Kenta Takatsu (University of Massachusetts Amherst) · David Jensen (University of Massachusetts Amherst) · Vikash Mansinghka (Massachusetts Institute of Technology)

T-Basis: a Compact Representation for Neural Networks
Anton Obukhov (ETH Zurich) · Maxim Rakhuba (ETH Zurich) · Menelaos Kanakis (ETH Zurich) · Stamatios Georgoulis (ETH Zurich) · Dengxin Dai (ETH Zurich) · Luc Van Gool (ETH Zurich)

Familywise Error Rate Control by Interactive Unmasking
Boyan Duan (Carnegie Mellon University) · Aaditya Ramdas (Carnegie Mellon University) · Larry Wasserman (Carnegie Mellon University)

Learning to Branch for Multi-Task Learning
Pengsheng Guo (Apple) · Chen-Yu Lee (Apple) · Daniel Ulbricht (Apple)

Augmenting Continuous Time Bayesian Networks with Clocks
Nicolai Engelmann (Darmstadt University of Technology) · Dominik Linzner (Technische Universität Darmstadt) · Heinz Koeppl (TU Darmstadt)

IPBoost – Non-Convex Boosting via Integer Programming
Sebastian Pokutta (ZIB) · Marc Pfetsch (TU Darmstadt)

On Efficient Constructions of Checkpoints
Yu Chen (College of William and Mary) · Zhenming LIU (College of William & Mary) · Bin Ren (College of William and Mary) · Xin Jin (Johns Hopkins University)

Feature Selection using Stochastic Gates
Yutaro Yamada (Yale University) · Ofir Lindenbaum (Yale) · Sahand Negahban (YALE) · Yuval Kluger (Yale School of Medicine)

How to train your Neural ODE
Chris Finlay (McGill University) · Joern-Henrik Jacobsen (Vector Institute and University of Toronto) · Levon Nurbekyan (UCLA) · Adam M Oberman (McGill University)

Evaluating Lossy Compression Rates of Deep Generative Models
Sicong Huang (University of Toronto) · Alireza Makhzani (University of Toronto) · Yanshuai Cao (Borealis AI) · Roger Grosse (University of Toronto and Vector Institute)

Mix-n-Match : Ensemble and Compositional Methods for Uncertainty Calibration in Deep Learning
Jize Zhang (Lawrence Livermore National Laboratory) · Bhavya Kailkhura (LLNL) · T. Yong-Jin Han (LLNL)

Learning Adversarially Robust Representations via Worst-Case Mutual Information Maximization
Sicheng Zhu (University of Virginia) · Xiao Zhang (University of Virginia) · David Evans (University of Virginia)

Stochastic Regret Minimization in Extensive-Form Games
Gabriele Farina (Carnegie Mellon University) · Christian Kroer (Columbia University) · Tuomas Sandholm (Carnegie Mellon University)

Simultaneous Inference for Massive Data: Distributed Bootstrap
Yang Yu (Purdue University) · Shih-Kang Chao (University of Missouri) · Guang Cheng (Purdue University)

Stabilizing Differentiable Architecture Search via Perturbation-based Regularization
Xiangning Chen (University of California, Los Angeles) · Cho-Jui Hsieh (UCLA)

Boosting Frank-Wolfe by Chasing Gradients
Cyrille Combettes (Georgia Institute of Technology) · Sebastian Pokutta (ZIB)

Concise Explanations of Neural Networks using Adversarial Training
Prasad Chalasani (MediaMath) · Jiefeng Chen (University of Wisconsin-Madison) · Amrita Roy Chowdhury (University of Wisconsin-Madison) · Xi Wu (Google) · Somesh Jha (University of Wisconsin, Madison)

Quantum Boosting
Srinivasan Arunachalam (IBM) · Reevu Maity (Oxford University)

Information-Theoretic Local Minima Characterization and Regularization
Zhiwei Jia (University of California, San Diego) · Hao Su (UCSD)

Kernel interpolation with continuous volume sampling
Ayoub Belhadji (Ecole Centrale de Lille) · Rémi Bardenet (CNRS and Univ. Lille) · Pierre Chainais (Centrale Lille / CRIStAL CNRS UMR 9189)

Efficient Identification in Linear Structural Causal Models with Auxiliary Cutsets
Daniel Kumor (Purdue University) · Carlos Cinelli (UCLA) · Elias Bareinboim (Columbia)

Partial Trace Regression and Low-Rank Kraus Decomposition
Hachem Kadri (Aix-Marseille University) · Stephane Ayache (AMU LIS) · Riikka Huusari (Aalto University) · alain rakotomamonjy (Universite de Rouen Normandie / Criteo AI Lab) · Ralaivola Liva (Criteo AI Lab)

Constant Curvature Graph Convolutional Networks
Gregor Bachmann (ETH Zurich) · Gary Becigneul (MIT) · Octavian Ganea (MIT)

Educating Text Autoencoders: Latent Representation Guidance via Denoising
Tianxiao Shen (MIT) · Jonas Mueller (Amazon Web Services) · Regina Barzilay (MIT CSAIL) · Tommi Jaakkola (MIT)

Generalization via Derandomization
Jeffrey Negrea (University of Toronto) · Daniel Roy (Univ of Toronto | Toronto) · Gintare Karolina Dziugaite (Element AI)

Inductive Relation Prediction by Subgraph Reasoning
Komal Teru (McGill University) · Etienne Denis (McGill) · Will Hamilton (McGill University and Mila)

Logarithmic Regret for Online Control with Adversarial Noise
Dylan Foster (MIT) · Max Simchowitz (UC Berkeley)

Multiresolution Tensor Learning for Efficient and Interpretable Spatial Analysis
Jung Yeon Park (Northeastern University) · Kenneth Carr (Northeastern University) · Stephan Zheng (Salesforce) · Yisong Yue (Caltech) · Rose Yu (Northeastern University)

Customizing ML Predictions for Online Algorithms
Keerti Anand (Duke University) · Rong Ge (Duke University) · Debmalya Panigrahi (Duke University)

Maximum Entropy Gain Exploration for Long Horizon Multi-goal Reinforcement Learning
Silviu Pitis (University of Toronto) · Harris Chan (University of Toronto, Vector Institute) · Stephen Zhao (University of Toronto) · Bradly Stadie (Vector Institute) · Jimmy Ba (University of Toronto)

Recht-Re Noncommutative Arithmetic-Geometric Mean Conjecture is False
Zehua Lai (University of Chicago) · Lek-Heng Lim (University of Chicago)

Predictive Multiplicity in Classification
Charles Marx (Haverford College) · Flavio Calmon (Harvard University) · Berk Ustun (Harvard University)

Word-Level Speech Recognition With a Letter to Word Encoder
Ronan Collobert (Facebook AI Research) · Awni Hannun (Facebook AI Research) · Gabriel Synnaeve (Facebook AI Research)

Reducing Sampling Error in Batch Temporal Difference Learning
Brahma Pavse (University of Texas at Austin) · Ishan Durugkar (University of Texas at Austin) · Josiah Hanna (University of Edinburgh) · Peter Stone (University of Texas at Austin)

Adaptive Sampling for Estimating Probability Distributions
Shubhanshu Shekhar (University of California, San Diego) · Tara Javidi (University of California San Diego) · Mohammad Ghavamzadeh (Facebook AI Research)

Adversarial Filters of Dataset Biases
Ronan Le Bras (Allen Institute for AI) · Swabha Swayamdipta (Allen Institute for AI) · Chandra Bhagavatula (AllenAI) · Rowan Zellers (University of Washington) · Matthew Peters (AI2) · Ashish Sabharwal (Allen Institute for AI) · Yejin Choi (University of Washington)

Black-Box Variational Inference as a Parametric Approximation to Langevin Dynamics
Matthew Hoffman (Google) · Yian Ma (Google)

Faster Graph Embeddings via Coarsening
Matthew Fahrbach (Georgia Institute of Technology) · Gramoz Goranci (University of Toronto) · Sushant Sachdeva (University of Toronto) · Richard Peng (Georgia Tech / MSR Redmond) · Chi Wang (Microsoft Research)

Efficient non-conjugate Gaussian process factor models for spike countdata using polynomial approximations
Stephen Keeley (Princeton University) · David Zoltowski (Princeton University) · Jonathan Pillow (Princeton University) · Spencer Smith (UC Santa Barbara) · Yiyi Yu (UNC)

Multigrid Neural Memory
Tri Huynh (The University of Chicago) · Michael Maire (University of Chicago) · Matthew Walter (Toyota Technological Institute at Chicago)

Cautious Adaptation For Reinforcement Learning in Safety-Critical Settings
Jesse Zhang (UC Berkeley) · Brian Cheung (UC Berkeley) · Chelsea Finn (Stanford) · Sergey Levine (UC Berkeley) · Dinesh Jayaraman (University of Pennsylvania)

Adversarial Nonnegative Matrix Factorization
lei luo (pitt) · yanfu Zhang (University of Pittsburgh) · Heng Huang (University of Pittsburgh)

Aligned Cross Entropy for Non-Autoregressive Machine Translation
Marjan Ghazvininejad (Facebook AI Research) · Vladimir Karpukhin (Facebook AI Research) · Luke Zettlemoyer (Facebook) · Omer Levy (Facebook)

Model-Agnostic Characterization of Fairness Trade-offs
Joon Sik Kim (Carnegie Mellon University) · Jiahao Chen (JPMorgan Chase) · Ameet Talwalkar (Carnegie Mellon University)

A Distributional Framework For Data Valuation
Amirata Ghorbani (Stanford) · Michael Kim (Stanford University) · James Zou (Stanford University)

Supervised Quantile Normalization for Low Rank Matrix Factorization
Marco Cuturi (Google and CREST/ENSAE) · Olivier Teboul (Google Brain) · Jonathan Niles-Weed (NYU) · Jonathan Weed (NYU) · Jean-Philippe Vert (Google)

AR-DAE: Towards Unbiased Neural Entropy Gradient Estimation
Jae Hyun Lim (Mila) · Aaron Courville (Université de Montréal) · Christopher Pal (École Polytechnique de Montréal) · Chin-Wei Huang (MILA)

Bridging the Gap Between f-GANs and Wasserstein GANs
Jiaming Song (Stanford) · Stefano Ermon (Stanford University)

“Other-Play” for Zero-Shot Coordination
Hengyuan Hu (FAIR) · Alexander Peysakhovich (Facebook) · Adam Lerer (Facebook AI Research) · Jakob Foerster (Facebook AI Research)

Correlation Clustering with Asymmetric Classification Errors
Jafar Jafarov (University of Chicago) · Sanchit Kalhan (Northwestern University) · Konstantin Makarychev (Northwestern University) · Yury Makarychev (TTIC)

An Optimistic Perspective on Offline Deep Reinforcement Learning
Rishabh Agarwal (Google Research, Brain Team) · Dale Schuurmans (Google / University of Alberta) · Mohammad Norouzi (Google Brain)

Neural Topic Modeling with Continual Lifelong Learning
Pankaj Gupta (Siemens AG) · Yatin Chaudhary (Siemens) · Thomas Runkler (Technical University of Munich) · Hinrich Schuetze (University of Munich (LMU))

Learning and Evaluating Contextual Embedding of Source Code
Aditya Kanade (Indian Institute of Science and Google Brain) · Petros Maniatis (Google Brain) · Gogul Balakrishnan (Google) · Kensen Shi (Google)

Uncertainty quantification for nonconvex tensor completion: Confidence intervals, heteroscedasticity and optimality
Changxiao Cai (Princeton University) · H. Vincent Poor (Princeton University) · Yuxin Chen (Princeton University)

Learning with Good Feature Representations in Bandits and in RL with a Generative Model
Gellért Weisz (DeepMind) · Tor Lattimore (DeepMind) · Csaba Szepesvari (DeepMind/University of Alberta)

Angular Visual Hardness
Beidi Chen (Rice University) · Weiyang Liu (Georgia Tech) · Zhiding Yu (NVIDIA) · Jan Kautz (NVIDIA) · Anshumali Shrivastava (Rice University) · Animesh Garg (University of Toronto, Vector Institute, Nvidia) · Anima Anandkumar (Caltech)

Cutting out the Middle-Man: Training and Evaluating Energy-Based Models without Sampling
Will Grathwohl (University of Toronto) · Kuan-Chieh Wang (University of Toronto) · Joern-Henrik Jacobsen (Vector Institute and University of Toronto) · David Duvenaud (University of Toronto) · Richard Zemel (Vector Institute)

Variance Reduction and Quasi-Newton for Particle-Based Variational Inference
Michael Zhu (Stanford University) · Chang Liu (Microsoft Research) · Jun Zhu (Tsinghua University)

Better depth-width trade-offs for neural networks through the lens of dynamical systems
Evangelos Chatziafratis (Stanford University) · Ioannis Panageas (Singapore University of Technology and Design) · Sai Ganesh Nagarajan (SUTD)

Stochastic Coordinate Minimization with Progressive Precision for Stochastic Convex Optimization
Sudeep Salgia (Cornell University) · Qing Zhao (Cornell University) · Sattar Vakili (Cornell University)

Fundamental Tradeoffs between Invariance and Sensitivity to Adversarial Perturbations
Florian Tramer (Stanford University) · Jens Behrmann (University of Bremen) · Nicholas Carlini (Google) · Nicolas Papernot (University of Toronto and Vector Institute) · Joern-Henrik Jacobsen (Vector Institute and University of Toronto)

Learning From Strategic Agents: Accuracy, Improvement, and Causality
Yonadav Shavit (Harvard University) · Benjamin Edelman (Harvard University) · Brian Axelrod (Stanford)

Causal Structure Discovery from Distributions Arising from Mixtures of DAGs
Basil Saeed (Massachusetts Institute of Technology) · Snigdha Panigrahi (University of Michigan) · Caroline Uhler (Massachusetts Institute of Technology)

Explainable and Discourse Topic-aware Neural Language Understanding
Yatin Chaudhary (Siemens) · Pankaj Gupta (Siemens AG) · Hinrich Schuetze (University of Munich (LMU))

Understanding Contrastive Representation Learning through Geometry on the Hypersphere
Tongzhou Wang (MIT) · Phillip Isola (MIT)

On Learning Language-Invariant Representations for Universal Machine Translation
Han Zhao (Carnegie Mellon University) · Junjie Hu (Carnegie Mellon University) · Andrej Risteski (CMU)

Compressive sensing with un-trained neural networks: Gradient descent finds a smooth approximation
Reinhard Heckel (Rice University) · Mahdi Soltanolkotabi (University of Southern California)

Representing Unordered Data Using Multiset Automata and Complex Numbers
Justin DeBenedetto (University of Notre Dame) · David Chiang (University of Notre Dame)

Mutual Transfer Learning for Massive Data
Ching-Wei Cheng (Purdue University) · Xingye Qiao (Binghamton University) · Guang Cheng (Purdue University)

The Differentiable Cross-Entropy Method
Brandon Amos (Facebook AI Research) · Denis Yarats (New York University)

A Sample Complexity Separation between Non-Convex and Convex Meta-Learning
Nikunj Umesh Saunshi (Princeton University) · Yi Zhang (Princeton University) · Mikhail Khodak (CMU) · Sanjeev Arora (Princeton University and Institute for Advanced Study)

On the Convergence of Nesterov's Accelerated Gradient Method in Stochastic Settings
Mahmoud Assran (McGill University; Mila; Facebook AI Research) · Michael Rabbat (Facebook)

The Buckley-Osthus model and the block preferential attachment model: statistical analysis and application
Wenpin Tang (UC Berkeley) · Xin Guo (University of California, Berkeley) · Fengmin Tang (UCLA)

Representations for Stable Off-Policy Reinforcement Learning
Dibya Ghosh (Google) · Marc Bellemare (Google Brain)

Piecewise Linear Regression via a Difference of Convex Functions
Ali Siahkamari (Boston University) · Aditya Gangrade (Boston University) · Brian Kulis (Boston University) · Venkatesh Saligrama (Boston University)

On the consistency of top-k surrogate losses
Forest Yang (Google Brain) · Sanmi Koyejo (Illinois / Google)

Collapsed Amortized Variational Inference for Switching Nonlinear Dynamical Systems
Zhe Dong (Google) · Bryan Seybold (Google) · Kevin Murphy (Google Brain) · Hung Bui (VinAI Research)

Boosting Deep Neural Network Efficiency with Dual-Module Inference
Liu Liu (University of California, Santa Barbara) · Lei Deng (University of California, Santa Barbara) · Zhaodong Chen (University of California, Santa Barbara) · yuke wang (ucsb) · Shuangchen Li (Alibaba Inc.) · Jingwei Zhang (Alibaba Inc.) · Yihua Yang (Alibaba Inc.) · Zhenyu Gu (Alibaba Inc.) · Yufei Ding (University of California, Santa Barbara) · Yuan Xie (University of California, Santa Barbara)

Time-Consistent Semi-Supervised Learning
Tianyi Zhou (University of Washington) · Shengjie Wang (University of Washington) · Jeff Bilmes (UW)

Selective Dyna-style Planning Under Limited Model Capacity
Muhammad Zaheer (University of Alberta) · Samuel Sokota (University of Alberta) · Erin Talvitie () · Martha White (University of Alberta)

A Pairwise Fair and Community-preserving Approach to k-Center Clustering
Brian Brubach (University of Maryland) · Darshan Chakrabarti (Carnegie Mellon University) · John P Dickerson (University of Maryland) · Samir Khuller (Northwestern University) · Aravind Srinivasan (University of Maryland College Park) · Leonidas Tsepenekas (University of Maryland, College Park)

How recurrent networks implement contextual processing in sentiment analysis
Niru Maheswaranathan (Google Brain) · David Sussillo (Google Brain, Google Inc.)

Smaller, more accurate regression forests using tree alternating optimization
Arman Zharmagambetov (UC Merced) · Miguel Carreira-Perpinan (University of California, Merced)

Divide and Conquer: Leveraging Intermediate Feature Representations for Quantized Training of Neural Networks
Ahmed T. Elthakeb (University of California, San Diego) · Prannoy Pilligundla (University of California, San Diego) · FatemehSadat Mireshghallah (University of California San Diego) · Alexander Cloninger (University of California San Diego) · Hadi Esmaeilzadeh (University of California, San Diego)

From Sets to Multisets: Provable Variational Inference for Probabilistic Integer Submodular Models
Aytunc Sahin (ETH Zurich) · Yatao Bian (Tencent AI Lab) · Joachim Buhmann (ETH Zurich) · Andreas Krause (ETH Zurich)

Empirical Study of the Benefits of Overparameterization in Learning Latent Variable Models
Rares-Darius Buhai (MIT) · Yoni Halpern (Google) · Yoon Kim (Harvard University) · Andrej Risteski (CMU) · David Sontag (Massachusetts Institute of Technology)

Improving the Gating Mechanism of Recurrent Neural Networks
Albert Gu (Stanford University) · Caglar Gulcehre (DeepMind) · Thomas Paine (DeepMind) · Matthew Hoffman (DeepMind) · Razvan Pascanu (DeepMind)

Efficient and Scalable Bayesian Neural Nets with Rank-1 Factors
Mike Dusenberry (Google Brain (AI Residency)) · Ghassen Jerfel (Google Brain) · Yeming Wen (University of Toronto) · Yian Ma (Google) · Jasper Snoek (Google Brain) · Katherine Heller (Google) · Balaji Lakshminarayanan (Google DeepMind) · Dustin Tran (Google)

Analyzing the effect of neural network architecture on training performance
Karthik Abinav Sankararaman (Facebook) · Soham De (DeepMind) · Zheng Xu (University of Maryland) · W. Ronny Huang (University of Maryland and EY LLP) · Tom Goldstein (University of Maryland)

Born-again Tree Ensembles
Thibaut Vidal (Pontifical Catholic University of Rio de Janeiro) · Maximilian Schiffer (TUM School of Management, Technical University of Munich)

Accountable Off-Policy Evaluation via a Kernelized Bellman Statistics
Yihao Feng (The University of Texas at Austin) · Tongzheng Ren (UT Austin) · Ziyang Tang (University of Texas at Austin) · Qiang Liu (UT Austin)

Improving Transformer Optimization Through Better Initialization
Xiao Shi Huang (Layer6 AI) · Juan Perez (Layer6 AI) · Jimmy Ba (University of Toronto) · Maksims Volkovs (Layer6 AI)

Learning to Simulate and Design for Structural Engineering
Kai-Hung Chang (Autodesk Research) · Chin-Yi Cheng (Autodesk Research)

Few-shot Relation Extraction via Bayesian Meta-learning on Task Graphs
Meng Qu (MILA) · Tianyu Gao (Tsinghua University) · Louis-Pascal Xhonneux (Mila / Université de Montréal) · Jian Tang (HEC Montreal & MILA)

Optimal Differential Privacy Composition for Exponential Mechanisms
Jinshuo Dong (University of Pennsylvania) · David Durfee (Georgia Tech) · Ryan Rogers (LinkedIn)

Scaling up Hybrid Probabilistic Inference with Logical and Arithmetic Constraints via Message Passing
Zhe Zeng (University of California, Los Angeles) · Paolo Morettin (University of Trento) · Fanqi Yan (UCAS) · Antonio Vergari (University of California, Los Angeles) · Guy Van den Broeck (University of California, Los Angeles)

Accelerating Large-Scale Inference with Anisotropic Vector Quantization
Ruiqi Guo (Google Research) · Quan Geng (Google) · David Simcha (Google) · Felix Chern (Google AI) · Philip Sun (Google) · Erik Lindgren (Google Research) · Sanjiv Kumar (Google Research, NY)

Convolutional dictionary learning based auto-encoders for natural exponential-family distributions
Bahareh Tolooshams (Harvard University) · Andrew Song (MIT) · Simona Temereanca (Brown University) · Demba Ba (Harvard)

Strength from Weakness: Fast Learning Using Weak Supervision
Joshua Robinson (MIT) · Stefanie Jegelka (MIT) · Suvrit Sra (MIT)

NADS: Neural Architecture Distribution Search for Uncertainty Awareness
Randy Ardywibowo (Texas A&M University) · Shahin Boluki (Texas A&M University) · Xinyu Gong (Texas A&M University) · Zhangyang Wang (Texas A&M University) · Xiaoning Qian (Texas A&M University)

Approximating Stacked and Bidirectional Recurrent Architectures with the Delayed Recurrent Neural Network
Javier Turek (Intel Labs) · Shailee Jain (The University of Texas at Austin) · Vy Vo (Intel Labs) · Mihai Capotă (Intel Labs) · Alexander Huth (The University of Texas at Austin) · Theodore Willke (Intel Labs)

Balancing Competing Objectives with Noisy Data: Score-Based Classifiers for Welfare-Aware Machine Learning
Esther Rolf (UC Berkeley) · Max Simchowitz (UC Berkeley) · Sarah Dean (UC Berkeley) · Lydia T. Liu (University of California Berkeley) · Daniel Bjorkegren (Brown University) · University of California Moritz Hardt (University of California, Berkeley) · Joshua Blumenstock (University of California, Berkeley)

Time-aware Large Kernel Convolutions
Vasileios Lioutas (Carleton University) · Yuhong Guo (Carleton University)

Amortised Learning by Wake-Sleep
Li Kevin Wenliang (Gatsby Unit, University College London) · Theodore Moskovitz (Gatsby Computational Neuroscience Unit) · Heishiro Kanagawa (Gatsby Unit, UCL) · Maneesh Sahani (Gatsby Unit, UCL)

Fair Generative Modeling via Weak Supervision
Kristy Choi (Stanford University) · Aditya Grover (Stanford University) · Trisha Singh (Stanford University) · Rui Shu (Stanford University) · Stefano Ermon (Stanford University)

Multi-Step Greedy Reinforcement Learning Algorithms
Manan Tomar (Indian Institute of Technology, Madras) · Yonathan Efroni (Technion) · Mohammad Ghavamzadeh (Facebook AI Research)

Linear Mode Connectivity and the Lottery Ticket Hypothesis
Jonathan Frankle (MIT CSAIL) · Gintare Karolina Dziugaite (Element AI) · Daniel Roy (Univ of Toronto | Toronto) · Michael Carbin (MIT)

Superpolynomial Lower Bounds for Learning One-Layer Neural Networks using Gradient Descent
Surbhi Goel (UT Austin) · Aravind Gollakota (University of Texas at Austin) · Zhihan Jin (Shanghai Jiao Tong University) · Sushrut Karmalkar (University of Texas at Austin) · Adam Klivans (University of Texas at Austin)

Learnable Group Transform For Time-Series
Romain Cosentino (Rice University) · Behnaam Aazhang (Rice University)

Optimistic bounds for multi-output learning
Henry Reeve (University of Birmingham) · Ata Kaban (University of Birmingham)

Detecting Out-of-Distribution Examples with Gram Matrices
Chandramouli Shama Sastry (Dalhousie University/Vector Institute) · Sageev Oore (Dalhousie University)

On Variational Learning of Controllable Representations for Text without Supervision
Peng Xu (Borealis AI) · Jackie Chi Kit Cheung (McGill University / Mila) · Yanshuai Cao (Borealis AI)

Model-Based Reinforcement Learning with Value-Targeted Regression
Zeyu Jia (Peking University) · Lin Yang (UCLA) · Csaba Szepesvari (DeepMind/University of Alberta) · Mengdi Wang (Princeton University) · Alex Ayoub (University of Alberta)

Robust and scalable credit assignment without weight symmetry
Daniel Kunin (Stanford University) · Aran Nayebi (Stanford University) · Javier Sagastuy-Brena (Stanford University) · Surya Ganguli (Stanford) · Jonathan Bloom (Broad Institute of MIT and Harvard) · Daniel Yamins (Stanford University)

Predicting deliberative outcomes
Vikas K Garg (Massachusetts Institute of Technology) · Tommi Jaakkola (MIT)

Black-box Certification and Learning under Adversarial Perturbations
Hassan Ashtiani (McMaster University) · Vinayak Pathak (Scotiabank) · Ruth Urner (York University)

When deep denoising meets iterative phase retrieval
Yaotian Wang (Princeton University) · Xiaohang Sun (Princeton University) · Jason Fleischer (Princeton University)

The Neural Tangent Kernel in High Dimensions: Triple Descent and a Multi-Scale Theory of Generalization
Ben Adlam (Google) · Jeffrey Pennington (Google Brain)

A Sequential Self Teaching Approach for Improving Generalization in Sound Event Recognition
Anurag Kumar (Facebook Reality Labs) · Vamsi Ithapu (Univresity of Wisconsin Madiso)

On the Global Convergence Rates of Softmax Policy Gradient Methods
Jincheng Mei (Google / University of Alberta) · Chenjun Xiao (Google / University of Alberta) · Csaba Szepesvari (DeepMind/University of Alberta) · Dale Schuurmans (University of Alberta)

Source Separation with Deep Generative Priors
Vivek Jayaram (University of Washington) · John Thickstun (University of Washington)

Non-Autoregressive Neural Text-to-Speech
Kainan Peng (Baidu Research) · Wei Ping (Baidu Research) · Zhao Song (Baidu Research) · Kexin Zhao (Baidu)

Amortized Population Gibbs Samplers with Neural Sufficient Statistics
Hao Wu (Northeastern University) · Heiko Zimmermann (Northeastern University) · Eli Sennesh (Northeastern University) · Tuan Anh Le (MIT) · Jan-Willem van de Meent (Northeastern University)

Neural Network Control Policy Verification With Persistent Adversarial Perturbation
Yuh-Shyang Wang (Argo AI) · Tsui-Wei Weng (MIT) · Luca Daniel (MIT)

Circuit-Based Intrinsic Methods to Detect Overfitting
Satrajit Chatterjee (Google) · Alan Mishchenko (UC Berkeley)

Inter-domain Deep Gaussian Processes with RKHS Fourier Features
Tim Rudner (University of Oxford) · Dino Sejdinovic (University of Oxford) · Yarin Gal (University of Oxford)

Estimating Q(s,s') with Deterministic Dynamics Gradients
Ashley Edwards (Uber AI) · Himanshu Sahni (Georgia Institute of Technology) · Rosanne Liu (Deep Collective) · Jane Hung (Uber) · Ankit Jain (Uber AI Labs) · Rui Wang (Uber AI) · Adrien Ecoffet (Uber AI) · Thomas Miconi (Uber AI Labs) · Charles Isbell (Georgia Institute of Technology) · Jason Yosinski (Uber Labs)

On conditional versus marginal bias in multi-armed bandits
Jaehyeok Shin (Carnegie Mellon University) · Aaditya Ramdas (Carnegie Mellon University) · Alessandro Rinaldo (Carnegie Mellon University)

Implicit competitive regularization in GANs
Florian Schaefer (Caltech) · Hongkai Zheng (Shanghai Jiao Tong University) · Anima Anandkumar (Caltech)

DrRepair: A Self-Supervised, Graph-Attentional Approach to Repairing Programs from Diagnostic Feedback
Michihiro Yasunaga (Stanford University) · Percy Liang (Stanford University)

Interpretable Off-Policy Evaluation in Reinforcement Learning by Highlighting Influential Transitions
Omer Gottesman (Harvard University) · Joseph Futoma (Harvard University) · Yao Liu (Stanford University) · Sonali Parbhoo (Harvard University) · Leo Celi (MIT) · Emma Brunskill (Stanford University) · Finale Doshi-Velez (Harvard University)

Communication-Efficient Federated Learning with Sketching
Daniel Rothchild (UC Berkeley) · Ashwinee Panda (UC Berkeley) · Enayat Ullah (Johns Hopkins University) · Nikita Ivkin (Amazon) · Vladimir Braverman (Johns Hopkins University) · Joseph Gonzalez (UC Berkeley) · Ion Stoica (UC Berkeley) · Raman Arora (Johns Hopkins University)

Learning Fair Policies in Multi-Objective (Deep) Reinforcement Learning with Average and Discounted Rewards
Umer Siddique (Shanghai Jiao Tong University) · Paul Weng (Shanghai Jiao Tong University) · Matthieu Zimmer (UM-SJTU JI)

Robust Black Box Explanations Under Distribution Shift
Himabindu Lakkaraju (Harvard) · Nino Arsov (Macedonian Academy of Arts and Sciences) · Osbert Bastani (University of Pennsylvania)

Distributed Online Optimization over a Heterogeneous Network
Nima Eshraghi (University of Toronto) · Ben Liang (University of Toronto)

ECLIPSE: An Extreme-Scale Linear Program Solver for Web-Applications
Kinjal Basu (LinkedIn Corporation) · Amol Ghoting (LinkedIn) · Rahul Mazumder (Massachusetts Institute of Technology) · Yao Pan (LinkedIn Corporation)

CURL: Contrastive Unsupervised Representation Learning for Reinforcement Learning
Michael Laskin (UC Berkeley) · Pieter Abbeel (UC Berkeley & Covariant) · Aravind Srinivas (UC Berkeley)

Confidence-Aware Learning for Deep Neural Networks
Sangheum Hwang (Seoul National University of Science and Technology) · Jooyoung Moon (Seoul National University of Science and Technology) · Jihyo Kim (Seoul National University of Science and Technology) · Younghak Shin (LGCNS)

Online Bayesian Moment Matching based SAT Solver Heuristics
Haonan Duan (University of Waterloo) · Saeed Nejati (University of Waterloo) · George Trimponias (Noah's Ark Lab) · Pascal Poupart (University of Waterloo and Borealis AI) · Vijay Ganesh (University of Waterloo, Electrical and Computer Engineering)

Retro*: Learning Retrosynthetic Planning with Neural Guided A* Search
Binghong Chen (Georgia Tech) · Chengtao Li (Galixir) · Hanjun Dai (Google Brain) · Le Song (Georgia Institute of Technology)

FedBoost: A Communication-Efficient Algorithm for Federated Learning
Jenny Hamer (Google Research) · Mehryar Mohri (Google Research and Courant Institute of Mathematical Sciences) · Ananda Theertha Suresh (Google Research)

Sharp Composition Bounds for Gaussian Differential Privacy via Edgeworth Expansion
Qinqing Zheng (Facebook) · Jinshuo Dong (University of Pennsylvania) · Qi Long (University of Pennsylvania) · Weijie Su (University of Pennsylvania)

Few-Shot Learning as Domain Adaptation: Algorithm and Analysis
Jiechao Guan (Renmin University of China) · Zhiwu Lu (Renmin University of China) · Tao Xiang (University of Surrey) · Ji-Rong Wen (Renmin University of China)

Fast and Three-rious: Speeding Up Weak Supervision with Triplet Methods
Daniel Fu (Stanford University) · Mayee Chen (Stanford University) · Frederic Sala (Stanford) · Sarah Hooper (Stanford University) · Kayvon Fatahalian (Stanford) · Christopher Re (Stanford)

Spectral Frank-Wolfe Algorithm: Strict Complementarity and Linear Convergence
Lijun Ding (Cornell University) · Yingjie Fei (Cornell University) · Qiantong Xu (Facebook) · Chengrun Yang (Cornell University)

Deep Molecular Programming: A Natural Implementation of Binary-Weight ReLU Neural Networks
Marko Vasic (The University of Texas at Austin) · Cameron Chalk (University of Texas at Austin) · Sarfraz Khurshid () · David Soloveichik (The University of Texas at Austin)

Generative Pretraining From Pixels
Mark Chen (OpenAI) · Alec Radford (OpenAI) · Rewon Child (OpenAI) · Jeffrey K Wu (OpenAI) · Heewoo Jun (OpenAI) · David Luan (OpenAI) · Ilya Sutskever (OpenAI)

Inferring DQN structure for high-dimensional continuous control
Andrey Sakryukin (National University of Singapore) · Chedy Raissi (INRIA) · Mohan Kankanhalli (National University of Singapore,)

Subspace Fitting Meets Regression: The Effects of Supervision and Orthonormality Constraints on Double Descent of Generalization Errors
Yehuda Dar (Rice University) · Paul Mayer (Rice University) · Lorenzo Luzi (Rice University) · Richard Baraniuk (OpenStax / Rice University)

Learning Selection Strategies in Buchberger's Algorithm
Dylan Peifer (Cornell University) · Michael Stillman (Cornell University) · Daniel Halpern-Leistner (Cornell University)

Estimating the Error of Randomized Newton Methods: A Bootstrap Approach
Miles Lopes (University of California, Davis) · Xiaotie Chen (UC Davis)

Spectral Subsampling MCMC for Stationary Time Series
Robert Salomone (University of New South Wales) · Matias Quiroz (University of Technology Sydney) · Robert kohn (UNSW) · Mattias Villani (Linkoeping University) · Minh-Ngoc Tran (U of Sydney)

Progressive Identification of True Labels for Partial-Label Learning
Jiaqi Lv (Southeast University) · Miao Xu (University of Queensland/ RIKEN AIP) · LEI FENG (Nanyang Technological University) · Gang Niu (RIKEN) · Xin Geng (Southeast University) · Masashi Sugiyama (RIKEN / The University of Tokyo)

R2-B2: Recursive Reasoning-Based Bayesian Optimization for No-Regret Learning in Games
Zhongxiang Dai (National University of Singapore) · Yizhou Chen (National University of Singapore) · Bryan Kian Hsiang Low (National University of Singapore) · Patrick Jaillet (MIT) · Teck-Hua Ho (National University of Singapore)

Graph Homomorphism Convolution
Hoang Nguyen (RIKEN AIP) · Takanori Maehara (RIKEN AIP)

Conditional Augmentation for Generative Modeling
Heewoo Jun (OpenAI) · Rewon Child (OpenAI) · Mark Chen (OpenAI) · John Schulman (OpenAI) · Aditya Ramesh (OpenAI) · Alec Radford (OpenAI) · Ilya Sutskever (OpenAI)

PDO-eConvs: Partial Differential Operator Based Equivariant Convolutions
Zhengyang Shen (Peking University) · Lingshen He (Peking University) · Zhouchen Lin (Peking University) · Jinwen Ma (Peking University)

Abstraction Mechanisms Predict Generalization in Deep Neural Networks
Alex Gain (Johns Hopkins University) · Hava Siegelmann (UMass Amherst; DARPA)

Revisiting Fundamentals of Experience Replay
William Fedus (University of Montreal/Google Brain) · Prajit Ramachandran (Google) · Rishabh Agarwal (Google Research, Brain Team) · Yoshua Bengio (Mila / U. Montreal) · Hugo Larochelle (Google Brain) · Mark Rowland (DeepMind) · Will Dabney (DeepMind)

Go Wide, Then Narrow: Efficient Training of Deep Thin Networks
Denny Zhou (Google Brain) · Mao Ye (PURDUE UNIVERSITY) · Chen Chen (Google) · Mingxing Tan (Google Brain) · Tianjian Meng (Google Brain) · Xiaodan Song (Google Brain) · Quoc Le (Google Brain) · Qiang Liu (UT Austin) · Dale Schuurmans (Google / University of Alberta)

Meta-learning for mixed linear regression
Weihao Kong (Stanford University) · Raghav Somani (University of Washington) · Zhao Song (UT-Austin & University of Washington) · Sham Kakade (University of Washington) · Sewoong Oh (University of Washington)

Efficiently Learning Adversarially Robust Halfspaces with Noise
Omar Montasser (TTI-Chicago) · Surbhi Goel (UT Austin) · Ilias Diakonikolas (USC) · Nati Srebro (Toyota Technological Institute at Chicago)

Bayesian Graph Neural Networks with Adaptive Connection Sampling
Arman Hasanzadeh (Texas A&M University) · Ehsan Hajiramezanali (Texas A&M University) · Shahin Boluki (Texas A&M University) · Nick Duffield (Texas A&M University) · Mingyuan Zhou (University of Texas at Austin) · Krishna Narayanan (Texas A&M University) · Xiaoning Qian (Texas A&M University)

On the Theoretical Properties of the Network Jackknife
Qiaohui Lin (University of Texas at Austin) · Robert Lunde (University of Texas at Austin) · Purnamrita Sarkar (UT Austin)

Thompson Sampling via Local Uncertainty
Zhendong Wang (University of Texas, Austin) · Mingyuan Zhou (University of Texas at Austin)

Decision Trees for Decision-Making under the Predict-then-Optimize Framework
Adam Elmachtoub (Columbia University) · Jason Cheuk Nam Liang (MIT) · Ryan McNellis (Amazon)

Representation Learning via Adversarially-Contrastive Optimal Transport
Anoop Cherian (MERL) · Shuchin Aeron (Tufts University)

Neuro-Symbolic Visual Reasoning: Disentangling "Visual" from "Reasoning"
Saeed Amizadeh (Microsoft) · Hamid Palangi (Microsoft Research) · Oleksandr Polozov (Microsoft Research) · Yichen Huang (Microsoft Research) · Kazuhito Koishida (Microsoft)

Two Simple Ways to Learn Individual Fairness Metric from Data
Debarghya Mukherjee (University of Michigan) · Mikhail Yurochkin (IBM Research AI) · Moulinath Banerjee (University of Michigan) · Yuekai Sun (University of Michigan)

A Simple Framework for Contrastive Learning of Visual Representations
Ting Chen (Google) · Simon Kornblith (Google Brain) · Mohammad Norouzi (Google Brain) · Geoffrey Hinton (Google)

The Implicit and Explicit Regularization Effects of Dropout
Colin Wei (Stanford University) · Sham Kakade (University of Washington) · Tengyu Ma (Stanford)

Variable-Bitrate Neural Compression via Bayesian Arithmetic Coding
Yibo Yang (University of California, Irivine) · Robert Bamler (University of California at Irvine) · Stephan Mandt (University of California, Irivine)

Orthogonalized SGD and Nested Architectures for Anytime Neural Networks
Chengcheng Wan (University of Chicago) · Henry (Hank) Hoffmann (The University of Chicago) · Shan Lu (University of Chicago) · Michael Maire (University of Chicago)

Evaluating Machine Accuracy on ImageNet
Vaishaal Shankar (UC Berkeley) · Rebecca Roelofs (Google) · Horia Mania (UC Berkeley) · Alex Fang (UC Berkeley) · Benjamin Recht (Berkeley) · Ludwig Schmidt (University of California, Berkeley)

Learning to Navigate in Synthetically Accessible Chemical Space Using Reinforcement Learning
Sai Krishna Gottipati (99andBeyond) · Boris Sattarov (99andBeyond) · Sufeng Niu (Linkedin) · Haoran Wei (University of Delaware) · Yashaswi Pathak (International Institute of Information Technology,Hyderabad) · Shengchao Liu (MILA-UdeM) · Shengchao Liu (Mila, Université de Montréal) · Simon Blackburn (Mila) · Karam Thomas (99andBeyond) · Connor Coley (MIT) · Jian Tang (HEC Montreal & MILA) · Sarath Chandar (Mila / University of Montreal) · Yoshua Bengio (Mila / U. Montreal)

Improved Bounds on Minimax Regret under Logarithmic Loss via Self-Concordance
Blair Bilodeau (University of Toronto) · Dylan Foster (MIT) · Daniel Roy (Univ of Toronto | Toronto)

Optimization Theory for ReLU Neural Networks Trained with Normalization Layers
Yonatan Dukler (UCLA) · Guido Montufar (UCLA) · Quanquan Gu (University of California, Los Angeles)

Improving Molecular Design by Stochastic Iterative Target Augmentation
Kevin Yang (UC Berkeley) · Wengong Jin (MIT) · Kyle Swanson (University of Cambridge) · Regina Barzilay (MIT CSAIL) · Tommi Jaakkola (MIT)

Don't Waste Your Bits! Squeeze Activations and Gradients for Deep Neural Networks via TinyScript
Fangcheng Fu (Peking University) · Yuzheng Hu (Peking University) · Yihan He (Peking University) · Jiawei Jiang (ETH Zurich) · Yingxia Shao (BUPT) · Ce Zhang (ETH Zurich) · Bin Cui (Peking University)

Robust One-Bit Recovery via ReLU Generative Networks: Near-Optimal Statistical Rate and Global Landscape Analysis
Shuang Qiu (University of Michigan) · Xiaohan Wei (University of Southern California) · Zhuoran Yang (Princeton University)

Multi-objective Bayesian Optimization using Pareto-frontier Entropy
Shinya Suzuki (Nagoya Institute of Technology) · shion takeno (Nagoya Institute of Technology) · Tomoyuki Tamura (National Institute for Material Science) · Kazuki Shitara (Osaka University) · Masayuki Karasuyama (Nagoya Institute of Technology)

Closing the convergence gap of SGD without replacement
Shashank Rajput (University of Wisconsin - Madison) · Anant Gupta (University of Wisconsin Madison) · Dimitris Papailiopoulos (University of Wisconsin-Madison)

Black-Box Methods for Restoring Monotonicity
Evangelia Gergatsouli (UW-Madison) · Brendan Lucier (Microsoft Research New England) · Christos Tzamos (UW-Madison)

Flexible and Efficient Long-Range Planning Through Curious Exploration
Aidan Curtis (Rice University) · Minjian Xin (Shanghai Jiao Tong University) · Dilip Arumugam (Stanford University) · Kevin Feigelis (Stanford University) · Daniel Yamins (Stanford University)

Sparse Convex Optimization via Adaptively Regularized Hard Thresholding
Kyriakos Axiotis (MIT) · Maxim Sviridenko (Yahoo! Research)

On Thompson Sampling with Langevin Algorithms
Eric Mazumdar (University of California Berkeley) · Aldo Pacchiano (UC Berkeley) · Yian Ma (Google) · Michael Jordan (UC Berkeley) · Peter Bartlett (UC Berkeley)

Strategic Classification is Causal Modeling in Disguise
John Miller (University of California, Berkeley) · Smitha Milli (UC Berkeley) · University of California Moritz Hardt (University of California, Berkeley)

Multi-fidelity Bayesian Optimization with Max-value Entropy Search and its Parallelization
shion takeno (Nagoya Institute of Technology) · Hitoshi Fukuoka (Nagoya University) · Yuhki Tsukada (Nagoya University) · Toshiyuki Koyama (Nagoya University) · Motoki Shiga (Gifu University) · Ichiro Takeuchi (Nagoya Institute of Technology / RIKEN) · Masayuki Karasuyama (Nagoya Institute of Technology)

Domain Aggregation Networks for Multi-Source Domain Adaptation
Junfeng Wen (University of Alberta) · Russell Greiner (U Alberta) · Dale Schuurmans (University of Alberta)

Improving Robustness of Deep-Learning-Based Image Reconstruction
Ankit Raj (University of Illinois at Urbana-Champaign) · Yoram Bresler (UIUC) · Bo Li (UIUC)

Outsourced Bayesian Optimization
Dmitrii Kharkovskii (National University of Singapore) · Zhongxiang Dai (National University of Singapore) · Bryan Kian Hsiang Low (National University of Singapore)

Learning Near Optimal Policies with Low Inherent Bellman Error
Andrea Zanette (Stanford University) · Alessandro Lazaric (Facebook AI Research) · Mykel Kochenderfer (Stanford University) · Emma Brunskill (Stanford University)

Message Passing Least Squares: A Unified Framework for Fast and Robust Group Synchronization
Yunpeng Shi (University of Minnesota) · Gilad Lerman (University of Minnesota)

Optimal Estimator for Unlabeled Linear Regression
hang zhang (Gatech Tech) · Ping Li (Baidu)

Recovery of sparse signals from a mixture of linear samples
Arya Mazumdar (University of Massachusetts Amherst) · Soumyabrata Pal (Umass Amherst)

Recurrent Hierarchical Topic-Guided RNN for Language Generation
Dandan Guo (National Laboratory of Radar Signal Processing, Xidian University) · Bo Chen (School of Electronic Engineering, Xidian University) · Ruiying Lu (xidian university) · Mingyuan Zhou (University of Texas at Austin)

Predictive Coding for Locally-Linear Control
Rui Shu (Stanford University) · Tung Nguyen (VinAI Research) · Yinlam Chow (Google) · Tuan Pham (VinAI) · Khoat Than (VinAI & HUST) · Mohammad Ghavamzadeh (Facebook) · Stefano Ermon (Stanford University) · Hung Bui (VinAI Research)

Near Input Sparsity Time Kernel Embeddings via Adaptive Sampling
Amir Zandieh (EPFL) · David Woodruff (CMU)

Near-optimal sample complexity bounds for learning Latent $k-$ polytopes and applications to Ad-Mixtures
Chiranjib Bhattacharyya (Indian Institute of Science) · Ravindran Kannan (Microsoft Research India)

Population-Based Black-Box Optimization for Biological Sequence Design
Christof Angermueller (Google) · David Belanger (Google) · Andreea Gane (Google) · Zelda Mariet (Google Inc.) · David Dohan (Google) · Kevin Murphy (Google Brain) · Lucy Colwell (Google) · D. Sculley (Google)

Emergence of Separable Manifolds in Deep Language Representations
Jonathan Mamou (Intel AI Lab) · Hang Le (MIT) · Miguel Del Rio (MIT) · Cory Stephenson (Intel Corporation) · Hanlin Tang (Intel AI) · Yoon Kim (Harvard University) · Sueyeon Chung (MIT)

Stochastic Hamiltonian Gradient Methods for Smooth Games
Nicolas Loizou (Mila, Université de Montréal) · Hugo Berard (Université de Montreal) · Alexia Jolicoeur-Martineau (Mila) · Pascal Vincent (U Montreal) · Simon Lacoste-Julien (Mila, University of Montreal) · Ioannis Mitliagkas (MILA, UdeM)

Understanding and Estimating the Adaptability of Domain-Invariant Representations
Ching-Yao Chuang (MIT) · Antonio Torralba (MIT) · Stefanie Jegelka (Massachusetts Institute of Technology)

Adversarial Mutual Information for Text Generation
Boyuan Pan (Zhejiang University) · Yazheng Yang (Zhejiang University) · Kaizhao Liang (University of Illinois, Urbana Champaign) · Bhavya Kailkhura (LLNL) · Zhongming Jin (Alibaba Group) · Xian-Sheng Hua (Alibaba Group) · Deng Cai (ZJU) · Bo Li (UIUC)

Bidirectional Model-based Policy Optimization
Hang Lai (Shanghai Jiao Tong University) · Jian Shen (Shanghai Jiao Tong University) · Weinan Zhang (Shanghai Jiao Tong University) · Yong Yu (Shanghai Jiao Tong University)

Input-Sparsity Low Rank Approximation in Schatten Norm
Yi Li (Nanyang Technological University) · David Woodruff (Carnegie Mellon University)

Do We Need Zero Training Loss After Achieving Zero Training Error?
Takashi Ishida (The University of Tokyo / RIKEN) · Ikko Yamane (The University of Tokyo) · Tomoya Sakai (NEC) · Gang Niu (RIKEN) · Masashi Sugiyama (RIKEN / The University of Tokyo)

Learning and sampling of atomic interventions from observations
Arnab Bhattacharyya (National University of Singapore) · Sutanu Gayen (National University of SIngapore) · Saravanan Kandasamy (Cornell University) · Ashwin Maran (Indian Institute of Science) · Vinodchandran N. Variyam (University of Nebraska, Lincoln)

Understanding and Mitigating the Tradeoff between Robustness and Accuracy
Aditi Raghunathan (Stanford) · Sang Michael Xie (Stanford University) · Fanny Yang (ETH) · John Duchi (Stanford University) · Percy Liang (Stanford University)

Combining Differentiable PDE Solvers and Graph Neural Networks for Fluid Flow Prediction
Filipe de Avila Belbute-Peres (Carnegie Mellon University) · Thomas Economon (SU2 Foundation) · Zico Kolter (Carnegie Mellon University / Bosch Center for AI)

From ImageNet to Image Classification: Contextualizing Progress on Benchmarks
Dimitris Tsipras (MIT) · Shibani Santurkar (MIT) · Logan Engstrom (MIT) · Andrew Ilyas (Massachusetts Institute of Technology) · Aleksander Madry (MIT)

On Implicit Regularization in β -VAEs
Abhishek Kumar (Google) · Ben Poole (Google Brain)

Data Amplification: Instance-Optimal Property Estimation
Yi Hao (University of California, San Diego) · Alon Orlitsky (UCSD)

Provable guarantees for decision tree induction: the agnostic setting
Guy Blanc (Stanford University) · Jane Lange (Stanford University) · Li-Yang Tan (Stanford University)

Statistical Bias in Dataset Replication
Logan Engstrom (MIT) · Andrew Ilyas (Massachusetts Institute of Technology) · Shibani Santurkar (MIT) · Dimitris Tsipras (MIT) · Jacob Steinhardt (University of California, Berkeley) · Aleksander Madry (MIT)

Towards Adaptive Residual Network Training: A Neural-ODE Perspective
chengyu dong (UCSD) · Liyuan Liu (University of Illinois at Urbana Champaign) · Zichao Li (University of California, San Diego) · Jingbo Shang (University of California, San Diego)

Overparameterization hurts worst-group accuracy with spurious correlations
Shiori Sagawa (Stanford University) · aditi raghunathan (stanford university) · Pang Wei Koh (Stanford University) · Percy Liang (Stanford University)

A Nearly-Linear Time Algorithm for Exact Community Recovery in Stochastic Block Model
Peng Wang (The Chinese University of Hong Kong) · Zirui Zhou (Hong Kong Baptist University) · Anthony Man-Cho So (The Chinese University of Hong Kong)

Online Multi-Kernel Learning with Graph-Structured Feedback
Pouya Mollaebrahim Ghari (University of California, Irvine) · Yanning Shen (University of California, Irvine)

Is Local SGD Better than Minibatch SGD?
Blake Woodworth (TTI-Chicago) · Kumar Kshitij Patel (Toyota Technological Institute at Chicago) · Sebastian Stich (EPFL) · Zhen Dai (University of Chicago) · Brian Bullins (TTI Chicago) · H. Brendan McMahan (Google) · Ohad Shamir (Weizmann Institute of Science) · Nati Srebro (Toyota Technological Institute at Chicago)

On Lp-norm Robustness of Ensemble Decision Stumps and Trees
Yihan Wang (Tsinghua University) · Huan Zhang (UCLA) · Hongge Chen (MIT) · Duane Boning (MIT) · Cho-Jui Hsieh (UCLA)

Sub-linear Memory Sketches for Near Neighbor Search on Streaming Data with RACE
Benjamin Coleman (Rice University) · Anshumali Shrivastava (Rice University) · Richard Baraniuk (OpenStax / Rice University)

Understanding Self-Training for Gradual Domain Adaptation
Ananya Kumar (Stanford University) · Tengyu Ma (Stanford) · Percy Liang (Stanford University)

Concept Bottleneck Models
Pang Wei Koh (Stanford University) · Thao Nguyen (Google) · Yew Siang Tang (Stanford University) · Stephen Mussmann (Stanford University) · Emma Pierson (Stanford) · Been Kim (Google) · Percy Liang (Stanford University)

Optimal Bounds between f-Divergences and Integral Probability Metrics
Rohit Agrawal (Harvard University) · Thibaut Horel (MIT)

Robustness to Spurious Correlations via Human Annotations
Megha Srivastava (Stanford University) · Tatsunori Hashimoto (Stanford) · Percy Liang (Stanford University)

DROCC: Deep Robust One-Class Classification
Sachin Goyal (Microsoft research) · Aditi Raghunathan (Stanford) · Moksh Jain (NIT Karnataka, Surathkal) · Harsha Vardhan Simhadri (Microsoft Research) · Prateek Jain (Microsoft Research)

Efficiently Solving MDPs with Stochastic Mirror Descent
Yujia Jin (Stanford University) · Aaron Sidford (Stanford)

Handling the Positive-Definite Constraint in the Bayesian Learning Rule
Wu Lin (UBC) · Mark Schmidt (University of British Columbia) · Mohammad Emtiyaz Khan (RIKEN)

A simpler approach to accelerated optimization: iterative averaging meets optimism
Pooria Joulani (DeepMind) · Anant Raj (Max-Planck Institute for Intelligent Systems) · Andras Gyorgy (DeepMind) · Csaba Szepesvari (DeepMind/University of Alberta)

Training Binary Neural Networks using the Bayesian Learning Rule
Xiangming Meng (Riken) · Roman Bachmann (EPFL) · Mohammad Emtiyaz Khan (RIKEN)

High-dimensional Robust Mean Estimation via Gradient Descent
Yu Cheng (University of Illinois at Chicago) · Ilias Diakonikolas (USC) · Rong Ge (Duke University) · Mahdi Soltanolkotabi (University of Southern California)

From Chaos to Order: Symmetry and Conservation Laws in Game Dynamics
Sai Ganesh Nagarajan (SUTD) · David Balduzzi (DeepMind) · Georgios Piliouras (Singapore University of Technology and Design)

Hierarchically Decoupled Morphological Transfer
Donald Hejna (UC Berkeley) · Lerrel Pinto (NYU/Berkeley) · Pieter Abbeel (UC Berkeley)

Puzzle Mix: Exploiting Saliency and Local Statistics for Optimal Mixup
Jang-Hyun Kim (Seoul National University) · Wonho Choo (Seoul National University) · Hyun Oh Song (Seoul National University)

Train Big, Then Compress: Rethinking Model Size for Efficient Training and Inference of Transformers
Zhuohan Li (UC Berkeley) · Eric Wallace (U.C. Berkeley) · Sheng Shen (University of California, Berkeley) · Kevin Lin (UC Berkeley) · Kurt Keutzer (UC Berkeley) · Dan Klein (UC Berkeley) · Joseph E. Gonzalez (UC Berkeley)

Interpolation between CNNs and ResNets
Zonghan Yang (Tsinghua University) · Yang Liu (Tsinghua University) · Chenglong Bao (Tsinghua University) · Zuoqiang Shi (Tsinghua University)

Online metric algorithms with untrusted predictions
Antonios Antoniadis (MPII) · Christian Coester (Centrum Wiskunde & Informatica) · Marek Elias (École polytechnique fédérale de Lausanne) · Adam Polak (Jagiellonian University) · Bertrand Simon (University of Bremen)

Collaborative Machine Learning with Incentive-Aware Model Rewards
Rachael Hwee Ling Sim (National University of Singapore) · Yehong Zhang (National University of Singapore) · Bryan Kian Hsiang Low (National University of Singapore) · Mun Choon Chan (NUS)

On Convergence-Diagnostic based Step Sizes for Stochastic Gradient Descent
Scott Pesme (EPFL) · Aymeric Dieuleveut (École polytechnique) · Nicolas Flammarion (EPFL)

The Performance Analysis of Generalized Margin Maximizers on Separable Data
Fariborz Salehi (California Institute of Technology) · Ehsan Abbasi (California Institute of Technology) · Babak Hassibi (Caltech)

Equivariant Flows: exact likelihood generative learning for symmetric densities.
Jonas Köhler (Freie Universität Berlin) · Leon Klein (Freie Universität Berlin) · Frank Noe (FU Berlin)

PoWER-BERT: Accelerating BERT Inference via Progressive Word-vector Elimination
Saurabh Goyal (IBM Research - India) · Anamitra Roy Choudhury (IBM Research-India) · Venkatesan Chakaravarthy ("IBM Research, India") · Saurabh Raje (IBM Research - India) · Yogish Sabharwal (IBM Research - India) · Ashish Verma (IBM Research)

Bayesian Sparsification of Deep C-valued networks
Ivan Nazarov (Skolkovo Institute of Science and Technology) · Evgeny Burnaev (Skoltech)

Minimally distorted Adversarial Examples with a Fast Adaptive Boundary Attack
Francesco Croce (University of Tuebingen) · Matthias Hein (University of Tübingen)

A distributional view on multi objective policy optimization
Abbas Abdolmaleki (Google DeepMind) · Sandy Huang (DeepMind) · Leonard Hasenclever (DeepMind) · Michael Neunert (Google DeepMind) · Martina Zambelli (DeepMind) · Murilo Martins (DeepMind) · Francis Song (DeepMind) · Nicolas Heess (DeepMind) · Raia Hadsell (DeepMind) · Martin Riedmiller (DeepMind)

On the Sample Complexity of Adversarial Multi-Source PAC Learning
Nikola Konstantinov (IST Austria) · Elias Frantar (TU Vienna) · Dan Alistarh (IST Austria & NeuralMagic) · Christoph H. Lampert (IST Austria)

Inducing and Exploiting Activation Sparsity for Fast Inference on Deep Neural Networks
Mark Kurtz (Neural Magic) · Justin Kopinsky (Neural Magic) · Rati Gelashvili (Neural Magic) · Alexander Matveev (Neural Magic) · John Carr (Neural Magic) · Michael Goin (Neural Magic) · William Leiserson (Neural Magic) · Sage Moore (Neural Magic) · Nir Shavit (Neural Magic) · Dan Alistarh (IST Austria & NeuralMagic)

Constructive universal distribution generation through deep ReLU networks
Dmytro Perekrestenko (ETH Zurich) · Stephan Müller (ETH Zurich) · Helmut Bölcskei (ETH Zurich)

Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks
Francesco Croce (University of Tuebingen) · Matthias Hein (University of Tübingen)

Multiclass Neural Network Minimization via Tropical Newton Polytope Approximation
Georgios Smyrnis (National Technical University of Athens) · Petros Maragos (National Technical University of Athens)

Finding trainable sparse networks through Neural Tangent Transfer
Tianlin Liu (Friedrich Miescher Institute) · Friedemann Zenke (Friedrich Miescher Institute)

Towards a General Theory of Infinite-Width Limits of Neural Classifiers
Eugene Golikov (Moscow Institute of Physics and Technology)

Controlling Overestimation Bias with Truncated Mixture of Continuous Distributional Quantile Critics
Arsenii Kuznetsov (Samsung) · Pavel Shvechikov (Samsung Artificial Intelligence Center) · Alexander Grishin (Higher School of Economics) · Dmitry Vetrov (Higher School of Economics, Samsung AI Center Moscow)

Learning to Learn Kernels with Variational Random Features
Xiantong Zhen (Inception Institute of Artificial Intelligence) · Haoliang Sun (Shandong University) · Yingjun Du (University of Amsterdam) · Jun Xu (Nankai University) · Yilong Yin (Shandong University) · Ling Shao (Inception Institute of Artificial Intelligence) · Cees Snoek (University of Amsterdam)

Efficient Robustness Certificates for Graph Neural Networks via Sparsity-Aware Randomized Smoothing
Aleksandar Bojchevski (Technical University of Munich) · Johannes Klicpera (Technical University Munich) · Stephan Günnemann (Technical University of Munich)

Learning to Simulate Complex Physics with Graph Networks
Alvaro Sanchez (DeepMind) · Jonathan Godwin (DeepMind) · Tobias Pfaff (DeepMind) · Rex (Zhitao) Ying (Stanford University) · Jure Leskovec (Stanford University) · Peter Battaglia (DeepMind)

Small Data, Big Decisions: Model Selection in the Small-Data Regime
Jorg Bornschein () · Francesco Visin (Deepmind) · Simon Osindero (DeepMind)

PolyGen: An Autoregressive Generative Model of 3D Meshes
Charlie Nash (DeepMind) · Yaroslav Ganin (DeepMind) · S. M. Ali Eslami (DeepMind) · Peter Battaglia (DeepMind)

XtarNet: Learning to Extract Task-Adaptive Representation for Incremental Few-Shot Learning
Sung Whan Yoon (Ulsan National Institute of Science and Technology (UNIST)) · Jun Seo (Korea Advanced Institute of Science and Technology(KAIST)) · Doyeon Kim (Korea Advanced Institute of Science and Technology) · Jaekyun Moon (KAIST)